

Piecewise-Affine Local Volatility

Calibrating a Dupire surface from sparse quotes: model, PDE, gradients, and speed

Vol-Fitter Technical Notes, No. 04

Vol-Fitter

June 29, 2026

Abstract

This note documents the Vol-Fitter’s full local-volatility model: a continuous piecewise-affine (P_1 finite-element) local-variance surface $\nu_\theta(t, x)$, calibrated directly to sparse vanilla and variance-swap quotes through the forward Dupire PDE as an implicit pricing map. Unlike the slice models (Notes 01–03) this is a genuine two-dimensional local-volatility surface: strict positivity of the nodal variances makes the implied prices butterfly- and calendar-clean up to scheme noise, which the model *measures* rather than assumes. We derive the normalized forward Dupire equation, the P_1 parametrization and its positivity guarantee, the implicit-Euler discretization and observation operator, the variance-swap log-contract replication and its faster source-PDE alternative, and the calibration objective with its roughness, convex-wing and front-tie regularizers. We then give the sensitivity PDE and the adjoint gradient, and — the engineering heart of the note — the numerical optimizations that cut the cold fit $\sim 3\text{--}6\times$ (from $\sim 86\text{ s}$ to $\sim 15\text{--}30\text{ s}$, scaling with grid size), with recalibrations near-instant: a compiled vectorized-Thomas Dupire march ($\sim 6.5\times$), a stall-based early-stop, a parametric Dupire warm seed, and the matrix-free Gauss–Newton solver that avoids the dense SVD. A synthetic round trip recovers a known surface to 0.6 vol bps, and the production fit on a static Bloomberg fixture reaches 2.8 vol bps on SPY and 11.7 on NVDA. The note closes with the short-dated-fit fixes and the convex-wing regression that the benchmark caught.

Contents

1	The idea	3
2	The forward Dupire equation	3
3	The piecewise-affine surface	3
3.1	Parametrization	3
3.2	Grid design	4
4	The pricing map	4
4.1	Discretization	4
4.2	No-arbitrage diagnostics	4
5	Variance swaps	4
6	The calibration objective	5

7	Sensitivities and the adjoint	5
7.1	The sensitivity PDE	5
7.2	The adjoint gradient	5
8	Solvers and the speed story	6
8.1	Where the time goes	6
8.2	What worked	6
9	Short-dated calibration	7
9.1	The convex-wing regression	7
10	Worked examples	7
10.1	Synthetic round trip	7
10.2	The Bloomberg benchmark	8
11	What is original, and limitations	9
A	Hyperparameter atlas	9
B	Performance notes	10

1 The idea

A slice model fits one expiry at a time and stitches expiries together with a calendar constraint. A *local-volatility* model does the opposite: it posits a single function $\sigma_{\text{loc}}(t, S)$ such that, under $dS_t = \sigma_{\text{loc}}(t, S_t) S_t dW_t$ (in forward-normalized units), the model reproduces *all* expiries and strikes at once. Dupire showed such a function exists and is unique given an arbitrage-free price surface. The Vol-Fitter follows the Andreasen–Huge philosophy: rather than *differentiate* a fitted price surface to extract σ_{loc} (numerically explosive in sparse data), it *parametrizes* σ_{loc} directly and prices *forward* through the Dupire PDE, calibrating the parameters to quotes.

Heuristic

Extracting local vol by the Dupire formula $\sigma_{\text{loc}}^2 = \partial_T C / (\frac{1}{2} K^2 \partial_{KK} C)$ divides two noisy numerical derivatives of a sparse surface — a recipe for garbage in the wings. Calibrating local vol *forward* instead never differentiates the data: it asks “what positive surface, run through the PDE, best reprices the quotes?” Positivity of the surface is imposed on the *parameters*, so every iterate is arbitrage-clean by construction — the same philosophy as LQD, one dimension up.

2 The forward Dupire equation

Work in forward-normalized strike $x = K/F_T$ and let $c(T, x)$ be the normalized call. Writing local variance $\nu(t, x) = \sigma_{\text{loc}}^2$, the forward Dupire equation is

$$\partial_T c(T, x) = \frac{1}{2} \nu(T, x) x^2 \partial_{xx} c(T, x), \quad c(0, x) = (1 - x)^+. \quad (1)$$

It is a forward parabolic PDE in (T, x) with the call payoff as initial data. Two facts make it the right object. First, $\partial_{xx} c \geq 0$ is exactly the butterfly (density) condition, and $\nu \geq 0$ keeps the equation parabolic, so a positive ν *propagates* non-negativity of the density forward in maturity. Second, monotonicity of c in T (the calendar condition) follows from (1) whenever $\nu \geq 0$, since the right-hand side has the sign of $\partial_{xx} c \geq 0$. Thus a strictly positive local-variance surface is *structurally* free of both static arbitrages, up to the discretization noise the model measures (section 4.2).

3 The piecewise-affine surface

3.1 Parametrization

The local variance is a continuous piecewise-affine (P_1) function on a tensor grid of $n_t \times n_x$ vertices (τ_i, ξ_j) ,

$$\nu_\theta(t, x) = \sum_{i,j} \theta_{ij} \phi_{ij}(t, x), \quad (2)$$

where ϕ_{ij} are the bilinear (or barycentric/Delaunay) hat functions and $\theta_{ij} = \nu(\tau_i, \xi_j) \geq 0$ are the *nodal local variances* — the calibration parameters. The flat vector is $\theta = \text{vec}(\theta)$, time-major.

Proposition 1 (Nodal positivity $\Rightarrow$ surface positivity). *A P_1 function is a convex combination of its vertex values on each cell. Hence if $\theta_{ij} \geq 0$ for all i, j , then $\nu_\theta(t, x) \geq 0$ everywhere, and $\nu_\theta > 0$*

if the nodal values are bounded below by a positive constant.

This is why the calibration enforces only *box bounds* on θ (default local-variance bounds $[0.005, 0.20]$, i.e. local vol in $[7\%, 45\%]$): positivity of the whole surface reduces to positivity of finitely many numbers. Beyond the lowest strike vertex the surface extrapolates the local variance *linearly* with a slope set by the first interior cell (`left_extrap_a`; Note 09), rising toward $x \rightarrow 0$; the right wing is flat-clamped.

3.2 Grid design

The vertex grid is delta-spaced in strike and floored per expiry so that even a narrow short-dated smile lands enough vertices on its curvature (`gridXNodes`, `gridXMinPerExpiry`; section 9). The time vertices (`gridTNodes`) carry the term structure. The grid is deliberately modest — a dozen strike vertices, ten in time — because each vertex is a free parameter and the PDE solve cost scales with the grid; the regularizer (section 6) supplies the smoothness that a coarse grid would otherwise lack.

4 The pricing map

4.1 Discretization

The forward map (1) is solved on a uniform spatial grid with step $\Delta x = 0.01$ on $x \in [0, \sim 2.5]$ (~ 250 nodes) by fully implicit (backward) Euler in maturity with step $\Delta T = 0.01$ (the local variance entering each step at the new time level t_{n+1}), optionally Rannacher-started with a couple of implicit start-up steps before Crank–Nicolson. (A separate log-strike validator, appendix A, runs the finer fixed grid.) Each step solves a tridiagonal M -matrix system

$$(I - \Delta T A_n) U^{n+1} = U^n, \quad (3)$$

where A_n encodes $\frac{1}{2}\nu x^2 \partial_{xx}$ at step n . The matrix is an M -matrix (diagonally dominant, non-positive off-diagonals), so the scheme is unconditionally stable and monotone — it cannot create spurious negative densities. The observation operator reads the normalized call at each quote’s (T, x) off the solution by interpolation.

4.2 No-arbitrage diagnostics

Because the surface is only *up to scheme noise* arbitrage-free, the model measures the residual: a butterfly tolerance on $\partial_{xx}c$ ($\gtrsim -10^{-6}$ density) and a calendar tolerance on $\partial_T c$ ($< -10^{-9}$ price decrement). The model is gated behind these diagnostics rather than assuming cleanliness — a deliberate design choice.

5 Variance swaps

A variance-swap quote is the fair strike of realized variance, which for a diffusion equals the expected total variance to T . Two pricers are available. The *static* method uses the log-contract replication

$$I(T) = 2 \int_0^1 \frac{P(T, x)}{x^2} dx + 2 \int_1^\infty \frac{C(T, x)}{x^2} dx, \quad (4)$$

a trapezoid over the priced surface (half-width and node count `VS_HALF_WIDTH`, `VS_POINTS`). The faster *source-PDE* method (default-off, gated by `varSwapMethod`) solves one backward source PDE $\partial_t g + \frac{1}{2}\nu x^2 \partial_{xx} g + \nu = 0$, $g(T, \cdot) = 0$, and reads $I(T) = g(0, 1)$ — with analytic sensitivities $\partial I / \partial \theta$ from the same multi-RHS machinery, and without the $x_{\max}$ -dependence of the replication tail. Note 08 develops both in full.

6 The calibration objective

The fit minimizes a stacked least-squares residual over θ ,

$$\min_{\theta \in [\nu_{lo}, \nu_{hi}]} \frac{1}{2} \left(\underbrace{\|r_{\text{opt}}\|^2}_{\text{quotes}} + \underbrace{\|r_{\text{vs}}\|^2}_{\text{var-swaps}} + \underbrace{\lambda \|L(\theta - \theta_{\text{ref}})\|^2}_{\text{roughness}} + \underbrace{w_{\text{cvx}} \|r_{\text{cvx}}\|^2}_{\text{convex wing}} + \underbrace{w_{\text{ft}} \|r_{\text{ft}}\|^2}_{\text{front tie}} \right), \quad (5)$$

with the blocks:

- **Option residuals** $r_{\text{opt},q} = (P_q(\theta) - y_q) / \eta_q$ in vega-normalized forward-call price units (mid), or the bid–ask/haireut band hinge with a small mid anchor when band edges are present (Note 07).
- **Var-swap residuals** $r_{\text{vs}} = (I(\theta) - z) / \zeta$ (Note 08).
- **Roughness** $\sqrt{\lambda} L(\theta - \theta_{\text{ref}})$ with L the spacing-aware second-difference operator and θ_{ref} a reference surface; the time-vs-strike balance is `gridRegRho`. This is the smoothness that lets a coarse grid generalize.
- **Convex wing** hinge: a one-sided penalty enforcing local convexity of the wing, confined to vertices *below the deepest quote* (the extrapolation tail only — see section 9.1).
- **Front tie**: a light penalty tying the unconstrained $t = 0$ time row to the first (data-identified) expiry row across strikes, anchoring the front end where the shortest expiry has little coverage.

7 Sensitivities and the adjoint

7.1 The sensitivity PDE

Each parameter θ_ℓ enters (1) linearly through ν_θ , so the price sensitivity $s_\ell = \partial c / \partial \theta_\ell$ obeys the same parabolic PDE with a *source*:

$$\partial_T s_\ell = \frac{1}{2} \nu_\theta x^2 \partial_{xx} s_\ell + \frac{1}{2} \phi_\ell x^2 \partial_{xx} c, \quad s_\ell(0, \cdot) = 0. \quad (6)$$

Discretized, every s_ℓ shares the *same* tridiagonal factor as the value solve (3), differing only in the source — so the Jacobian columns are a single multi-RHS back-substitution against one factorization. This is the forward-sensitivity route, $O(N_t N_x m)$ for m active parameters.

7.2 The adjoint gradient

When only the *gradient* of the scalar objective is needed (not the full Jacobian), the discrete adjoint computes it in one forward value solve plus one backward adjoint sweep, $O(N_t N_x)$ *independent of m* . The Vol-Fitter uses the forward sensitivities for the Gauss–Newton Jacobian (it needs the columns) and keeps the adjoint as the route for very large vertex counts.

8 Solvers and the speed story

The cold fit of (5) was, at ~ 533 vertices, an ~ 86 second affair hitting the 200-evaluation cap — the single heaviest computation in the app. Driving it down is most of the engineering in this model, and the lessons are worth stating because several plausible ideas *failed*.

8.1 Where the time goes

On a cold fit the cost splits roughly as: optimizer linear algebra / dense SVD $\sim 52\%$, PDE sensitivity march $\sim 32\%$, Jacobian assembly $\sim 14\%$, value solve $\sim 2\%$. No single per-evaluation lever moves the total, because the cost is distributed — which is why four separate “faster per evaluation” attempts (coarse grid, first Numba try, Rannacher, and the first matrix-free GN) each won only $\sim 1.1\text{--}1.2\times$ and were shelved.

8.2 What worked

Performance

The cold fit is now $\sim 3\text{--}6\times$ faster than the original banded baseline (scaling with grid size); recalibrations are near-instant via warm start. The levers, in order of impact:

1. **Stall-based early-stop.** Track the option-block misfit and stop at the best iterate once it stalls. Because fewer evaluations multiply march, assembly *and* optimizer together, this scales the *whole* fit: measured $1.45\times$ (SPY) to $3.3\times$ (NVDA), at $+0.1\text{--}0.25$ vol bp.
2. **Compiled vectorized-Thomas march.** A @njit no-pivot Thomas solver with the sensitivity columns as the contiguous SIMD inner loop and a fused source: $\sim 6.5\times$ the LAPACK banded march, numerically exact to $\sim 10^{-15}$. The first Numba attempt (column-outer scalar Thomas) won only $1.2\times$ — the real lever was the loop order.
3. **Matrix-free Gauss–Newton.** The default solver avoids `trf`’s dense SVD (52% of an evaluation) by solving the GN step with preconditioned LSMR on matrix-free $Jv/J^\top w$ products. Once the march is cheap, the no-SVD evaluations win: $\sim 1.3\text{--}1.65\times$ over TRF. Gated to the smooth mid-fit target.
4. **Parametric Dupire warm seed.** Seed θ from the parametric surface’s local variance ($\sim 1.3\text{--}1.8\times$ on cold fits) and from the previous surface on recalibration ($\sim 38\times$).

The order-of-magnitude wins are spent; remaining levers are incremental.

Caution

Removing the dense SVD made fits *slower* at ≤ 440 vertices, because there the bottleneck is the PDE march, not the SVD — the SVD wall is a $\gtrsim 1000$ -vertex (future non-tensor “bowtie” grid) phenomenon. The clean synthetic benchmark (zero residual, in bounds) hid this: GN converged in 8 evaluations there but did *not* converge within the cap on stiff real data. The honest lesson: measure on real chains, and distrust a synthetic-only speed-up.

9 Short-dated calibration

A genuine one-week smile fit *catastrophically* (108 vol bp RMS vs the parametric ~ 47) while normal expiries fit well. A measure-first diagnosis (per-expiry side-channel counting in-range vertices, vega flooring and PDE steps) ruled out the usual suspects — vega floor, time steps, local-vol cap, prior leakage — and isolated *strike under-resolution*: the delta strike axis is sized to the *longest* expiry and clipped to the global range, so a narrow short smile landed only ~ 3 vertices on its sharpest curvature. Two fixes:

1. **Per-expiry coverage floor** (`gridXMinPerExpiry`, default 8): after the delta axis is built, split the widest in-range gaps until each expiry has at least $m_{\min}$ vertices inside *its own* traded range — adding nodes only to under-covered short fronts.
2. **Short-expiry-aware PDE step**: refine the shared PDE x -step to $0.3\times$ the smallest ATM $\sigma\sqrt{\tau}$, snapped so $x = 1$ stays a node, clamped to $[1/400, 0.01]$. Normal surfaces stay on 0.01 (byte-identical).

The one-week fit fell from $108 \rightarrow 23.5$ vol bp (now *better* than the parametric), with the long expiries byte-identical.

9.1 The convex-wing regression

The benchmark also caught a subtler bug. The convex-wing constraint originally selected *every* vertex in the deep wing regardless of data; at a user's saved `gridXNodes = 20` it stacked convexity penalties onto densely-quoted put strikes and forced the *wrong* wing on low-vol SPY (NVDA's naturally convex wing had hidden it). The fix confines `convex_cols` to vertices *below the deepest quote* — the extrapolation tail only. SPY went $25.7 \rightarrow 2.6$ vol bp. The lesson echoes Note 09: a shape constraint must not be imposed where the data already speaks.

10 Worked examples

10.1 Synthetic round trip

A known skewed local-variance surface is priced through the forward PDE, quotes are generated at four expiries, and the surface is recovered from a *flat* seed by the production calibrator. Figure 1 shows the recovered surface and the implied-vol fit: the calibrator recovers the truth to a maximum of 0.6 vol bps in 33 evaluations (1.01s wall, 28 vertices), with the skew (higher local vol at low strike) faithfully reconstructed.




Figure 1: Synthetic round trip through the real calibrator: the skewed surface is recovered to 0.6 vol bps from a flat seed.

10.2 The Bloomberg benchmark

On the static Bloomberg fixture the production fit reaches the quality in table 1 and fig. 2: 2.8 vol bps surface RMS on SPY and 11.7 on NVDA, with no node at a bound on the easy expiries. The harder NVDA short-dated expiries carry the larger error — the regime section 9 addresses.

Table 1: Production local-vol fit on the Bloomberg fixture (fresh run).

Name	surface RMS (bp)	worst expiry (bp)
SPY	2.8	4.4
NVDA	11.7	12.4




Figure 2: Per-expiry weighted RMS for SPY and NVDA on the fixture.

11 What is original, and limitations

The P_1 -positivity guarantee and the forward-calibration philosophy are Andreasen–Huge; the contributions here are engineering depth. The *vectorized-Thomas Numba march* (loop order as the lever, $6.5\times$), the *stall early-stop* that scales the whole fit, the *matrix-free GN* correctly scoped to where the SVD actually dominates, the *source-PDE var-swap* with analytic sensitivities, and the *measure-first* short-dated diagnosis are each non-obvious and were validated against a real benchmark that twice caught regressions a synthetic would have missed. Limitations: the model is a pure diffusion (no jumps), so a true short-dated event spike is only approximated; the tensor grid wastes vertices in the corners (the future non-tensor bowtie grid is where GN’s no-SVD advantage would finally dominate); and the fit, though much faster, remains the heaviest single computation in the app.

A Hyperparameter atlas

Table 2: Local-volatility hyperparameters.

Knob	Default	Role
<i>Surfaced (OptionsSettings)</i>		
<code>gridXNodes</code>	12	Strike vertices of the variance surface.
<code>gridXMinPerExpiry</code>	8	Min in-range strike vertices per expiry (section 9).
<code>gridTNodes</code>	10	Maturity vertices.
<code>gridRegLambda λ</code>	10^{-2}	Roughness penalty strength.
<code>gridRegRho ρ</code>	1.0	Time-vs-strike roughness balance.
<code>convexWingWeight</code>	10^3	Convex-wing hinge weight (tail only).
<code>frontTieWeight</code>	10^{-2}	Front-end tie-to-ATM weight.
<code>lvVolCapMult</code>	3.0	Local-vol cap as a multiple of max implied vol.
<code>leftWingSlopeMult</code>	1.5	Left local-variance extrapolation slope multiple.
<code>varSwapMethod</code>	<code>static</code>	Var-swap pricer: replication or source-PDE.
<code>lvSolver</code>	<code>gn</code>	Gauss–Newton (default) or <code>trf</code> .
<code>lvFastKernel</code>	<code>true</code>	Numba vectorized-Thomas march.
<code>lvEarlyStop</code>	<code>true</code>	Stall-based early stop.
<code>timeScheme</code>	<code>implicit</code>	Implicit Euler or Rannacher/CN.
<i>Hidden (affine-calibration PDE / solver)</i>		
affine PDE step	$\Delta x = 0.01$	Spatial step on $x \in [0, \sim 2.5]$ (~ 250 nodes).
affine time step	$\Delta T = 0.01$	Implicit-Euler maturity step.
<code>rannacher_steps</code>	2	Implicit start-up steps before CN (opt-in scheme).
short-dated step	$0.3 \sigma \sqrt{\tau}$	Refined Δx floor, clamped $[1/400, 0.01]$ (section 9).
<code>VAR_FLOOR</code>	10^{-6}	Dupire local-variance floor.
variance bounds	$[0.005, 0.20]$	Nodal local-variance box (vol $[7\%, 45\%]$).
<code>x_scale</code>	<code>jac</code>	Optimizer column scaling.
<code>xtol/ftol/gtol</code>	10^{-8}	Optimizer tolerances.
<code>max_nfev</code>	200	Evaluation cap.

Knob	Default	Role
<code>stall_window/rtol</code>	12–18 / $\sim 3\text{--}5 \times 10^{-3}$	Early-stop window/tolerance.
<i>Hidden (separate log-strike validator, <code>pde.py</code>)</i>		
<code>DEFAULT_N_K</code>	1201	Validator spatial nodes.
<code>DEFAULT_DT_MAX</code>	1/400	Validator max time step.
<code>N_RANNACHER</code>	4	Validator Rannacher quarter-steps.
<code>SPAN_SD</code> / <code>SPAN_MIN</code>	7.5 / 0.6	Validator grid span (sd / floor).

B Performance notes

1. **Numba march** (`lvFastKernel`). No-pivot factor-once Thomas with the m sensitivity columns as the contiguous inner SIMD loop and fused source: $6.1\text{--}6.9\times$ vs LAPACK `dgbstv`, exact to $\sim 10^{-15}$. Graceful banded fallback if Numba is absent.
2. **Early stop** (`lvEarlyStop`). Best-iterate stop on option-block stall; scales the whole fit, $1.45\times$ (SPY) to $3.3\times$ (NVDA), $+0.1\text{--}0.25$ vol bp. `stall_window=0` is byte-identical.
3. **Matrix-free GN** (`lvSolver=gn`). Preconditioned LSMR on $Jv/J^\top w$; avoids the dense SVD. $1.3\text{--}1.65\times$ over TRF at tensor-grid sizes; gated to the smooth mid target.
4. **Warm starts**. Previous-surface seed ($\sim 38\times$ recal) and parametric Dupire cold seed ($1.3\text{--}1.8\times$).
5. **Shelved (documented)**: coarse grid (biases θ), first Numba attempt (wrong loop order), Rannacher ($1.1\times$ + arb risk), GN for the non-smooth band objective, thread/process parallelism (GIL). Four distributed-cost dead-ends confirming the cost is spread, not localized.

References

- [1] B. Dupire. Pricing with a smile. *Risk*, 7(1):18–20, 1994.
- [2] J. Andreasen and B. Huge. Volatility interpolation. *Risk*, March 2011.
- [3] J. Gatheral. *The Volatility Surface*. Wiley, 2006.
- [4] M. Giles and R. Carter. Convergence analysis of Crank–Nicolson and Rannacher time-marching. *J. Comput. Finance*, 9(4), 2006.