

Wings, Lee Bounds and Tail Control

How four models keep their extrapolated wings arbitrage-admissible

Vol-Fitter Technical Notes, No. 09

Vol-Fitter

July 11, 2026

Abstract

A smile model earns its keep in the *wings* — the region beyond the quoted strikes, where the fit becomes pure extrapolation and where mispriced tails leak into var-swaps, risk-reversals and butterflies. This cross-cutting note collects the wing treatment shared by the Vol-Fitter’s models. The universal constraint is Lee’s moment formula: the asymptotic slope of total implied variance, $\beta = \limsup_{k \rightarrow \pm\infty} w(k)/|k|$, cannot exceed 2 and is tied to the moment critical exponents by the explicit map $\beta = \psi(p)$. We show how each model realizes or enforces it — LQD structurally, SVI by a soft cap, the Multi-Core Sigmoid (MCS) by wing-neutral hats — and prove the zero-wing inheritance the MCS design relies on. The finite wing is a separate battle: the asymptotic slope can be admissible while a hat still dents the Durrleman density out where quotes are sparse, which is why the default-on put-wing regularizer samples *g past* the traded range (measured: worst violation $-10.2 \rightarrow -0.008$ on the synthetic slice, a $\sim 400\times$ median repair on real illiquid wings, byte-identical on clean slices). A case file replays the backtest finding that motivated it and the ablation verdict that the de-Americanization repair and the penalty are complementary, not redundant (92 $\rightarrow$ 25 bp with the input repair alone; 749 bp with the output penalty alone; 225 bp with both). The note closes with the arc’s centerpiece: the *confinement principle* — which no-arbitrage constraints must be sampled only where data lives, and which must extend into the wings — with its four production instances.

Contents

1	The production problem	3
2	The universal constraint: Lee’s moment formula	3
3	Four models, four wing mechanisms	4
3.1	LQD: Lee-consistent by construction	4
3.2	SVI: linear wings with a soft cap	5
3.3	Multi-Core Sigmoid (MCS): inherited wings	5
3.4	Local volatility: controlled linear extrapolation	5
3.5	Side by side	5
4	Policing the finite wing: the MCS put-wing regularizer	7
5	The confinement principle	8

6	What is original here	9
7	Traceability	9
A	Hyperparameter atlas	10
B	Performance notes	10
C	Reference implementation	11

1 The production problem

Nobody quotes the far wings, yet everything downstream reads them: var-swap replication integrates $1/K^2$ -weighted prices into the tails, risk metrics difference the wings, and the graph extrapolator (Note 14) inherits whatever tail the fitted model draws. Wing errors are invisible in the fit RMS — there are no quotes to miss — and show up later as a mispriced tail somewhere else. Two distinct things can go wrong: the *asymptotic* slope can violate Lee’s model-free bound (an arbitrage against the existence of any law), and the *finite* wing can lose density positivity even with an admissible slope (a butterfly arbitrage in the extrapolated region).

Invariant

1. Every displayed wing is Lee-admissible: $\beta \leq 2$, structurally where possible, by a soft cap otherwise.
2. Single-slice (butterfly) constraints *extend* into the wings; two-slice (calendar) constraints are *confined* to the traded range (section 5).
3. Every wing guard is strictly additive: on an already-clean slice it is byte-identical to no guard.
4. The tail is a property of the model’s extrapolation law, not of the quotes; where the desk needs control (LV deep puts), the slope is an explicit knob or a calibrated variable.

2 The universal constraint: Lee’s moment formula

Let $w(k, T) = \sigma_{BS}^2(k, T) T$ be total implied variance at log-moneyness $k = \log(K/F_T)$.

Theorem 1 (Lee moment formula [1]). *Define the wing slopes*

$$\beta_R = \limsup_{k \rightarrow +\infty} \frac{w(k, T)}{k}, \quad \beta_L = \limsup_{k \rightarrow -\infty} \frac{w(k, T)}{|k|},$$

and let $p_R^* = \sup\{p : \mathbb{E}[S_T^{1+p}] < \infty\}$ and $q_L^* = \sup\{q : \mathbb{E}[S_T^{-q}] < \infty\}$ be the moment critical exponents. Then

$$\beta_R = \psi(p_R^*), \quad \beta_L = \psi(q_L^*), \quad \psi(p) = 2 - 4(\sqrt{p^2 + p} - p). \quad (1)$$

Lemma 1 (Properties of the Lee map). *ψ is a strictly decreasing bijection from $[0, \infty]$ onto $[0, 2]$, with $\psi(0) = 2$, $\psi(\infty) = 0$, and explicit inverse $p = \frac{1}{2\beta} + \frac{\beta}{8} - \frac{1}{2}$. In particular $0 \leq \beta \leq 2$ for every arbitrage-free smile — the ceiling $\beta = 2$ is model-free.*

Proof. Write $\psi(p) = 2(\sqrt{p+1} - \sqrt{p})^2$: the parenthesis is positive and strictly decreasing (its derivative $\frac{1}{2}(1/\sqrt{p+1} - 1/\sqrt{p}) < 0$), with limits 1 at $p = 0$ and 0 at ∞ ; squaring and doubling gives the range $(0, 2]$ plus the endpoint $\psi(\infty) = 0$. Solving $\beta = 2(\sqrt{p+1} - \sqrt{p})^2$ for p gives the stated inverse by elementary algebra. $\square$

The map, and the way LQD reads its slopes off the endpoint scales $A_L = e^{g(0)}$, $A_R = e^{g(1)}$ of Note 01, are a few lines. Appendix C verifies this distilled listing against the production `lee_psi` and `lee_slopes` to 10^{-15} (the production function takes the LQD parameter object and computes the endpoint scales itself):

Listing 1: Lee’s ψ and the LQD structural wing slopes (distilled from `models/lqd/basis.py`).

```

1 def lee_psi(p):                                # beta = psi(p); decreasing, psi(0)=2
2     return 2 - 4*(np.sqrt(p*p + p) - p)
3
4 def lee_slopes(A_L, A_R):                       # exponential tails -> exact exponents
5     beta_L = lee_psi(1/A_L)                     # left: p* = 1/
        A_L
6     beta_R = 2.0 if A_R >= 1 else lee_psi(1/A_R - 1) # right: p* = 1/
        A_R - 1
7     return beta_L, beta_R

```




Figure 1: The Lee map $\beta = \psi(p)$: heavier tails (smaller moment exponent p) force steeper variance wings, capped at the model-free ceiling $\beta = 2$.

Heuristic

The wing slope and the tail heaviness are two faces of one coin. A fat right tail (few finite positive moments) *must* show up as a steep right variance wing; a model that draws a flat wing over a fat-tailed density is internally inconsistent and will misprice far-OTM calls. Lee’s bound is not a nuisance constraint but a consistency condition between the smile and the law it implies — which is why the LQD model, built from the law outward, satisfies it automatically.

3 Four models, four wing mechanisms

3.1 LQD: Lee-consistent by construction

The LQD model (Note 01) parametrizes the log quantile density with logarithmic endpoints, giving *exponential* return tails with scales A_L, A_R read off the endpoint coefficients. The moment exponents are exact, $p_R^* = 1/A_R - 1$ and $q_L^* = 1/A_L$, and the wing slopes follow from theorem 1:

$$\beta_L = \psi(1/A_L), \quad \beta_R = \psi(1/A_R - 1), \quad A_R < 1. \quad (2)$$

No penalty enforces this — it is structural. The only wing-related constraint is the hard $A_R < 1$ (finite forward), handled by Note 01’s soft barrier. The benchmark fit yields $(\beta_L, \beta_R) = (0.0971, 0.0358)$.

3.2 SVI: linear wings with a soft cap

Raw SVI (Note 02) has exactly linear wings, $w(k) \sim b(1 \mp \rho)|k|$, so $\beta_L = b(1 - \rho)$ and $\beta_R = b(1 + \rho)$ in closed form ($= (0.1093, 0.0364)$ on the benchmark parameters). Lee’s bound is the explicit inequality $b(1 + |\rho|) \leq 2$, enforced as a soft hinge $\max(b(1 + |\rho|) - \beta_{\max}, 0)$ at weight 10^3 alongside the min-variance penalty (cap `leeSlopeMax` = 2). On any admissible smile both penalties are exactly zero; they only fence the optimizer out of the inadmissible region.

3.3 Multi-Core Sigmoid (MCS): inherited wings

The Multi-Core Sigmoid (MCS) slice (Note 03) is a convex base plus up to two *zero-wing* hats.

Proposition 1 (Zero-wing inheritance). *Let $v(z) = v_{\text{base}}(z) + \sum_r h_r(z)$ in standardized moneyness z , where each hat h_r and its first two derivatives vanish as $z \rightarrow \pm\infty$. Then the asymptotic variance slopes of v equal those of the base, $S_0 \mp 2K_0/\kappa_{P,C}$, for any number of hats and any hat parameters; in total-variance k -space, with $k = \sigma_{\text{ref}}\sqrt{t}z$ and $w = tv$,*

$$\beta_{L,R} = \frac{\sqrt{t}}{\sigma_{\text{ref}}} \left| S_0 \mp \frac{2K_0}{\kappa_{P,C}} \right|. \quad (3)$$

Proof. $v'(z) = v'_{\text{base}}(z) + \sum_r h'_r(z) \rightarrow v'_{\text{base}}(z)$ as $|z| \rightarrow \infty$ since each $h'_r \rightarrow 0$; the base is asymptotically linear with the stated slopes. The k -space conversion is the chain rule through $k = \sigma_{\text{ref}}\sqrt{t}z$ and $w = tv$. $\square$

On the benchmark, $(\beta_L, \beta_R) = (0.1057, 0.0297)$. Adding cores cannot change the wings — the entire point of the zero-wing design. Two caveats bound what the proposition buys. First, it pins only the asymptotic *slope*: a hat can still dent the Durrleman g below zero at *finite* $|z|$, out in the sparsely quoted put wing — a residual failure mode that is real, measured, and policed by the put-wing regularizer of section 4. Second, wing-*neutral* is not wing-*admissible*: the inheritance theorem preserves whatever slopes the base fitted, and nothing in the MCS calibration bounds forces those base slopes to satisfy Lee’s sign and cap conditions, $0 \leq \beta_{L,R} \leq 2$. Unlike LQD (structural) and SVI (soft cap), MCS admissibility of the wings is *diagnosed*, not enforced.

3.4 Local volatility: controlled linear extrapolation

The piecewise-affine local-vol surface (Note 04) has no closed-form implied-vol wing. Beyond the deepest strike vertex it extrapolates the *local* variance linearly with slope `left_extrap_a` times the first interior cell’s slope; the right wing is flat-clamped. The knob mapping is deliberate: the slope is 0 (flat clamp) on the hot path, seeded from `leftWingSlopeMult` = 1.5 when the convex-wing hinge is on, and becomes a *free calibration variable* (bounds $[0, 20]$) exactly when a var-swap quote is present to identify the deep-put tail — the one instrument that genuinely observes it (Note 08).

3.5 Side by side

Figure 2 fits the same narrow window with the three parametric models and extrapolates far past it. The wings agree in sign (the equity put skew, $\beta_L > \beta_R$) and order of magnitude but diverge in detail — each tail is a property of *its* extrapolation law, not of the identical quotes. That divergence is not a defect to be averaged away: past the last quote a wing is a modelling *choice*,

not an estimate — the far wing is exactly where a model extrapolates with confidence it does not have — so what the fitter guarantees is not agreement but a stated, tested *wing contract* per model. LQD draws exponential return tails with structural slopes $\psi(1/A_L), \psi(1/A_R - 1)$; SVI draws linear variance wings $b(1 \mp \rho)$ under the soft Lee cap; the MCS base alone owns the tails, hats neutral (proposition 1); LV extrapolates local variance flat on the call side and controlled-linearly on the put side. Choosing a model *is* choosing an extrapolation law — the differences are the point. The shared promise is that each contract's Lee status is *stated and checked*, with per-model teeth: structural for LQD, a soft cap for SVI, bounded local variance for LV, and *diagnosis only* for the MCS base slopes. Table 1 collects the analytic slopes.




Figure 2: Three models fitted to the same window (shaded) and extrapolated. Inside the window they agree to vol bps; outside, each draws its own Lee-admissible wing. The divergence *is* the model choice.

Table 1: Analytic asymptotic total-variance wing slopes $\beta = \lim w(k)/|k|$ on the benchmark window; all far below the Lee ceiling 2.

Model	β_L (put)	β_R (call)	mechanism
LQD	0.0971	0.0358	structural via $\psi(1/A)$
SVI	0.1093	0.0364	linear $b(1 \mp \rho)$, soft Lee cap
MCS	0.1057	0.0297	base wings, hats neutral (proposition 1)

4 Policing the finite wing: the MCS put-wing regularizer

An admissible slope does not make an admissible wing. The backtest’s arb census located the MCS’s failures precisely: 64% of butterfly violations sit in the *put* wing (median worst-violation location $z \approx -3.2$), against 4% at the money — hats fitted to sparse illiquid quotes buy in-sample RMS by denting the density where no quote objects. The shipped defence (the R6 fix) has three parts:

1. **A capacity cap.** `nCores` is clamped to ≤ 2 at the schema (and to $(n_{\text{quotes}} - 6)/4$ in the calibrator): most historical violations came from high-core fits chasing noise.
2. **A one-sided Durrleman hinge.** The refine stage carries residual rows $\sqrt{\lambda_w} \max(-g(z_m), 0)$ on a 49-point grid extending *two standardized-moneyness units past the traded range* on both sides, with the put side weighted $2\times$ (where the census says the risk lives). The weight is $\lambda_w = (\text{sivWingPenaltyPct}/100) \cdot 10^3$, default 100%.
3. **A hybrid Jacobian.** The g rows are finite-differenced (49 points, one column at a time) while the quote/ridge/calendar blocks keep their analytic Jacobian — the penalty costs little and perturbs nothing when inactive.

Measured: the synthetic arb slice’s worst g went $-10.2 \rightarrow -0.008$; on real illiquid wings (EEM) the median worst violation improved $\sim 400\times$ ($-7.86 \rightarrow -0.019$) at +79 bp in-sample cost; on clean liquid slices the penalty rows are zero and the fit is byte-identical to 3×10^{-17} . Note that sampling g *past* the quoted range here is correct, not a violation of confinement — see section 5.

Case file: the MCS that invented a put wing

Setup. The three-regime backtest sweep (F4): the Multi-Core Sigmoid on illiquid single names, where the put wing has a handful of wide quotes.

Failure mode. MCS-2 improved in-sample RMS over MCS-0 — and carried genuine butterfly arbitrage on 76% of arb-prone nodes, almost all in the put wing. The model was buying RMS with a fake weekend wing: a hat parked where no quote could object.

Diagnosis. The analytic (FD-free) arb metric (R2) separated real violations from finite-difference noise: with exact w', w'' per model, the apparent LQD “violations” vanished entirely ($28\% \rightarrow 0$ — structural positivity confirmed) while MCS-2’s survived (76%, genuine). The census then localized them: put wing, $z \approx -3.2$.

Fix. The 2-core cap plus the put-weighted Durrleman hinge above (output side), composing with the de-Americanization wing repair of Note 05 (input side).

Verdict (ablation, test-locked). On the arb-prone illiquid population, with neither fix: worst g -30.4 , arbitrage on 100%, in-sample 92 bp. Input repair alone: arbitrage cut $\sim 3\times$ and RMS improved to 25 bp — the MCS had been chasing arbitrated de-Am inputs. Output penalty alone: arbitrage essentially eliminated ($g \geq -0.02$) but at 749 bp — brutal, because the model must fight its own corrupted inputs. Both: the penalty’s arbitrage removal at 225 bp. (Those are medians across the 38 arb-prone census nodes; one real node makes it concrete — EFA in the Aug-2024 vol spike, 11 days out, 22 quotes: 75 bp at $g = -116$ with neither defence, 18 bp but $g = -12$ with the repair alone, arb-free at 726 bp with the penalty alone, arb-free at 34 bp with both.) *Complementary, not redundant:* the input repair makes the output penalty affordable. On liquid names both are byte-identical no-ops.

5 The confinement principle

Four production incidents — the phantom calendar violations (Note 10’s case file), the local-vol convex-wing regression (Note 04), the de-Am global repair revert (Note 05), and the MCS put-wing penalty above — resolve into one rule.

Heuristic

The confinement principle. *A constraint that compares two curves must be sampled only where data pins both; a constraint intrinsic to one curve must hold everywhere, wings included.* A two-curve constraint evaluated where both sides are extrapolation can only compare two fictions — any “violation” there is an artefact of the extrapolation laws, and enforcing it lets the fiction corrupt the fit where data *does* live. A single-curve constraint (density positivity, call convexity) is a property of the law itself: it has no rival extrapolation to be spuriously violated against, and abandoning it in the wings is exactly how fake tails are born.

Constraint	Curves	Treatment	Where
Calendar floor $w_{\text{far}} \geq w_{\text{near}}$	two	confined to traded range	Note 10 (variance_floor_grid_from)
LV convex-wing hinge	two ¹	confined to the extrapolation tail below the deepest quote	Note 04
MCS Durrleman $g \geq 0$	one	extended ± 2 ATM-std past quotes	section 4
De-Am call convexity	one	extended to the wings, but band-constrained and ATM-core-fixed	Note 05

The de-Am row shows the principle’s refinement rather than an exception: the constraint (convexity) rightly extends into the wings, but the *correction* it licenses is confined — to the quoted bid–ask band and away from the ATM core — because a repair, unlike a constraint, moves prices. Extending a one-curve constraint is safe; extending a repair’s *authority* unboundedly is how a 27% put wing became 104% (Note 05).

Remark 1 (Open problem: softer two-curve enforcement in the wings). Confinement as stated is binary: a two-curve constraint is either sampled (inside the traded range) or ignored (outside). But the extrapolated region is not uniformly worthless — just past the last quote the options still carry genuine premium, and there a calendar inversion between two *stated* wing contracts (section 3) is a real inconsistency even if no quote witnesses it. Whether two-curve constraints deserve a *softer* enforcement in the near-extrapolation region — the penalty weight decaying with distance past the last quote, or with remaining time value, so gross contract disagreements are still leaned on while deep-wing phantoms stay harmless — is open (flagged 2026-07). It must be settled jointly with this note’s wing contracts and Note 10’s floor-grid construction; see the matching open problem there. *Phase 1 (2026-07-09): the Quality report now measures this region — worst Durrleman g , calendar crossings and wing-slope order over the time-value envelope, advisory only (never gating readiness).*

¹The LV wing hinge compares the model’s wing against the *quotes*’ implied shape when sampled inside the quoted range — the regression of Note 04 — so it behaves as a two-curve constraint and is confined; in the pure extrapolation tail it degenerates to a one-curve convexity preference, which is where it is allowed to act.

Phase 2 (same day): opt-in tapered enforcement (`extrapEnforce`, OFF by default) adds exactly these three hinges to the SVI/MCS overlay fits, budgeted in vol units at a fraction of the data weight so it leans without outvoting quotes; the far-field calendar order is enforced through the SCALAR wing-slope condition, never pointwise. Note 10’s matching remark carries the mechanism and the measured lean/cost trade. Phase 3 (2026-07-10): the publish-time wing-only projection — exported curve samples lifted onto the discrete arbitrage-free cone in OTM-price space, traded core pinned byte-identically, fits and in-app views untouched (`volfit/models/projection.py`); the projection is a repair, so its authority is confined to the wings per section 5 even though the constraints it restores extend there. Note 10’s remark carries the mechanism.

6 What is original here

The Lee theory is classical; the contribution is the *uniform* treatment. The same slope β is computed three structurally different ways ((2), the SVI closed forms, (3)) and shown to agree on one window, so the model choice is a choice of *how* to extrapolate, never *whether* the extrapolation is admissible. The put-wing regularizer is a census-guided constraint — placed where the measured violations live, weighted by their measured asymmetry, and validated by an ablation that also proved it composes with the input-side repair. And the confinement principle of section 5 — with its two-curve / one-curve criterion — turns four hard-won production incidents into one transferable design rule.

7 Traceability

Table 2: Claims in this note and the code/tests that lock them.

Claim	Object	Code anchor	Test anchor
LQD structural Lee slopes; (2) underflow-safe		<code>models/lqd/basis.py</code>	<code>test_lqd_basis.py::</code> <code>test_svi_fit_lee_</code> <code>slopes_match_note,</code> <code>::test_lee_slopes_handle_</code> <code>underflowed_endpoint_</code> <code>scales</code>
SVI Lee cap enforced	section 3	<code>models/svi_jw/</code> <code>calibrate.py</code>	<code>test_svi_calibrate.py::</code> <code>test_respects_lee_wing_</code> <code>bound</code>
MCS hats are zero-wing	proposition 1	<code>models/sigmoid/sigmoid.</code> <code>py, kernels.py</code>	<code>test_sigmoid_model.py</code>
2-core cap	section 4	<code>api/schemas.py, models/</code> <code>sigmoid/calibrate.py</code>	<code>test_siv_wing_penalty.</code> <code>py::test_cores_clamped_</code> <code>to_two</code>
Put-wing hinge removes wing arb; byte-identical on clean slices	section 4	<code>models/sigmoid/</code> <code>calibrate.py</code>	<code>test_siv_wing_penalty.</code> <code>py::test_penalty_removes_</code> <code>wing_arb, ::test_penalty_</code> <code>byte_identical_on_arb_</code> <code>free_slice</code>
Analytic FD-free arb metric	section 4	<code>backend/backtest/</code> <code>dispatch.py, models/</code> <code>sigmoid/kernels.py</code>	<code>test_backtest_arb.py</code>

Claim	Object	Code anchor	Test anchor
R3×R6 ablation: complementary	complemen- case file	backend/backtest/ablation_arb.py	test_ablation_arb.py
De-Am wing repair (input side)	section 5	calib/convex_deam.py	test_convex_deam.py
Calendar floor confined (two-curve)	section 5	calib/calendar.py	test_overlay_calendar.py::test_wide_grid_breaks_svi_but_data_grid_does_not
LV convex wing confined to tail	section 5	api/affine_fit.py	test_affine_grid_design.py::test_convex_wing_confined_to_quoted_extrapolation

A Hyperparameter atlas

Table 3: Wing-related knobs (cross-model).

Knob	Default	Role
<i>Surfaced</i>		
sivWingPenaltyPct	100	Put-wing Durrleman hinge strength (% of the 10^3 base weight); 0 disables.
nCores	2 (max 2)	MCS core count, schema-clamped (section 4).
leeSlopeMax	2.0	SVI Lee cap $b(1 + \rho) \leq \beta_{\max}$.
sviPenaltyWeight	10^3	Weight of the SVI Lee / min-variance penalties.
leftWingSlopeMult	1.5	LV left extrapolation slope seed; free variable under a var-swap (Note 04).
lvVolCapMult	3.0	LV local-vol cap multiple (Note 04).
<i>Hidden</i>		
_WING_PAD	2.0	The g -grid extends ± 2 standardized units past the traded range.
_WING_GRID	49	Points on the g -grid.
_WING_PUT_FACTOR	2.0	Extra weight on the put side (the census asymmetry).
EPS_AR	10^{-6}	LQD $A_R < 1$ buffer; the only LQD wing constraint (Note 01).
calendar floor grid	traded range	Confined per the principle of section 5 (Note 10).

B Performance notes

Wing handling is essentially free: the LQD slopes are a by-product of the endpoint scales; SVI/MCS slopes are closed forms of the fitted parameters; the LV wing is part of the surface solve. The put-wing hinge adds 49 residual rows in the refine stage only, with a hybrid Jacobian (analytic blocks untouched, g rows by finite differences); on a clean slice the rows are zero and the optimum

is byte-identical. The analytic arb metric (R2) replaced a double `np.gradient` on 201 points with closed-form w', w'' per model — more accurate *and* cheaper. Confinement (section 5) reduces work by not over-constraining the extrapolation region.

C Reference implementation

The listing of section 2 is verified against `models/lqd/basis.py`: `lee_psi` is character-identical, and the distilled `lee_slopes` (taking the endpoint scales directly) agrees with the production function (which takes the LQD parameter vector and computes the endpoint scales itself) to 10^{-15} across the benchmark fits, including the underflow guards at extreme scales.

References

- [1] R. W. Lee. The moment formula for implied volatility at extreme strikes. *Mathematical Finance*, 14(3):469–480, 2004.
- [2] S. Benaim and P. Friz. Regular variation and smile asymptotics. *Mathematical Finance*, 19(1):1–12, 2009.
- [3] J. Gatheral and A. Jacquier. Arbitrage-free SVI volatility surfaces. *Quantitative Finance*, 14(1):59–71, 2014.