

Wings, Lee Bounds and Tail Control

How four models keep their extrapolated wings arbitrage-admissible

Vol-Fitter Technical Notes, No. 09

Vol-Fitter

July 3, 2026

Abstract

A smile model earns its keep in the *wings* — the region beyond the quoted strikes, where the fit becomes pure extrapolation and where mispriced tails leak into var-swaps, risk-reversals and butterflies. This cross-cutting note collects the wing treatment shared by the Vol-Fitter’s models. The universal constraint is Lee’s moment formula: the asymptotic slope of total implied variance, $\beta = \limsup_{k \rightarrow \pm\infty} w(k)/|k|$, cannot exceed 2, and it is tied to the number of finite moments of the underlying by an explicit map $\beta = \psi(p)$. We show how each model realizes or enforces this: LQD reads the slopes off its endpoint scales and is Lee-consistent *by construction*; SVI has linear wings and enforces the bound with a soft penalty; Multi-Core SIV inherits its base wings because the hats are wing-neutral; and the local-vol surface extrapolates with a controlled linear slope. A single calibration window, fitted by all three parametric models and extrapolated far, shows their wings agreeing in sign and order of magnitude while differing in detail — and we explain the subtle bug in which a fixed wide arbitrage-floor grid let a model’s own linear extrapolation manufacture phantom violations, and the confinement fix.

Contents

1	The universal constraint: Lee’s moment formula	2
2	Four models, four wing mechanisms	3
2.1	LQD: Lee-consistent by construction	3
2.2	SVI: linear wings with a soft cap	3
2.3	Multi-Core SIV: inherited wings	3
2.4	Local volatility: controlled linear extrapolation	4
2.5	Side by side	4
3	A wing-extrapolation pitfall: phantom calendar violations	5
4	What is original here	5
A	Hyperparameter atlas	5

1 The universal constraint: Lee's moment formula

Let $w(k, T) = \sigma_{BS}^2(k, T) T$ be total implied variance at log-moneyness $k = \log(K/F_T)$. Lee's theorem bounds how fast the wings may grow.

Theorem 1 (Lee moment formula). *Define the right tail slope $\beta_R = \limsup_{k \rightarrow +\infty} w(k, T)/k$ and the left slope $\beta_L = \limsup_{k \rightarrow -\infty} w(k, T)/|k|$. Let $p_R^* = \sup\{p : \mathbb{E}[S_T^{1+p}] < \infty\}$ and $q_L^* = \sup\{q : \mathbb{E}[S_T^{-q}] < \infty\}$ be the right/left moment critical exponents. Then*

$$\beta_R = \psi(p_R^*), \quad \beta_L = \psi(q_L^*), \quad \psi(p) = 2 - 4(\sqrt{p^2 + p} - p), \quad (1)$$

and in particular $0 \leq \beta \leq 2$. Equivalently $\frac{1}{2\beta} + \frac{\beta}{8} - \frac{1}{2} = p$.

The map and the way LQD reads its wing slopes off the endpoint scales $A_L, A_R = e^{g(0)}, e^{g(1)}$ (section 2) are a few lines (matches production to $10^{-13}/10^{-10}$):

Listing 1: Lee's ψ and the LQD structural wing slopes (distilled from `models/lqd/basis.py`).

```
1 def lee_psi(p):                                     # beta = psi(p); decreasing,
    psi(0)=2, psi(inf)=0
    return 2 - 4*(np.sqrt(p*p + p) - p)
2
3
4 def lqd_wing_slopes(A_L, A_R):                       # exponential tails -> exact
    moment exponents
    beta_L = lee_psi(1/A_L)                           # left: p* =
    1/A_L
    beta_R = 2.0 if A_R >= 1 else lee_psi(1/A_R - 1)   # right: p* =
    1/A_R - 1 (A_R < 1)
5
6
7 return beta_L, beta_R
```

Figure 1 plots ψ . It is decreasing, from $\psi(0) = 2$ (no finite moment beyond the first $\Rightarrow$ steepest admissible wing) to $\psi(\infty) = 0$ (all moments finite $\Rightarrow$ flat wing). The ceiling $\beta = 2$ is model-free: any smile with a steeper asymptotic variance slope admits arbitrage.




Figure 1: The Lee map $\beta = \psi(p)$: heavier tails (smaller moment exponent p) force steeper variance wings, capped at the model-free ceiling $\beta = 2$.

Heuristic

The wing slope and the tail heaviness are two faces of one coin. A fat right tail (few finite positive moments) *must* show up as a steep right variance wing; a model that draws a flat wing over a fat-tailed density is internally inconsistent and will misprice far-OTM calls. Lee’s bound is therefore not a nuisance constraint but a consistency condition between the smile and the law it implies — which is why the LQD model, built from the law outward, satisfies it automatically.

2 Four models, four wing mechanisms

2.1 LQD: Lee-consistent by construction

The LQD model (Note 01) parametrizes the log quantile density with logarithmic endpoints, giving *exponential* return tails with scales A_L, A_R read off the endpoint coefficients. The moment exponents are then exact, $p_R^* = 1/A_R - 1$ and $q_L^* = 1/A_L$, and the wing slopes follow from theorem 1,

$$\beta_L = \psi(1/A_L), \quad \beta_R = \psi(1/A_R - 1), \quad A_R < 1. \quad (2)$$

No penalty enforces this — it is structural. The only wing-related constraint is the hard $A_R < 1$ (a finite forward), handled by the soft barrier of Note 01. The benchmark fit yields $(\beta_L, \beta_R) = (0.0971, 0.0358)$.

2.2 SVI: linear wings with a soft cap

Raw SVI (Note 02) has exactly linear wings, $w(k) \sim b(1 \mp \rho)|k|$, so its slopes are $\beta_L = b(1 - \rho)$ and $\beta_R = b(1 + \rho)$ in closed form ($= (0.1093, 0.0364)$ here, though the figure’s fitted slice differs slightly). Lee’s bound is the explicit inequality $b(1 + |\rho|) \leq 2$, enforced as one of the two soft no-arbitrage penalties (weight `sviPenaltyWeight`, cap `leeSlopeMax`=2). On any admissible smile the penalty is zero; it only fences the optimizer out of the steep-wing region.

2.3 Multi-Core SIV: inherited wings

The Multi-Core SIV slice (Note 03) is a convex base plus R *zero-wing* hats. Because each hat and its first two derivatives vanish in both tails, the asymptotic wings depend only on the base, with z -space variance slopes $S_0 \mp 2K_0/\kappa_{P,C}$. Converting to total-variance k -space via $k = \sigma_{\text{ref}}\sqrt{t}z$, $w = tv$ gives

$$\beta_{L,R} = \frac{\sqrt{t}}{\sigma_{\text{ref}}} \left| S_0 \mp \frac{2K_0}{\kappa_{P,C}} \right|, \quad (3)$$

$= (0.1057, 0.0297)$ on the benchmark. Adding cores cannot change the wings — the entire reason for the zero-wing design. The zero-wing property pins the asymptotic *slope*; it does not by itself keep the finite-wing *curvature* admissible, since a hat can still dent the Durrleman g below zero out in the sparsely quoted put wing. That residual is what the default-on put-wing penalty of Note 03 discourages, sampling g a fixed padding past the quoted range.

2.4 Local volatility: controlled linear extrapolation

The piecewise-affine local-volatility surface (Note 04) has no closed-form implied-vol wing; beyond the deepest quoted strike $x_{\min}$ it extrapolates the *local* variance linearly with a slope tied to the first interior cell (parameter `leftWingSlopeMult`), rising toward $x \rightarrow 0$. The Dupire map then turns that into an implied-vol wing. The slope multiplier and, when a var-swap quote is present, a free left-extrapolation coefficient give the desk direct control of the put wing where there are no strikes to pin it.

2.5 Side by side

Figure 2 fits the same narrow calibration window with all three parametric models and extrapolates far past it. The wings agree in sign (the equity put skew: $\beta_L > \beta_R$) and order of magnitude, but diverge in detail — each model’s tail is a property of *its* extrapolation law, not of the (identical) quotes. Table 1 collects the analytic slopes.




Figure 2: Three models fitted to the same window (shaded) and extrapolated. Inside the window they agree to vol bps; outside, each draws its own Lee-admissible wing.

Table 1: Analytic asymptotic total-variance wing slopes $\beta = \lim w(k)/|k|$ on the benchmark window; all are far below the Lee ceiling 2.

Model	β_L (put)	β_R (call)	mechanism
LQD	0.0971	0.0358	structural via $\psi(1/A)$
SVI	0.1093	0.0364	linear $b(1 \mp \rho)$, soft Lee cap
MC-SIV	0.1057	0.0297	base wings, hats neutral

3 A wing-extrapolation pitfall: phantom calendar violations

Wing extrapolation interacts dangerously with the *calendar* no-arbitrage machinery (Note 10). Calendar-freedom requires the later expiry’s total variance to dominate the nearer one’s at every k ; the Vol-Fitter enforces this with a soft floor evaluated on a grid of log-moneyness points. The subtlety:

Caution

If the calendar floor is evaluated on a *fixed, wide* grid (say $k \in [-1, 1]$) that extends far beyond the traded strikes, then a model’s own *linear-wing extrapolation* — which is unconstrained by data out there — can dip below the nearer expiry purely as an artefact of extrapolation, creating a *phantom* calendar violation. The optimizer then flattens the far fit to “cure” a problem the data never posed. This was observed live: SVI overlays on NVDA and SPY had their far wings flattened by phantom violations manufactured on the fixed wide grid.

The fix is *confinement*: the calendar floor is built only on the *traded* log-moneyness range (`variance_floor_grid_from` the actual quotes), so the constraint bites where there is data and the wings are left to each model’s own Lee-admissible extrapolation. The LQD and sigmoid math were untouched; only the grid over which the floor is sampled changed. This is the single most important practical lesson about wings: *do not let a constraint sample the curve where only extrapolation lives*. Note 10 develops the calendar machinery in full.

Remark 1 (Confinement versus intrinsic constraints). The confinement lesson is specific to constraints that compare *two* extrapolations. A calendar floor pits one expiry’s wing against another’s, so sampling it where neither is pinned by data manufactures a phantom violation — hence confine it to the traded range. A *butterfly* constraint is different: it is an intrinsic property of a *single* slice ($g \geq 0$, equivalently a convex call curve) that must hold everywhere, wings included. So the Multi-Core SIV put-wing penalty (Note 03) deliberately samples g *past* the quoted range, and the de-Americanization convexity repair (Note 05) fixes the same wing butterfly at the *input* — neither contradicts confinement, because a single-curve constraint has no rival extrapolation to be spuriously violated against.

4 What is original here

The Lee theory is classical; the contribution is the *uniform* treatment. The same wing slope β is computed three structurally different ways ((2), the SVI linear slopes, (3)) and shown to agree, so the choice of model is a choice of *how* to extrapolate, not *whether* the extrapolation is admissible. And the confinement fix of section 3 — recognizing that an arbitrage constraint must not be sampled in the extrapolation region — is a genuinely non-obvious engineering insight that the cross-model perspective makes visible.

A Hyperparameter atlas

Table 2: Wing-related knobs (cross-model).

Knob	Default	Role
<code>leeSlopeMax</code>	2.0	SVI Lee wing-slope cap $b(1 + \rho) \leq \beta_{\max}$.
<code>sviPenaltyWeight</code>	10^3	Weight of the SVI Lee/min-variance penalties.
<code>leftWingSlopeMult</code>	1.5	LV left local-variance extrapolation slope multiple (Note 04).
<code>lvVolCapMult</code>	3.0	LV local-vol cap as a multiple of max implied vol (Note 04).
<code>EPS_AR</code> (hidden)	10^{-6}	LQD $A_R < 1$ buffer; the only LQD wing constraint (Note 01).
calendar floor grid	traded range	Confined to quoted k to avoid phantom wing violations (Note 10).

Performance. Wing handling adds no measurable cost: LQD slopes are a by-product of the endpoint scales already computed; SVI/SIV slopes are closed forms of the fitted parameters; the LV wing is part of the surface solve. The only performance-relevant choice is the confinement of section 3, which *reduces* work by not over-constraining the extrapolation region.

References

- [1] R. W. Lee. The moment formula for implied volatility at extreme strikes. *Mathematical Finance*, 14(3):469–480, 2004.
- [2] S. Benaim and P. Friz. Regular variation and smile asymptotics. *Mathematical Finance*, 19(1):1–12, 2009.
- [3] J. Gatheral and A. Jacquier. Arbitrage-free SVI volatility surfaces. *Quantitative Finance*, 14(1):59–71, 2014.