

The Volatility Surface as a Field of Probability Laws

A working draft

Thibaut Delahaye

August 2026

Contents

1	Introduction	1
I	Models	3
2	The log-quantile-density model	5
2.1	A smile is a probability law	5
2.2	From ranks to LQD coordinates	8
2.3	Validity and scope	12
2.4	Tails, moments, and wings	14
2.5	The ledger calculus	17
2.6	Calibration	20
2.7	Numerical realization	23
2.8	Calendar order	24
2.9	Examples and the model contract	26
2.A	Proofs deferred from the main text	30
2.B	Numerical implementation	33
2.C	Data and reproducibility	37
3	Families in the volatility chart	39
3.1	Working in the volatility chart	39
3.2	One hyperbola	40
3.3	Wings are tails: Lee’s bound inside the density factor	42
3.4	The belly escapes	44
3.5	Trader coordinates: the jump-wing handles	45
3.6	Calibration in the chart	48
3.7	Superposition: capacity without touching the wings	50
3.8	Three families on the same quotes	53
3.A	Comparison protocol and numerical notes	56
4	Local volatility	59
4.1	One diffusion for the whole surface	59
4.2	The equation the surface must obey	60
4.3	The wrong direction	62
4.4	A sheet with finitely many parameters	64
4.5	Pricing on the lattice	66
4.6	What the quotes determine	68
4.7	The objective and its derivatives	70
4.8	Examples and the contract	71
4.A	Numerical facts and the chapter protocol	75

Chapter 1

Introduction

[This chapter is written last. It will state the problem — finitely many noisy option prices, from which a surface of probability laws must be published — and the thesis that organizes every chapter: the separation of what follows from probability, what the market determines, and what the modeler must choose and disclose.]

Part I

Models

Chapter 2

The Log-Quantile-Density Model of the Volatility Smile

A volatility smile is a probability law expressed through the Black formula. This chapter constructs that law directly. The log-forward return is written as a monotone transport of a logistic variable, and the logarithm of the transport speed is expanded in Legendre polynomials of the percentile rank. Monotonicity, positive density, and unit mass are then automatic; a scalar shift imposes the martingale condition, and the sole remaining requirement is that the right endpoint speed be smaller than one. The endpoint speeds give the exact moment strip and the asymptotic Lee slopes. A single tail integral, the upper-share ledger, yields option prices, the density, the analytic calibration Jacobian, the variance-swap strike, and a complete test for calendar order. The construction separates what follows from probability, what is learned from quoted options, and what must be supplied by modeling convention. Theory, computation, and market examples are developed from the same representation.

2.1 A smile is a probability law

A listed option market gives finitely many prices. A smile model must turn them into a complete curve: a price and an implied volatility at every strike, including strikes that were not quoted. The completed curve will be differentiated for digitals and densities, integrated for variance contracts, and compared with curves at other expiries. It is therefore not enough to draw a smooth line through the quotes.

The obstruction is elementary. The second strike derivative of a call price is a probability density. A curve that fits every quote can nevertheless make this derivative negative between two quotes. It then assigns a negative price to a butterfly and a negative mass to an interval of terminal prices. No numerical routine has malfunctioned; the curve was allowed to leave the space of probability laws.

The useful question is thus not “which curve fits well?” but “in which coordinates is every fitted curve already a law?” The answer developed here uses quantiles. In rank coordinates, unit mass is automatic and positive density reduces to monotonicity. Taking the logarithm of the quantile speed removes even that constraint. The resulting log-quantile-density (LQD) family provides unconstrained coordinates for a broad class of mean-one laws with exponential log-return tails.

2.1.1 Normalized market

Fix one expiry T . Let F be its forward and D its discount factor. Under the T -forward measure define

$$X = \log \frac{S_T}{F}, \quad Y = e^X = \frac{S_T}{F}, \quad \mathbb{E}[Y] = 1. \quad (2.1)$$

The last equality is the martingale condition. Write $k = \log(K/F)$ for log-moneyness. Normalized call and put prices are

$$c(k) = \mathbb{E}[(Y - e^k)^+], \quad P(k) = \mathbb{E}[(e^k - Y)^+], \quad c(k) - P(k) = 1 - e^k. \quad (2.2)$$

A dollar price is DF times its normalized price.

Let τ denote variance time: calendar time, possibly adjusted for the desk's event-time convention. Total implied variance is $w(k) = \sigma(k)^2 \tau$. The normalized Black call is

$$B(k, w) = \Phi(d_+) - e^k \Phi(d_-), \quad d_{\pm} = -\frac{k}{\sqrt{w}} \pm \frac{\sqrt{w}}{2}, \quad (2.3)$$

and $w(k)$ is the root of $B(k, w(k)) = c(k)$ [4]. Rates and carry have now disappeared from the single-slice problem.

2.1.2 The static constraint set

Use normalized strike $y = e^k$. If Y has density f_Y , differentiation of the payoff gives

$$\frac{\partial c}{\partial y} = -\mathbb{P}(Y > y), \quad \frac{\partial^2 c}{\partial y^2} = f_Y(y) \geq 0. \quad (2.4)$$

This is the Breeden–Litzenberger identity [6]. It identifies the geometry of the call curve with the law of Y : decreasing slope is a digital probability and convexity is positive density.

The identity can be read using traded portfolios. For a strike spacing δ , the symmetric butterfly

$$c(y - \delta) - 2c(y) + c(y + \delta)$$

is the price of a long call below, a long call above, and two short calls at the center. Dividing by δ^2 and shrinking the spacing recovers $f_Y(y)$. A negative second difference is therefore not merely undesirable curvature; it is a negative price for a nonnegative terminal payoff. The probability statement and the static-arbitrage statement are the same.

Definition 2.1 (Arbitrage-free slice). An arbitrage-free slice is a call curve of the form (2.2) for a positive random variable Y with $\mathbb{E}[Y] = 1$. When Y has a density, this is equivalent to a decreasing convex function of strike y satisfying

$$(1 - y)^+ \leq c(y) \leq 1, \quad -1 \leq c'(y) \leq 0, \quad c(y) \rightarrow 0 \quad (y \rightarrow \infty),$$

with $c'(y) \rightarrow -1$ as $y \downarrow 0$ when there is no mass at zero.

Each boundary condition has a probabilistic source. The upper bound follows from $(Y - y)^+ \leq Y$ and $\mathbb{E}[Y] = 1$. The lower bound is Jensen's inequality, $(\mathbb{E}[Y] - y)^+ \leq \mathbb{E}[(Y - y)^+]$. The slope is a survival probability and must lie between -1 and zero. Finally, calls vanish at infinite strike because the integrable random variable Y carries no mean mass arbitrarily far in its tail. A valid call curve is thus exactly a compact encoding of a mean-one law.

Convexity belongs to y , not to k . Indeed,

$$\frac{\partial^2 c}{\partial k^2} = e^k \{e^k f_Y(e^k) - \mathbb{P}(Y > e^k)\},$$

which is negative far in the left wing for any ordinary density. A convexity test performed in log-strike will therefore reject correct call curves.

The same constraint becomes less transparent in volatility coordinates. Differentiating $c(k) = B(k, w(k))$ twice yields

$$f_X(k) = g_D(k) \frac{\varphi(d_-(k, w(k)))}{\sqrt{w(k)}}, \quad g_D(k) = \left(1 - \frac{k w'}{2w}\right)^2 - \frac{(w')^2}{4} \left(\frac{1}{w} + \frac{1}{4}\right) + \frac{w''}{2}. \quad (2.5)$$

Thus a volatility curve is butterfly-free precisely when $g_D(k) \geq 0$ for all k [13]. The condition is global, nonlinear, and involves two derivatives. Direct volatility models must enforce it; the representation below makes it an identity.

This explains why good residuals alone are weak evidence. Quotes constrain only finitely many values of w . The function g_D also depends on the slope and curvature between them, and may become negative in a narrow unquoted interval. Increasing the strike grid reduces but does not remove the logical problem: the optimizer still moves in a space containing invalid curves. The quantile construction reverses the order of work. It first chooses a law and only then asks which smile the Black formula assigns to it.

Two checks fix the normalization in (2.5). If the smile is flat, $w' = w'' = 0$, so $g_D \equiv 1$ and the remaining factor is exactly the Black density of X . At the money, the identities derived later in section 2.5 give $g_D(0) = \sqrt{w_0} f_X(0) / \varphi(\sqrt{w_0}/2)$. Thus $g_D \geq 0$ is not a penalty formula: it is positive density written in the volatility chart.

2.1.3 Notation ledger

The construction uses three horizontal axes. Strike y is the axis on which calls are convex; log-moneyness $k = \log y$ is the market axis; rank u and log-odds z are the model axes. For a rank $u \in (0, 1)$, the *odds* are $u/(1 - u)$ and the *log-odds* are

$$z = \log \frac{u}{1 - u}.$$

The function taking rank to log-odds is called the *logit*: $\text{logit}(u) = \log\{u/(1 - u)\}$. It maps $(0, 1)$ bijectively onto $\mathbb{R}$; its inverse is the logistic function $\Lambda(z) = 1/(1 + e^{-z})$. Thus u and z record the same percentile on two different scales. The recurring notation is collected below.

$X, Y = e^X$	log and gross forward return	$u, z = \text{logit } u$	rank and log-odds
$k, y = e^k$	log and normalized strike	Λ, ρ	logistic CDF and density
c, P	normalized call and put	$Q, q = Q'$	quantile and quantile density
$w = \sigma^2 \tau$	total implied variance	g	smooth log-speed in rank
x, m	transport and martingale shift	G	upper-share ledger
L, R, a_n	polynomial coordinates	λ_-, λ_+	endpoint speeds
u_k, z_k	strike rank and log-odds	u_*, f_*, Δ	ATM rank, density, digital gap
β_L, β_R	Lee wing slopes	$N, Z_{\max}$	basis order and grid half-width

Table 2.1: Notation used throughout the chapter. A prime denotes an ordinary derivative of a one-variable function; ∂ denotes a partial derivative. The time τ is always variance time.

2.1.4 What the chapter separates

Three questions recur, and they should not be confused. *Validity* asks whether a parameter vector defines a probability law. *Information* asks which features of that law are determined by the quoted strip. *Convention* selects the remainder: basis order, regularization, tail coordinates, and the scope of calendar control. The LQD coordinates settle the first question structurally. The later sections measure the second and state the third explicitly.

The development is progressive. Section 2.2 constructs the coordinates. Sections 2.3 and 2.4 establish their probabilistic content. Section 2.5 derives prices and local smile quantities from one integral. Section 2.6 turns the representation into a fitter, section 2.7 into a computation, and section 2.8 into a surface. Section 2.9 then shows what the construction says on market and synthetic data. Proofs and implementation details are collected in the appendices.

2.2 From ranks to LQD coordinates

2.2.1 The speed of a quantile

Let $Q(u)$ be the quantile function of X , with rank $u \in (0, 1)$. If U is uniform, then $Q(U)$ has the desired law. Hence probability one is built into the coordinate u , and the only remaining requirement is that Q increase.

Rank is a universal clock for distributions. The value $Q(0.01)$ is the first-percentile return whether the law is symmetric, skewed, or multimodal. Changing the law changes the returns assigned to the clock, not the clock itself. This is the main simplification: normalization of probability is handled once by the uniform variable rather than reimposed on every candidate density.

Write $q(u) = Q'(u)$ for the quantile density. Differentiating $F_X(Q(u)) = u$ gives

$$f_X(Q(u)) = \frac{1}{q(u)}. \quad (2.6)$$

The function q is the speed with which return space is traversed as rank increases. Slow motion packs probability into a narrow return interval; fast motion thins it. Since Q is increasing exactly when $q > 0$, an arbitrary real function may be used for $\log q$.

An unbounded distribution forces q to diverge at both endpoints. Here a *tail* means the probability lying beyond a remote threshold. We say that X has exponential tails, with positive scales λ_- and λ_+ , when

$$F_X(x) \sim C_- e^{x/\lambda_-} \quad (x \rightarrow -\infty), \quad 1 - F_X(x) \sim C_+ e^{-x/\lambda_+} \quad (x \rightarrow +\infty), \quad (2.7)$$

for constants $C_-, C_+ > 0$. Their logarithms therefore decay linearly in $|x|$. By contrast, a Gaussian tail has log-probability proportional to $-x^2$: it decays quadratically on the log-probability scale and hence faster than every exponential tail. For the exponential class, the leading endpoint singularities of q are $1/u$ and $1/(1-u)$. Remove them explicitly:

$$\log q(u) = -\log u - \log(1-u) + g(u). \quad (2.8)$$

The remainder g is bounded for the exponential-tail class considered here.

The subtraction is not arbitrary. The left relation in (2.7) implies $Q(u) \sim \lambda_- \log u$ as $u \downarrow 0$ and $q(u) \sim \lambda_-/u$. The right endpoint is analogous. More precisely, if $X \sim N(0, s^2)$, then

$$F_X(x) \sim \frac{s}{|x|\sqrt{2\pi}} e^{-x^2/(2s^2)} \quad (x \rightarrow -\infty),$$

with the symmetric formula on the right. Its quantile density is $q(u) = s/\varphi(\Phi^{-1}(u))$ and, by the normal tail equivalent,

$$q(u) \sim \frac{s}{u\sqrt{2\log(1/u)}} \quad (u \downarrow 0).$$

Its leading $1/u$ singularity is removed, but the remainder still tends to $-\infty$. Exponential tails therefore lie inside the LQD chart, whereas Gaussian tails lie on a different boundary.

Now set

$$z = \text{logit } u = \log \frac{u}{1-u}, \quad u = \Lambda(z) = \frac{1}{1+e^{-z}}, \quad \rho(z) = \Lambda(z)\Lambda(-z). \quad (2.9)$$

The function Λ is the logistic CDF and $\rho = \Lambda'$ its density. Because $du/dz = u(1-u)$, the singularities in (2.8) cancel:

$$\frac{dQ}{dz} = q(u)u(1-u) = e^{g(u)}. \quad (2.10)$$

Thus the entire distribution is generated by integrating a positive, ordinary function on the log-odds line.

The coordinate also places every percentile at a finite location. The median is $z = 0$; the tenth and ninetieth percentiles are $z = \mp \log 9$; moving one unit in z multiplies the odds $u/(1-u)$ by e . Meanwhile $\rho(z) \sim e^{-|z|}$, so rank integrals become ordinary integrals on $\mathbb{R}$ with exponentially decaying weight. This is why the same coordinate is convenient both analytically and numerically.




Figure 2.1: The transport construction for constant speed. Left: the logistic score and five percentile marks. Center: an affine map from log-odds to log-return, with two strike roots. Right: the implied-volatility smile priced from that law. The first two panels define the distribution; the last merely expresses it in market coordinates.

2.2.2 The LQD slice

Choose g to be a low-degree polynomial in rank. The shifted Legendre basis is convenient because it is orthogonal under the uniform rank measure. The constant and linear terms are written as an endpoint chord, leaving higher modes to describe the body.

Definition 2.2 (The LQD slice). Fix $N \geq 2$ and $\theta = (L, R, a_2, \dots, a_N) \in \mathbb{R}^{N+1}$. Let Z have CDF Λ and density ρ . The LQD slice is the law of

$$\boxed{X = x(Z), \quad x(z) = m + \int_0^z e^{g(\Lambda(t))} dt, \quad g(u) = (1-u)L + uR + \sum_{n=2}^N a_n P_n(1-2u),} \quad (2.11)$$

where P_n is the Legendre polynomial of degree n and m is chosen so that $\mathbb{E}[e^X] = 1$.

For every finite coefficient vector, $x'(z) = e^{g(\Lambda(z))} > 0$. Consequently

$$F_X(x(z)) = \Lambda(z), \quad f_X(x(z)) = \frac{\rho(z)}{x'(z)} = \Lambda(z)\Lambda(-z)e^{-g(\Lambda(z))} > 0. \quad (2.12)$$

The quantile is $Q(u) = x(\text{logit } u)$, and (2.10) recovers

$$q(u) = \frac{e^{g(u)}}{u(1-u)}. \quad (2.13)$$

Equations (2.11) and (2.13) are the transport and log-quantile-density descriptions of the same law.

Since a polynomial g is bounded on $[0, 1]$, its exponential is bounded above and away from zero. Therefore $x(z)$ tends to $-\infty$ and $+\infty$ with z and is a diffeomorphism of the line. The construction cannot fold, stop, or reverse. Positivity of the density in (2.12) is thus a geometric fact about the coordinate map, not a pointwise inequality checked after the fit.

The two endpoint values are

$$g(0) = L + \sum_{n \geq 2} a_n, \quad g(1) = R + \sum_{n \geq 2} (-1)^n a_n, \quad (2.14)$$

and their exponentials

$$\lambda_- = e^{g(0)}, \quad \lambda_+ = e^{g(1)} \quad (2.15)$$

are the endpoint speeds. If $\bar{x}(z) = \int_0^z e^{g(\Lambda(t))} dt$, then Lipschitz continuity of g and exponential convergence of Λ to its endpoints imply

$$\bar{x}(z) = b_- + \lambda_- z + O(e^z) \quad (z \rightarrow -\infty), \quad \bar{x}(z) = b_+ + \lambda_+ z + O(e^{-z}) \quad (z \rightarrow +\infty). \quad (2.16)$$

The endpoint speeds will therefore govern both tails.




Figure 2.2: Three Legendre perturbations of the same reference law. Left: their action on log-speed g . Center: the response of implied volatility. Right: the response of the density. An even mode moves both endpoints in the same direction; an odd mode moves them in opposite directions. In all cases the density remains positive because the transport speed remains an exponential.

The modes have a direct interpretation. Raising g stretches returns near the affected ranks and therefore lowers density there by (2.6). Even modes act symmetrically; odd modes transfer shape from one side to the other; higher modes work on finer rank scales. The raw coefficients are not pure body controls, since (2.14) couples each of them to one or both tails. The endpoint chart introduced in section 2.6 removes that coupling without changing the family.

Legendre polynomials are a coordinate choice, not part of the probability argument. Rank integrals use uniform weight on $[0, 1]$, under which Legendre modes are orthogonal. Their columns remain distinguishable in least-squares normal equations as order rises. Monomials span the same space but inherit Hilbert-matrix conditioning. Chebyshev modes would also be reasonable; the endpoint chord would remain useful because it names the two tail controls.

2.2.3 Martingale normalization and the finite-forward wall

Only a translation remains. Let

$$I = \int_{-\infty}^{\infty} e^{\bar{x}(z)} \rho(z) dz.$$

Proposition 2.3 (Martingale normalization). *If $I < \infty$, then $m = -\log I$ is the unique shift for which $\mathbb{E}[e^X] = 1$. Moreover, $I < \infty$ if and only if $\lambda_+ < 1$.*




Figure 2.3: A fitted SPY December 2026 slice in the model's coordinates. Left: the universal boundary skeleton and the fitted $\log q$; their difference is g . Right: the density produced by the same transport. The quoted strikes occupy the central ranks; the continuation to the endpoints is extrapolation.

Proof. Since $X = m + \bar{x}(Z)$, $\mathbb{E}[e^X] = e^m I$, giving the unique shift. At the right endpoint, (2.16) and $\rho(z) \sim e^{-z}$ make the integrand asymptotic to a positive constant times $e^{(\lambda_+ - 1)z}$; it is integrable exactly when $\lambda_+ < 1$. At the left endpoint it behaves as $e^{(1 + \lambda_-)z}$, which is integrable for every $\lambda_- > 0$. $\square$

The set of admissible raw parameters is therefore

$$\Theta_N = \{\theta \in \mathbb{R}^{N+1} : \lambda_+ < 1\}. \quad (2.17)$$

It is open and has a single boundary. A logistic coordinate for λ_+ will make that boundary unreachable.

2.2.4 The exactly solvable slice

Set $L = R = \log s$, $a_n = 0$, with $0 < s < 1$. Then $x(z) = m + sz$ and X is logistic with scale s . Its moment-generating function is the beta integral

$$\mathbb{E}[e^{tZ}] = \int_0^1 \left(\frac{u}{1-u} \right)^t du = B(1+t, 1-t) = \Gamma(1+t)\Gamma(1-t) = \frac{\pi t}{\sin(\pi t)}, \quad |t| < 1. \quad (2.18)$$

Hence

$$m = -\log \frac{\pi s}{\sin(\pi s)}, \quad \text{Var}(X) = \frac{\pi^2 s^2}{3}.$$

Matching a target ATM total variance w_0 gives the cold start

$$s = \frac{\sqrt{3w_0}}{\pi}. \quad (2.19)$$

This seed matches variance, not ATM price. As $s \downarrow 0$, $m = -\pi^2 s^2/6 + O(s^4)$. Moreover,

$$\mathbb{E}[Z^+] = \int_0^\infty z \rho(z) dz = \sum_{n \geq 1} \frac{(-1)^{n-1}}{n} = \log 2,$$

where the logistic density is expanded as a geometric series on the positive half-line. Thus the LQD ATM call is $s \log 2 + O(s^2)$. Matching it to the small-variance Black price gives

$$\frac{w_{\text{implied}}}{w_0} \longrightarrow \frac{6(\log 2)^2}{\pi} \approx 0.918.$$

The seed is therefore about eight percent low in ATM variance, a known bias removed by calibration.




Figure 2.4: The solvable slice as an end-to-end numerical test. Left: errors in the computed transport and martingale shift against the closed forms from (2.18); the normalization becomes harder as s approaches the wall. Right: the variance-matched cold start and its derived small- s bias.

2.3 Validity and scope

The construction has two distinct claims. First, every admissible parameter vector defines an arbitrage-free slice. Second, increasing polynomial order fills out a specified class of laws. The first is exact at every order; the second determines the scope of the coordinates.

2.3.1 One-expiry validity

Proposition 2.4 (One-expiry no-arbitrage). *Let $\theta \in \Theta_N$ and price calls from the LQD slice (2.11). Then:*

- (i) *X has a strictly positive continuous density on $\mathbb{R}$, and $Y = e^X$ has a strictly positive density on $(0, \infty)$ with $\mathbb{E}[Y] = 1$;*
- (ii) *as a function of strike y , the call is strictly decreasing and strictly convex, satisfies $(1 - y)^+ \leq c(y) \leq 1$, has $c'(y) = -\mathbb{P}(Y > y) \in (-1, 0)$, and tends to zero as $y \rightarrow \infty$;*
- (iii) *every nondegenerate finite-width butterfly has positive price, and put-call parity holds identically;*
- (iv) *as a function of log-strike k , the call is concave sufficiently far in the money.*

Proof. Equation (2.12) gives a positive continuous density; the shift of proposition 2.3 gives unit mean. Rank integration gives

$$c(y) = \int_0^1 (e^{Q(u)} - y)^+ du.$$

Each integrand is decreasing and convex in y , and differentiation yields $c'(y) = -\mathbb{P}(Y > y)$. Positivity of the density makes both monotonicity and convexity strict. Jensen gives $(1 - y)^+ \leq c(y)$, while $c(y) \leq \mathbb{E}[Y] = 1$ and dominated convergence gives the right boundary. Butterfly positivity and parity follow immediately. Finally,

$$\frac{\partial^2 c}{\partial k^2} = e^k \{f_X(k) - \mathbb{P}(Y > e^k)\},$$

which is negative as $k \rightarrow -\infty$ because $f_X(k) \rightarrow 0$ and the probability tends to one. $\square$

The proposition is stronger than a dense-grid test. It covers every strike, including strikes between grid nodes and beyond the displayed smile, because the call is an expectation under

a positive law. Calibration may move the coefficient vector arbitrarily inside Θ_N without approaching a hidden butterfly boundary. The numerical implementation must still preserve this continuous geometry; that separate question is treated by the certificates of section 2.7.2.




Figure 2.5: Validity on the fitted SPY December 2026 slice. Left: the call curve is convex in normalized strike and lies below its secant chord. Right: finite butterflies, divided by squared half-width, converge to the density. The geometry is produced by the law, not imposed during calibration.

2.3.2 Why the wall cannot be removed

Proposition 2.5 (Finite-forward wall). *Within the LQD family the martingale normalizer exists if and only if $\lambda_+ < 1$. More generally, if a law satisfies $\mathbb{P}(X > x) \geq \kappa e^{-x}$ for some $\kappa > 0$ and all sufficiently large x , then $\mathbb{E}[e^X] = \infty$.*

Proof. The first statement is proposition 2.3. For the second, at any large x_0 ,

$$\mathbb{E}[e^X] \geq \int_{e^{x_0}}^{\infty} \mathbb{P}(X > \log t) dt \geq \kappa \int_{e^{x_0}}^{\infty} \frac{dt}{t} = \infty.$$

□

The boundary is therefore probabilistic, not parametric. A right log-return tail with scale at least one cannot support a finite forward. There is no left counterpart because e^X is bounded on the negative half-line.

2.3.3 Coverage of the coordinates

The natural target class can be described without reference to polynomials. For a differentiable quantile Q^* define

$$g^*(u) = \log Q^{*'}(u) + \log u + \log(1 - u).$$

Continuity of g^* on $[0, 1]$ is exactly the condition that the law have a positive continuous density and exponential tails with finite nonzero endpoint scales.

Proposition 2.6 (Universality, with exclusions). *Let X^* be a mean-one law for which g^* extends continuously to $[0, 1]$, with $e^{g^*(1)} < 1$. Then:*

- (i) *There exist admissible LQD slices with $g_N \rightarrow g^*$ uniformly. Their quantiles converge uniformly on compact subsets of $(0, 1)$, their call curves converge uniformly over all strikes, and their gross-return laws converge in Wasserstein- p distance for every $1 \leq p < e^{-g^*(1)}$.*

- (ii) *No finite-order LQD slice has an atom, a zero-density interval, bounded support, mass at zero, or Gaussian-or-faster tails. These exclusions persist under uniformly convergent log-speeds. More precisely, let admissible g_N converge uniformly to g_∞ . If $g_\infty(1) < 0$, the limiting law remains in the class of part (i). If $g_\infty(1) = 0$, weak convergence can fail because mass escapes to $-\infty$; if a weak limit nevertheless exists, it has full support and a strictly positive continuous density, but sits at the finite-forward boundary. Its moments of order $r > 1$ diverge, and its mean may be smaller than one.*

Here the one-dimensional Wasserstein metric has the concrete quantile form

$$W_p(\mu, \nu)^p = \int_0^1 |Q_\mu(u) - Q_\nu(u)|^p du.$$

The proof is in appendix 2.A. Its approximation step is Weierstrass: polynomials approximate the continuous g^* uniformly. Since $g^*(1) < 0$, sufficiently accurate approximants remain inside the open admissible set. The tail margin supplies a common integrable envelope for normalizers, calls, and moments.

Uniform convergence of g is the correct strength for this result. It controls the transport speed at every rank, including both endpoints. Local approximation on the quoted ranks alone would say nothing about normalization or tail moments. Once a strict margin $g^*(1) < 0$ is fixed, the same exponential envelope controls all approximating normalizers and supplies uniform integrability of gross returns. This is what upgrades pointwise shape approximation to uniform convergence of call prices.

The exclusions follow from the same geometry. A positive continuous transport speed makes the quantile strictly increasing and unbounded. An atom would require a quantile jump; a gap, a flat interval; bounded support, a finite endpoint; and a Gaussian tail, $g(u) \rightarrow -\infty$ at an endpoint. None is available to a finite polynomial, nor to a uniform limit of bounded log-speeds.

Multimodality is not excluded. By (2.6), density modes are created where the transport slows and are separated where it accelerates. A positive speed can vary many times without ever becoming zero. Thus the family may represent a smooth double-hump density, but not a literal gap of zero density between the humps. Basis order determines how sharply the speed can vary, as the synthetic example in section 2.9 demonstrates.

Thus the LQD family is not a coordinate system for every probability law. It is a coordinate system for a flexible exponential-tail class. That class is large enough to approximate smooth unimodal and multimodal smile laws, while its boundary and exclusions remain explicit.

2.4 Tails, moments, and wings

The endpoint speeds introduced in (2.15) have an exact probabilistic meaning. This section follows the chain

$$\text{endpoint speed} \longrightarrow \text{tail scale} \longrightarrow \text{finite moments} \longrightarrow \text{asymptotic smile slope}.$$

2.4.1 Endpoint speeds are tail scales

Proposition 2.7 (Exponential tails). *For every admissible LQD slice there are constants $K_\pm > 0$ such that*

$$\mathbb{P}(X > x) \sim K_+ e^{-x/\lambda_+} \quad (x \rightarrow +\infty), \quad \mathbb{P}(X < x) \sim K_- e^{x/\lambda_-} \quad (x \rightarrow -\infty). \quad (2.20)$$

Equivalently, $\mathbb{P}(Y > y) \sim K_+ y^{-1/\lambda_+}$ and $\mathbb{P}(Y < y) \sim K_- y^{1/\lambda_-}$.

Proof. On the right, (2.16) gives $x(z) = m + b_+ + \lambda_+ z + O(e^{-z})$. Its inverse satisfies

$$z(x) = \frac{x - m - b_+}{\lambda_+} + O(e^{-z(x)}).$$

Since $1 - \Lambda(z) \sim e^{-z}$,

$$\mathbb{P}(X > x) = 1 - \Lambda(z(x)) \sim e^{(m+b_+)/\lambda_+} e^{-x/\lambda_+}.$$

The left tail is identical after reversing signs. $\square$

The gross return therefore has power tails, while the log-return has exponential tails. The two endpoint coefficients control extrapolation in a directly readable unit: log-return distance per e -fold reduction in tail probability.

2.4.2 The moment strip

Every expectation may be written either in rank or log-odds:

$$\mathbb{E}[h(X)] = \int_0^1 h(Q(u)) du = \int_{-\infty}^{\infty} h(x(z)) \rho(z) dz. \quad (2.21)$$

Proposition 2.8 (Moment strip). *For an admissible LQD slice,*

$$\mathbb{E}[e^{rX}] < \infty \iff -\frac{1}{\lambda_-} < r < \frac{1}{\lambda_+}. \quad (2.22)$$

Consequently the last finite positive moment of Y beyond the first and the last finite inverse moment are

$$r_+^* = \frac{1}{\lambda_+} - 1, \quad r_-^* = \frac{1}{\lambda_-}.$$

Proof. In (2.21), the right-tail integrand for $h(x) = e^{rx}$ is asymptotic to $e^{(r\lambda_+-1)z}$; the left-tail integrand is asymptotic to $e^{(r\lambda_-+1)z}$. Integrability gives exactly (2.22), with divergence at both boundary values. $\square$

The martingale condition is the point $r = 1$. It lies inside the strip exactly when $\lambda_+ < 1$, so normalization, the finite-forward wall, and the first moment are three forms of the same fact.

The asymmetry of the strip is financially important. A large λ_- may eliminate low-order inverse moments of Y without threatening the forward; negative log-returns contribute at most one to Y . A comparable right-tail scale would destroy the first positive moment and hence the forward itself. For the long NVDA example, $\lambda_- > 1$, so even $\mathbb{E}[Y^{-1}]$ diverges, while $\mathbb{E}[Y] = 1$ remains exact.

2.4.3 Closed Lee slopes

Lee's moment formula [19] relates the critical moments to the asymptotic slopes of total implied variance:

$$\beta_R = \limsup_{k \rightarrow +\infty} \frac{w(k)}{k} = \Psi(r_+^*), \quad \beta_L = \limsup_{k \rightarrow -\infty} \frac{w(k)}{|k|} = \Psi(r_-^*), \quad (2.23)$$

where

$$\Psi(r) = 2 - 4(\sqrt{r^2 + r} - r), \quad r \geq 0.$$

Proposition 2.9 (Speed-to-slope maps). *The LQD wing slopes are*

$$\beta_R = \frac{2\lambda_+}{(1 + \sqrt{1 - \lambda_+})^2}, \quad \beta_L = \frac{2\lambda_-}{(1 + \sqrt{1 + \lambda_-})^2}. \quad (2.24)$$

Both maps increase with their endpoint speed. For small speeds, $\beta_{\pm} \sim \lambda_{\pm}/2$; as $\lambda_+ \uparrow 1$, $\beta_R \uparrow 2$. For LQD slices the limsups in (2.23) are ordinary limits.

Proof. Substitute $r_+^* = 1/\lambda_+ - 1$ and $r_-^* = 1/\lambda_-$ into Ψ and rationalize. For example, with $s = \sqrt{1 - \lambda_+}$,

$$\Psi(1/\lambda_+ - 1) = 2 - \frac{4s}{1+s} = \frac{2\lambda_+}{(1+s)^2}.$$

The left formula is obtained with $t = \sqrt{1 + \lambda_-}$. Monotonicity and limits follow from the final expressions. By proposition 2.7, the return tails are regularly varying; the Benaim–Friz refinement of Lee’s formula therefore gives convergence rather than only a limsup [2]. $\square$

For completeness, the left algebra is

$$\Psi(1/\lambda_-) = 2 - \frac{4(\sqrt{1 + \lambda_-} - 1)}{\lambda_-} = \frac{2\lambda_-}{(1 + \sqrt{1 + \lambda_-})^2}.$$

The formulas expose two useful scales. When λ is small, the wing slope is about half the endpoint speed. On the right, the map reaches Lee’s model-free ceiling 2 exactly when the first moment reaches its boundary. The tail coordinate and the no-arbitrage wing bound meet at the same point.




Figure 2.6: The tail chain. Left: endpoint speed determines the last finite moment; the right map reaches the finite-forward wall at $\lambda_+ = 1$. Center: Lee’s moment-to-slope function. Right: the closed compositions in (2.24). Markers show the fitted SPY December 2026 and NVDA December 2027 slices.

The marked market slices occupy only a small part of the admissible range. For SPY December 2026 the endpoint speeds are $\lambda_- = 0.432$ and $\lambda_+ = 0.066$, giving slopes $\beta_L = 0.179$ and $\beta_R = 0.034$. The long NVDA slice has $\lambda_- = 1.421$ and $\lambda_+ = 0.225$, with slopes 0.435 and 0.128. The heavier single-name tails appear as larger endpoint speeds before the smile is ever plotted.

2.4.4 The limit and the traded wing

The numbers $\beta_{\pm}$ describe limits, not quoted strikes. If $w(k) = \beta|k| + b + o(1)$, then $w(k)/|k| = \beta + b/|k| + o(1/|k|)$ and may approach its limit from either side. The sign of the intercept b is not fixed.

No finite option strip identifies an asymptotic decay rate by itself. In a fit, $\lambda_{\pm}$ are selected jointly by the outer quotes, basis order, regularization, and endpoint coordinates. The model makes this choice low-dimensional and interpretable, but it does not turn extrapolation into observation.

The finite-strike discrepancy is large enough to matter. On the SPY slice in fig. 2.7, the effective right slope is 1.69 times its limit at ten delta and 1.05 times at one delta. On the left, the corresponding ratios are 0.70 and 0.70, both below one. A limiting slope can therefore overstate one traded wing and understate the other on the same fitted law.




Figure 2.7: Effective variance slopes for the fitted SPY December 2026 slice. The left wing lies below its Lee limit at the marked deltas, while the right wing lies above. The asymptotic slopes are boundary conditions for extrapolation; they are not substitutes for prices at finite strikes.

2.5 The ledger calculus

Once the transport x is known, a single tail integral contains nearly every quantity needed by the fitter.

2.5.1 The upper-share ledger

Definition 2.10 (Upper-share ledger). For an admissible slice define

$$G(z) = \int_z^\infty e^{x(t)} \rho(t) dt = \mathbb{E}[Y \mathbf{1}\{Z > z\}]. \quad (2.25)$$

The name is literal: $G(z)$ is the fraction of the unit forward carried by outcomes above log-odds level z . It satisfies

$$G(-\infty) = 1, \quad G(+\infty) = 0, \quad G'(z) = -e^{x(z)} \rho(z) < 0. \quad (2.26)$$

For strike k , let z_k be the unique root and u_k its rank:

$$x(z_k) = k, \quad u_k = \Lambda(z_k) = F_X(k). \quad (2.27)$$

Proposition 2.11 (Ledger pricing). For every admissible slice,

$$c(k) = G(z_k) - e^k(1 - u_k), \quad (2.28)$$

$$P(k) = e^k u_k - (1 - G(z_k)), \quad (2.29)$$

$$c'(k) = -e^k(1 - u_k), \quad (2.30)$$

$$c''(k) = -e^k(1 - u_k) + e^k \frac{\rho(z_k)}{x'(z_k)}, \quad (2.31)$$

$$f_X(k) = e^{-k} \{c''(k) - c'(k)\} = \frac{\rho(z_k)}{x'(z_k)}. \quad (2.32)$$

Proof. Because x increases, exercise is the event $Z > z_k$. Splitting asset and cash legs gives

$$c(k) = \mathbb{E}[Y \mathbf{1}\{Z > z_k\}] - e^k \mathbb{P}(Z > z_k) = G(z_k) - e^k(1 - u_k),$$

and the put follows on the complementary event. Differentiating the call, the moving-root term is

$$\{G'(z_k) + e^k \rho(z_k)\} \frac{dz_k}{dk} = 0$$

because $G'(z_k) = -e^{x(z_k)}\rho(z_k) = -e^k\rho(z_k)$. This yields (2.30). Since $dz_k/dk = 1/x'(z_k)$, a second derivative gives (2.31); subtracting (2.30) gives (2.32). $\square$

Pricing is thus the difference between the share delivered on exercise and the strike paid on the same event. The digital probability is $1 - u_k$, and the density is one division of two arrays already present in the transport. No price differentiation is required in the implementation.

The same rank change of variables prices any integrable terminal payoff:

$$\mathbb{E}[H(Y)] = \int_{-\infty}^{\infty} H(e^{x(z)})\rho(z) dz. \quad (2.33)$$

The ledger is the specialization in which $H(y) = y\mathbf{1}\{y > e^{x(z_0)}\}$ and the threshold z_0 is allowed to vary. Calls are unusually cheap because their payoff splits into this asset tail and an ordinary probability tail. Digitals retain only the probability tail; power and log contracts become moments of the transport. The whole payoff calculus therefore remains on one fixed rank measure.

The constant-speed slice gives a concrete first calculation. Choose its scale so that the six-month ATM volatility is exactly twenty percent. The solution is $s = 0.081250$ and the closed-form shift is $m = -0.010883$. For the call with $k = 0.10$, the strike rank is $u_k = 79.65\%$. The upper-share entry is $G(z_k) = 0.247221$, while the cash entry is $e^k(1 - u_k) = 0.224876$. Hence

$$c(0.10) = 0.247221 - 0.224876 = 0.022345,$$

corresponding to 20.60% implied volatility. The symmetric law produces a smile because its exponential return tails carry more remote-strike value than the lognormal matched at the money.

The fitted law also gives the log-contract variance directly:

$$w_{\text{vs}} = -2\mathbb{E}[\log Y] = -2 \int_{-\infty}^{\infty} x(z)\rho(z) dz. \quad (2.34)$$

This is the direct-law form of the familiar strip replication: the classical construction synthesizes $-\log Y$ from calls and puts over all strikes [7, 21], while here the law is already in hand and the same payoff is integrated against its rank measure. Under the usual continuous-path, continuous-monitoring identity this is the fair variance-swap strike; jumps and discrete sampling take their standard separate corrections.

2.5.2 The ATM microscope

At the money, the smile is determined locally by three quantities from the law. Let z_* solve $x(z_*) = 0$, and write

$$c_0 = G(z_*) - (1 - u_*), \quad u_* = \Lambda(z_*), \quad f_* = \frac{\rho(z_*)}{x'(z_*)}. \quad (2.35)$$

The ATM Black variance is

$$w_0 = 4 \left[\Phi^{-1} \left(\frac{1 + c_0}{2} \right) \right]^2, \quad d = \frac{\sqrt{w_0}}{2},$$

and define the digital mismatch

$$\Delta = u_* - \Phi(d). \quad (2.36)$$

Proposition 2.12 (Exact ATM handles). *The total-variance smile $B(k, w(k)) = c(k)$ satisfies*

$$w(0) = w_0, \quad w'(0) = \frac{2\sqrt{w_0}}{\varphi(d)}\Delta, \quad (2.37)$$

$$w''(0) = \frac{2\sqrt{w_0}}{\varphi(d)} f_* - 2 + \left(\frac{1}{2w_0} + \frac{1}{8} \right) w'(0)^2. \quad (2.38)$$

In volatility units, $\sigma(k) = \sqrt{w(k)/\tau}$,

$$\sigma(0) = \sqrt{w_0/\tau}, \quad \sigma'(0) = \frac{\Delta}{\varphi(d)\sqrt{\tau}}, \quad (2.39)$$

$$\sigma''(0) = \frac{1}{\sqrt{\tau}} \left[\frac{f_*}{\varphi(d)} - \frac{1}{\sqrt{w_0}} + \frac{\sqrt{w_0}}{4} \left(\frac{\Delta}{\varphi(d)} \right)^2 \right]. \quad (2.40)$$

The full differentiation is given in section 2.A.2. The first derivative contains the main idea. At $(k, w) = (0, w_0)$,

$$\partial_k B = -(1 - \Phi(d)), \quad \partial_w B = \frac{\varphi(d)}{2\sqrt{w_0}}, \quad c'(0) = -(1 - u_*).$$

Implicit differentiation therefore gives

$$w'(0) = \frac{c'(0) - \partial_k B}{\partial_w B} = \frac{2\sqrt{w_0}}{\varphi(d)} \{u_* - \Phi(d)\}. \quad (2.41)$$

The model digital is $1 - u_*$; the matched flat-Black digital is $1 - \Phi(d)$. Hence Δ is Black digital minus model digital. A negative equity skew is exactly a model digital richer than the flat-Black digital. Level, digital, and density determine level, skew, and curvature.

Three sign checks make the interpretation concrete. For the matched Black law, $X \sim \mathcal{N}(-w_0/2, w_0)$, so $u_* = \mathbb{P}(X \leq 0) = \Phi(d)$ and the skew vanishes. If $\sigma'(0) < 0$, then $u_* < \Phi(d)$ and the model digital exceeds the Black digital. Conversely, an upward skew is exactly a cheaper model digital. The comparison is with $\Phi(d)$, not with one half: the martingale shift may place the forward rank on either side of the median even when the skew is small.

The curvature formula adds one new local fact: f_* . Holding ATM price and digital fixed, a larger density at the forward raises $w''(0)$. This matches the geometry of a peaked distribution: more mass close to the forward makes nearby options cheap relative to options a little farther away, and the smile bends upward more rapidly. The other terms in (2.38) are the Black-coordinate correction and a quadratic skew contribution. Price level, first CDF value, and first density value form a hierarchy for the three smile handles.

The identities also run in reverse. A target ATM level fixes c_0 ; a target skew fixes u_* through (2.37); and a target curvature fixes f_* through (2.38). These are local constraints on the law rather than finite differences of a plotted smile.

2.5.3 Envelope cancellation

Calibration requires derivatives with respect to every parameter. The call depends on a parameter both through G and through the moving root z_k . The same cancellation used for (2.30) removes the latter.

Theorem 2.13 (Envelope cancellation). *For every coordinate θ_j ,*

$$\frac{\partial c(k; \theta)}{\partial \theta_j} = \frac{\partial G(z; \theta)}{\partial \theta_j} \Big|_{z=z_k(\theta)}. \quad (2.42)$$

Proof. Differentiate (2.28):

$$\frac{\partial c}{\partial \theta_j} = \frac{\partial G}{\partial \theta_j} \Big|_{z_k} + \{G'(z_k) + e^k \rho(z_k)\} \frac{\partial z_k}{\partial \theta_j}.$$

The bracket vanishes by (2.26) and $x(z_k) = k$. □

This is an envelope theorem: the exercise boundary is the point where the payoff integrand vanishes, so its displacement has no first-order value.

Because g is affine in the coefficients, all columns are obtained by two cumulative passes. With $\bar{x} = x - m$ and $I = \int e^{\bar{x}} \rho$, let $b_j(u) = \partial g(u) / \partial \theta_j$. Then

$$\frac{\partial \bar{x}(z)}{\partial \theta_j} = \int_0^z e^{g(\Lambda(t))} b_j(\Lambda(t)) dt, \quad (2.43)$$

$$\frac{\partial m}{\partial \theta_j} = -\frac{1}{I} \int_{-\infty}^{\infty} e^{\bar{x}(z)} \frac{\partial \bar{x}(z)}{\partial \theta_j} \rho(z) dz, \quad (2.44)$$

$$\frac{\partial G(z)}{\partial \theta_j} = \int_z^{\infty} e^{x(t)} \left(\frac{\partial m}{\partial \theta_j} + \frac{\partial \bar{x}(t)}{\partial \theta_j} \right) \rho(t) dt. \quad (2.45)$$

The basis columns are fixed functions,

$$b_L(u) = 1 - u, \quad b_R(u) = u, \quad b_{a_n}(u) = P_n(1 - 2u).$$

The first pass differentiates the transport speed and integrates it. The normalization dot product removes the component that would change $\mathbb{E}[e^X]$. The reverse pass differentiates the upper-share tail. Every stage follows the direction of the corresponding primal build, so the derivative is obtained by array operations rather than a separate root-and-price routine for each coefficient.

The variance-swap row is obtained from the same sensitivity of x :

$$\frac{\partial w_{vs}}{\partial \theta_j} = -2 \int_{-\infty}^{\infty} \left(\frac{\partial m}{\partial \theta_j} + \frac{\partial \bar{x}(z)}{\partial \theta_j} \right) \rho(z) dz. \quad (2.46)$$

Both analytic and finite-difference Jacobians scale as $O(Pn_{\text{grid}})$ for $P = N + 1$ parameters, but the analytic route avoids $P + 1$ complete rebuilds. In the reference implementation it reaches the same optimum, with costs agreeing to 1.7×10^{-10} relative and parameters to 3.7×10^{-7} , in 1.41 – $1.76 \times$ less time. The worst relative disagreement with an independent central difference is 3.7×10^{-6} .

The theorem is exact for the continuous curves. Numerically, x and G are interpolated separately, so between nodes the price interpolant is not algebraically identical to the integral of the quantile interpolant. The analytic derivative inherits this fourth-order interpolation error. The finite-difference comparison in fig. 2.8 measures the discrete discrepancy instead of assuming the continuous cancellation transfers without loss.

It is worth counting the shared work. One build stores (x, x', G, G') on a fixed grid. A price is one root and two ledger entries; a density is one division; the log-contract variance is one weighted sum; the Jacobian reads ledger sensitivities at the same roots; and calendar order, in section 2.8, will compare two such arrays. These are not five approximations to the law; they are five queries of the same law.

2.6 Calibration

Validity has already been settled by the coordinates. Calibration may therefore concentrate on information: selecting, among admissible laws, one that represents the quoted strip without inventing unsupported detail.

2.6.1 A one-expiry objective

Suppose a slice contains quotes (k_i, σ_i) with $w_i = \sigma_i^2 \tau$. The reference fit minimizes

$$\sum_i \left(\frac{c(k_i; \theta) - B(k_i, w_i)}{\partial_\sigma B(k_i, w_i) + \eta} \right)^2 + \alpha \sum_{n \geq 4} n^{2\nu} a_n^2. \quad (2.47)$$




Figure 2.8: The analytic Jacobian on the fitted SPY December 2026 slice. Left: the normalized price-sensitivity pattern of every parameter across strikes. Right: columnwise disagreement with an independent central finite difference. The comparison tests the entire sensitivity pass by a separate route.

Each element of (2.47) answers one practical problem. The residual is computed in price, so implied-volatility inversion stays outside the optimization loop, where it cannot fail on an exploratory trial point. Division by the Black vega $\partial_\sigma B = \varphi(d_+) \sqrt{\tau}$ makes a central residual a volatility error — the unit in which quotes are judged — while the floor η lets the metric degrade to a scaled price error where vega vanishes, instead of assigning unbounded weight to a negligible option. The $n^{2\nu}$ ridge makes fine modes progressively more expensive but leaves a_2 and a_3 , the main curvature and asymmetry modes, unpenalized.

The reference objective also contains a smooth row $\log(1 + e^{50(\lambda_+ - 0.90)})$. The admissibility chart already prevents the wall from being reached; this row is instead an explicit prior against a nearly exhausted right-tail moment budget.

A quoted option is an interval, not necessarily a point. Three targets are useful. A *mid* fit uses (2.47) directly. A *band* fit uses the outside-band hinge together with a weak mid anchor. If $[p_i^{\text{lo}}, p_i^{\text{hi}}]$ is the price band and ϕ_i the vega normalizer, its two residuals are

$$\frac{(c_i - p_i^{\text{hi}})^+ + (p_i^{\text{lo}} - c_i)^+}{\phi_i}, \quad \sqrt{0.05} \frac{c_i - p_i^{\text{mid}}}{\phi_i}. \quad (2.48)$$

The squared hinge is continuously differentiable at a band edge, and the weak mid row — its weight deliberately small — selects one point when a wide band leaves a flat valley of equally admissible laws. A *haircut* fit first shrinks the volatility band by h :

$$\sigma_i^{\text{lo}} = \min(\sigma_i^{\text{bid}} + h, \sigma_i^{\text{mid}}), \quad \sigma_i^{\text{hi}} = \max(\sigma_i^{\text{mid}}, \sigma_i^{\text{ask}} - h). \quad (2.49)$$

If the original spread is narrower than $2h$, the band collapses to the mid instead of changing orientation. This occurs for every SPY quote in the frozen sample; the wider NVDA spreads retain live bands.

The two underliers illustrate both regimes. The SPY December node has median spread 4 vol bp, far inside the default $2h = 100$ vol-bp width, and none of its haircut bands survives. Its figures therefore show raw bid–ask intervals around a mid fit. On the featured NVDA nodes, 47% and 52% of the bands remain live. The same rule becomes a mid target on a liquid index and a genuine interval target on the wider single-name book.

The informational division can be read by location. Between liquid quotes a small change in the law moves observed prices and is strongly penalized: the strip determines the call curve there, and with it the ATM level, skew, curvature, and low-order moments. Near the edge of the

strip, data and regularization share control; beyond the outer quote, the endpoint speeds and the ridge supply most of the continuation. The curve itself is smooth — what changes gradually is the source of information, a transition made visible by plotting quote locations with the density as in fig. 2.3. Tail coordinates, ridge, and basis order should be reported with a fit because they make these choices.

2.6.2 Order and starting values

At order N the model has $N + 1$ parameters. The reference implementation uses the effective order

$$N_{\text{eff}} = \max \left(\min \left(N, \left\lfloor \frac{n_q}{2} \right\rfloor - 1 \right), \min(N, 6) \right), \quad (2.50)$$

where n_q is the quote count. The cap keeps the parameter count well below the number of observations; the floor leaves thin strips enough shape to form a usable smile. The rule is an identification convention, not a validity condition. Changing it changes inter-quote density detail while every resulting slice remains arbitrage-free.

Away from the floor, (2.50) requires roughly two quotes per body degree of freedom. This is not a theorem about optimal smoothing; it is a guard against inferring one oscillatory direction from each noisy observation. On the shortest NVDA slice, 17 quotes give $N_{\text{eff}} = 7$. The fit remains visibly less certain there rather than borrowing high-order detail from the basis.

Cold starts use the constant-speed value (2.19). In a maturity sweep, the previous expiry supplies a warm start. A multi-start check on the featured SPY and NVDA nodes used 10 randomized starts per node and found 1 basin, with worst fit-quality spread 0.00 vol bp. This supports the starting rule on those slices; event-dominated books may still require a broader search.

2.6.3 Coordinates for optimization

Three charts describe the same polynomial family.

1. The *raw chart* is $(L, R, a_2, \dots, a_N)$. It is simple, but (2.14) couples body modes to both tails.
2. The *endpoint chart* is $(\log \lambda_-, \log \lambda_+, a_2, \dots, a_N)$. Solving (2.14) for (L, R) is linear. Equivalently, each body basis function is replaced by an endpoint-vanishing combination. Body moves then hold both tail scales fixed.
3. The *logistic chart*, the reference default, writes $\lambda_+ = \text{expit } \xi$ with $\xi \in \mathbb{R}$. Every finite coordinate vector is admissible, and

$$\frac{d \log \lambda_+}{d \xi} = 1 - \lambda_+. \quad (2.51)$$

Steps shrink smoothly as the wall is approached.

The endpoint-vanishing body functions are explicit:

$$\tilde{P}_n(u) = P_n(1 - 2u) - (1 - u) - (-1)^n u, \quad n \geq 2. \quad (2.52)$$

They satisfy $\tilde{P}_n(0) = \tilde{P}_n(1) = 0$. A tail policy can therefore set (λ_-, λ_+) while calibration adjusts the body in a subspace that cannot move either endpoint.

The objective is identical in all three charts. Fits in the frozen sample agree across charts to 2.2×10^{-7} in the recovered raw parameters. In floating point, $\text{expit } \xi$ eventually rounds to one, so the implementation also enforces $\lambda_+ \leq 1 - 10^{-6}$.

2.6.4 Calibration sequence

The complete one-slice calculation is short enough to state as an algorithm.

1. Convert strikes and prices to (k_i, c_i) using the node's forward and discount factor. Map bid and ask volatilities through Black once to construct price bands and compute the quote vegas used for normalization.
2. Choose the effective order from (2.50) and select mid, band, or haircut residuals. This fixes the data part of the objective before optimization begins.
3. Construct a cold start from the ATM variance using (2.19), or map the previous expiry into the endpoint chart for a warm start. In either case, convert to logistic coordinates so all finite trial vectors are admissible.
4. At each trust-region step, build (x, x', G, G') on the fixed grid of section 2.7, price every quote by (2.28), and obtain the whole residual Jacobian from (2.43)–(2.45). Failed exponent-budget trials return a large residual rather than an undefined price.
5. After convergence, rebuild the accepted parameters on the dense grid. Evaluate ATM handles, tail moments, Lee slopes, the log-contract variance, and the numerical certificates (section 2.7.2) from this one final object.
6. If the slice belongs to a maturity sequence, add calendar rows against the accepted nearer slice and repeat. Report any residual common-support gap in normalized price units.

This order matters. Market conventions are fixed before the solve; probability validity holds during the solve; diagnostics are computed after the solve from the same law that produced the quoted prices. No separate post-processing smoother is inserted between calibration and publication.

2.7 Numerical realization

The theorems of the preceding sections concern the continuous objects x and G ; the fitter computes with discrete realizations of them. One grid build produces every array used above, analytic corrections handle what lies beyond the grid, and two certificates verify that discretization has preserved what the theory guarantees.

2.7.1 One grid build

All integrals use a symmetric log-odds grid $[-Z_{\max}, Z_{\max}]$, with $Z_{\max} = 40$. Optimization uses 2001 nodes; the accepted slice is rebuilt on 8001 nodes. A build performs four operations: evaluate g , integrate e^g cumulatively to form $\bar{x}$, compute the martingale shift, and integrate $e^x \rho$ backward to form G . Exact nodal derivatives $x' = e^g$ and $G' = -e^x \rho$ accompany the values. Cubic Hermite interpolation then gives fourth-order off-grid evaluation and a Newton-polished strike root.

The missing logistic tails are integrated analytically. To leading order,

$$\int_{Z_{\max}}^{\infty} e^{\bar{x}} \rho \, dz \simeq \frac{e^{\bar{x}(Z_{\max}) - Z_{\max}}}{1 - \lambda_+}, \quad \int_{-\infty}^{-Z_{\max}} e^{\bar{x}} \rho \, dz \simeq \frac{e^{\bar{x}(-Z_{\max}) - Z_{\max}}}{1 + \lambda_-}. \quad (2.53)$$

For a strike beyond the right return grid, continue the root along the tail asymptote,

$$z_R(k) = Z_{\max} + \frac{k - x(Z_{\max})}{\lambda_+}.$$

Both the ledger and cash leg are then single exponentials. Their difference is

$$c(k) \simeq e^{k-z_R(k)} \frac{\lambda_+}{1-\lambda_+}. \quad (2.54)$$

The left-wing continuation follows from parity.

To derive the factor in (2.54), restart the right-tail correction at $z_R(k)$, where $x = k$. The asset leg is approximately $e^{k-z_R(k)}/(1-\lambda_+)$ and the cash leg is $e^{k-z_R(k)}$. Subtraction leaves $\lambda_+/(1-\lambda_+)$. The continuation is therefore the tail form of the same asset-minus-cash identity (2.28), not an independent wing model.

Three floating-point rules are essential. First, compute $\rho(z) = \text{expit}(z) \text{expit}(-z)$ rather than $u(1-u)$, since u rounds to one for large positive z . Compute the cash leg as $\exp\{k - \log(1 + e^{z_k})\}$ using a stable log-sum-exp. Second, reject a trial if an interior value of g or $\bar{x}$ exceeds the exponent budget, even when the endpoints are admissible. Third, evaluate the Lee function at large r in the rationalized form

$$\Psi(r) = 2 - \frac{4}{\sqrt{1 + 1/r} + 1}, \quad (2.55)$$

which avoids subtracting nearly equal large numbers. The full build and sensitivity pseudocode appear in appendix 2.B.

2.7.2 Certificates

The continuous theorems concern x and G ; numerical prices use their interpolants. Two inexpensive certificates connect the two levels. On each Hermite segment, the Fritsch–Carlson inequalities require both endpoint derivatives to lie between zero and three times the secant slope [10]. Applying them to x and $-G$ certifies monotone interpolation; segments below round-off are declared flat to a stated tolerance. Across the traded strike span, the analytic derivatives of $w(k)$ are also inserted in (2.5) on a dense grid and required to satisfy $g_D \geq -10^{-4}$.

An independent randomized battery prices 27 admissible slices across several orders and stress classes at sub-grid strikes. Its worst errors in call bounds, butterflies, digitals, and agreement with a four-times-finer grid are respectively 1.1×10^{-11} , 8.2×10^{-13} , 1.5×10^{-12} , and 2.3×10^{-9} in normalized price. The exact constant-speed test of fig. 2.4 independently checks transport, normalization, tail correction, and root evaluation. These numerical tests do not create validity; they verify that the discrete realization has preserved it to the reported tolerance.

The two-grid rebuild supplies a final resolution check. Parameters fitted on 2001 nodes and refitted on 8001 nodes agree to 1.4×10^{-6} . The accepted vector is rebuilt once on the dense grid, so prices, densities, ATM handles, and calendar ledgers all read the same final discrete object.

These diagnostics answer different questions and should not be collapsed into one flag. Quote residuals measure agreement with the strip; effective order and ridge describe the asserted inter-quote structure; endpoint speeds describe the extrapolation and its moment budget; monotonicity and belly margins certify that discretization preserved validity; calendar gaps describe compatibility across expiries. A fit can be excellent on one axis and weak on another.

2.8 Calendar order

Butterfly freedom is a property of one marginal law. A surface also requires later expiries to dominate earlier ones in convex order. Under deterministic rates and carry, this comparison has a simple form in normalized units.

Lemma 2.14 (Calendar ordering at equal log-moneyness). *Let $F_{0,t}$ be the deterministic forward curve and suppose $M_t = S_t/F_{0,t}$ is a positive martingale. For $T_1 < T_2$ and every k , the normalized calls satisfy*

$$c_1(k) \leq c_2(k). \quad (2.56)$$

Proof. Since $M_{T_i} = S_{T_i}/F_i = Y_i$, conditional Jensen gives

$$\begin{aligned} c_2(k) &= \mathbb{E}[(M_{T_2} - e^k)^+] \\ &\geq \mathbb{E}[(\mathbb{E}[M_{T_2} | \mathcal{F}_{T_1}] - e^k)^+] = \mathbb{E}[(M_{T_1} - e^k)^+] = c_1(k). \end{aligned}$$

□

Equal k means strikes $K_i = F_i e^k$, which are generally different dollar strikes. The comparison is a forward-scaled calendar relation, not the usual listed equal-strike spread. In probability language, the laws of Y_i have equal means and increase in convex order.

The scaling is essential. A violation at k is monetized by shorting $1/(D_1 F_1)$ calls at $K_1 = F_1 e^k$ and buying $1/(D_2 F_2)$ calls at $K_2 = F_2 e^k$. After normalization both payoffs are convex functions of the common martingale M_t . Without deterministic carry, one must first specify the numeraire and the objects being compared; equal log-moneyness alone is not a model-free calendar relation.

2.8.1 Calls and ledgers are conjugate

Lemma 2.15 (Ledger–call conjugacy). *For any admissible slice,*

$$\begin{aligned} c(k) &= \max_{z \in [-\infty, \infty]} \{G(z) - e^k(1 - \Lambda(z))\}, \\ G(z) &= \min_{k \in [-\infty, \infty]} \{c(k) + e^k(1 - \Lambda(z))\}. \end{aligned} \tag{2.57}$$

The optimizers are $z = z_k$ and $k = x(z)$, respectively.

Proof. For fixed k ,

$$G(z) - e^k(1 - \Lambda(z)) = \int_z^\infty (e^{x(t)} - e^k) \rho(t) dt.$$

The integrand is negative before z_k and positive after it. Starting the integral at z_k therefore maximizes it and gives the call price. The first identity implies $c(k) + e^k(1 - \Lambda(z)) \geq G(z)$ for every k , with equality at $k = x(z)$, proving the second. □

Geometrically, the first identity scans all possible exercise ranks. Before z_k , the integral includes outcomes where the payoff is negative; after z_k , it discards outcomes where the payoff is positive. The strike root is the unique point of tangency between call and ledger descriptions. The second identity reconstructs the ledger from the full call curve, so no information is lost in changing representation.

Theorem 2.16 (Calendar order is ledger order). *For two admissible mean-one slices,*

$$c_1(k) \leq c_2(k) \text{ for all } k \iff G_1(z) \leq G_2(z) \text{ for all } z. \tag{2.58}$$

Proof. If $G_1 \leq G_2$, the maximum representation in (2.57) gives $c_1 \leq c_2$. If $c_1 \leq c_2$, the minimum representation gives $G_1 \leq G_2$. Both implications are pointwise. □

Thus the continuum of calendar inequalities is equivalent to ordering two arrays already produced by the slice builds. This is also the classical convex-order criterion: by the results of Strassen and Kellerer, a family of equal-mean laws increasing in convex order can be realized as the marginals of a martingale [18, 24].

This is an existence statement. Ordered ledgers make the marginal laws compatible with some martingale dynamics; they neither select that dynamics nor determine forward smiles. The calendar problem treated here is exactly the static marginal problem.

2.8.2 A support-confined control

The theorem is an exact diagnostic. A finite calibration requires a chosen enforcement rule. When fitting a farther expiry after a nearer one, the reference objective adds

$$\sqrt{w_{\text{cal}}} \zeta_m \{c_{\text{near}}(k_m) - c_{\text{far}}(k_m)\}^+ \quad (2.59)$$

on a grid of strikes in the intersection of their quoted spans. The weights ζ_m taper to zero with a $\cos^2$ profile over the outer fifteen percent of that span. The number of nodes is $\max(49, 4N + 1)$.

The rows are imposed in price space, not ledger space. A ledger value at one rank integrates the entire upper tail, so constraining it couples the quoted region of one expiry to the unquoted continuation of another: the ledger is ideal for diagnosis precisely because it sees the whole law, and unsuitable as a constraint for the same reason. In a reproduced stress comparison, a full-grid ledger floor moved the worst quote error from 10.6 to 1095.0 vol bp; fixed-strike call rows, each a specified payoff in the common quoted region, localize the control. The weight is finite, and any remaining violation is reported as a price gap and as a fraction of the local bid–ask width. Calendar order is thus controlled on common quoted support; outside it, the ledger comparison remains a diagnostic of the extrapolations.




Figure 2.9: Calendar order in the frozen SPY sample. Left: total variance is ordered through maturity at representative log-moneyness values. Right: the two shortest expiries are ordered on their common quoted span, while a volatility-space crossing appears only in a numerically empty far wing. The price-space gap there is at round-off scale.

For the two shortest SPY expiries, a full-grid volatility audit flags a 1840.6 vol-bp crossing at $k = +0.981$, far outside the common quote span $[-0.028, +0.016]$. The corresponding calls are 5.4×10^{-107} and 1.6×10^{-26} of forward. A direct price audit finds a worst violation of only 2.2×10^{-15} . The crossing is a disagreement between extrapolations in a region where Black inversion is ill-conditioned, not an economically available calendar trade.

2.9 Examples and the model contract

All market examples use one frozen after-hours snapshot of SPY and NVDA listed options from the 2026-08-03 US session. Each underlier has eight expiries, for 16 fitted nodes in total. The common recipe is order $N = 16$ subject to (2.50), the logistic chart, the haircut target (2.49), ridge $\alpha = 10^{-6}$, and the calendar rows (2.59). No node is tuned separately. Across the sample the median quote error is 5.9 vol bp and the worst is 24.6 vol bp.




Figure 2.10: Eight SPY expiries fitted with one recipe. Each panel shows raw bid-ask quotes, mids, the fitted LQD slice, and rms error. The short end is steeper and more localized; the long end broadens smoothly. Beyond the last quote, each curve is determined by the fitted exponential-tail convention.

2.9.1 SPY: a term structure and one slice

The gallery separates recurring market shape from the completion chosen by the model. Through the quoted region the slices follow the term structure of level and skew; outside it the two endpoint speeds continue the curve. The boundary is most visible at the short end, where quoted spans are narrow in log-moneyness: the sharp local skew is supported by quotes, while the remote upward turn is the exponential-tail completion. At longer maturities the span widens and a larger fraction of the displayed smile is directly informed by prices. The median error across the eight slices is 3.6 vol bp; the worst is 7.5 vol bp.




Figure 2.11: The SPY December 2026 slice. Left: quotes and fitted volatility. Center: signed residuals in vol bp. Right: the fitted density against an ATM-matched normal. The density view makes the negative skew visible as a transfer of mass toward the left tail.

The December 2026 slice contains 94 quotes and fits them to 3.3 vol bp rms. Its exact ATM quantities from proposition 2.12 are volatility 14.76%, skew -0.430 , and curvature 4.40. The

forward rank is 41.31%%, below the matched Black rank; hence $\Delta < 0$, in agreement with the negative skew. The log-contract variance corresponds to 19.84% annualized volatility.

The three panels answer different questions. The first tests the nonlinear map from law to implied volatility. The second asks whether residuals retain systematic strike structure. The third removes the Black coordinate and shows the fitted object itself: a left-heavy density. The ATM-matched normal has the same central price level but not the same digital or tail allocation, and therefore cannot reproduce the skew.

The ledger also makes an individual option transparent. For the listed December 800 call, $k = 0.0409$ and the strike rank is 65.27%%. The upper-share and cash entries are 0.378366 and 0.361834, so

$$c(0.0409) = 0.378366 - 0.361834 = 0.016532.$$

This is \$12.70 per share and corresponds to 13.37% implied volatility.

2.9.2 NVDA and the order guard




Figure 2.12: Two NVDA slices. Left: the shortest expiry, with 17 quotes and effective order $N_{\text{eff}} = 7$. Right: the December 2027 expiry at full order. Wider NVDA spreads leave live haircut bands and permit visibly larger residual variation than SPY.

The shortest node fits at 24.6 vol bp rms; the long node, with 69 quotes, at 7.4 vol bp. The latter has heavier fitted tails than SPY. In particular its left endpoint speed exceeds one, which is admissible: only the right endpoint is constrained by the finite-forward wall. The consequence is a last finite inverse-moment order 0.70 below one.

The shortest node is also the clearest example of the order guard. A thin strip cannot distinguish sixteen smooth modes, even though an optimizer can assign values to them. Reducing N_{eff} converts that lack of information into visible residual uncertainty instead of invisible oscillation between quotes. The long node has enough observations to use the full order and separate body curvature from tail continuation.

2.9.3 What basis order asserts

The synthetic benchmark is a martingale-normalized mixture of two lognormal components. It has a genuinely bimodal density but only gentle shoulder structure in implied volatility.




Figure 2.13: The double-hump benchmark at $N = 16$. Left: the fitted smile matches the synthetic quotes to 4.8 vol bp rms. Right: the fitted density recovers both modes, the trough, and the asymmetry of the target.




Figure 2.14: The same target at orders 6 and 16. Left: the implied-volatility curves nearly coincide over the quoted region. Right: the densities do not; the lower-order fit displaces the modes and fills part of the trough. Basis order controls how much inter-quote structure the model is allowed to assert.

At $N = 16$ the fitted mode locations are -0.099 and 0.088, against true locations -0.101 and 0.089. At $N = 6$ both modes remain, but their shape is distorted. The L^1 density error rises from 0.047 to 0.108, even though the two implied-volatility curves differ by at most 41.8 vol bp on the comparison grid. Option quotes can be nearly indifferent between laws that price density-sensitive structures differently.

The lesson is not that high order is always preferable. At $N = 16$ the synthetic target is known to contain fine structure, so extra modes recover a feature known to be real. In market data the truth is not supplied. High order then grants the model permission to assert finer inter-quote density shape. Digitals and narrow butterflies are sensitive to that permission; broad vanilla prices may not be. Basis order should be chosen for the downstream instruments, not solely for the in-sample volatility residual.

2.9.4 Position among existing models

The construction joins several classical ideas. Breeden–Litzenberger identify option prices with a density [6]; Lee and Benaim–Friz connect moments with wing slopes [2, 19]; quantile modeling supplies the rank coordinate [14, 22]; and Petersen–Müller use the same log-quantile-density transform to place density-valued data in an unconstrained function space [23]. The metalog family is a related direct quantile expansion [17].

SVI and SSVI work in volatility coordinates and characterize admissible parameter regions [11, 13]; arbitrage-free smoothing works directly with constrained prices [9]. LQD instead makes the probability law primary. Its contribution is the combined architecture: structural one-slice validity, readable tail coordinates, and one ledger for pricing, differentiation, density, variance, and calendar order.

2.9.5 The contract

Two things remain: how to read a fitted slice, and what precisely is claimed.

A fitted slice should be read in the same order as the theory. First inspect quote residuals over the retained span. Then read the three ATM handles as local level, digital, and density information. Next inspect $g(u)$ and the density over the ranks occupied by quotes. Only after that should one read $\lambda_{\pm}$, the moment strip, and Lee slopes, since these describe the extrapolated endpoints. Finally compare the ledger with neighboring expiries. This sequence prevents a tail convention or a calendar diagnostic from being mistaken for a directly observed quote.

Structural	Numerically controlled	Outside the claim
Positive density, unit mean, and butterfly-free calls for every admissible slice	Interpolation monotonicity, traded-range density, and randomized price audits	Atoms, mass at zero, density gaps, bounded support, and Gaussian-or-faster tails
Exact exponential tail scales, moment strip, and Lee limits	Tail quadrature, exponent budgets, grid convergence, and finite-strike wing prices	Tail parameters as observations, or asymptotic slopes as finite-strike quotes
Exact ATM handles, analytic Jacobian, and ledger characterization of calendar order	Quote fit, basis-order control, multi-start checks, and calendar order on common quoted support	Dynamics, jump decomposition, realized-variance corrections, and calendar structure outside the controlled support

Table 2.2: The model contract: theorem, numerical control, and excluded scope.

The central distinction is now visible in one place. Probability supplies validity. Quotes supply information over a finite strip. Basis, ridge, tails, and calendar scope supply convention elsewhere. The value of the LQD representation is not that convention disappears, but that it can no longer be mistaken for a theorem or for market information.

2.A Proofs deferred from the main text

2.A.1 Derivation of the volatility-space density factor

Let $c(k) = B(k, w(k))$. For general (k, w) , the identity $e^k \varphi(d_-) = \varphi(d_+)$ gives

$$\partial_k B = -e^k \Phi(d_-), \quad \partial_w B = \frac{\varphi(d_+)}{2\sqrt{w}}.$$

The second partials required by the chain rule are

$$\partial_{kk} B = -e^k \Phi(d_-) + \frac{\varphi(d_+)}{\sqrt{w}},$$

$$\begin{aligned}\partial_{kw}B &= \frac{d_+\varphi(d_+)}{2w}, \\ \partial_{ww}B &= -\frac{d_+\varphi(d_+)}{2\sqrt{w}}\partial_w d_+ - \frac{\varphi(d_+)}{4w^{3/2}}, \quad \partial_w d_+ = \frac{k}{2w^{3/2}} + \frac{1}{4\sqrt{w}}.\end{aligned}$$

Differentiate $c(k)$ twice and subtract the first derivative:

$$c'' - c' = (\partial_{kk}B - \partial_k B) + (2\partial_{kw}B - \partial_w B)w' + \partial_{ww}B(w')^2 + \partial_w B w''.$$

Factor $\varphi(d_+)/\sqrt{w}$. Since $d_+ = -k/\sqrt{w} + \sqrt{w}/2$,

$$2\partial_{kw}B - \partial_w B = -\frac{k\varphi(d_+)}{w^{3/2}},$$

and a direct multiplication gives

$$d_+\partial_w d_+ = -\frac{k^2}{2w^2} + \frac{1}{8}.$$

Therefore

$$\begin{aligned}c'' - c' &= \frac{\varphi(d_+)}{\sqrt{w}} \left[1 - \frac{kw'}{w} + \left(\frac{k^2}{4w^2} - \frac{1}{4w} - \frac{1}{16} \right) (w')^2 + \frac{w''}{2} \right] \\ &= \frac{\varphi(d_+)}{\sqrt{w}} \left[\left(1 - \frac{kw'}{2w} \right)^2 - \frac{(w')^2}{4} \left(\frac{1}{w} + \frac{1}{4} \right) + \frac{w''}{2} \right].\end{aligned}$$

Finally, $f_X(k) = e^{-k}(c'' - c')$ and $e^{-k}\varphi(d_+) = \varphi(d_-)$, which proves (2.5).

2.A.2 Proof of the ATM formulas

The smile is defined by $B(k, w(k)) = c(k)$. Put $w_0 = w(0)$ and $d = \sqrt{w_0}/2$. The identity $e^k\varphi(d_-) = \varphi(d_+)$ follows from $d_+^2 - d_-^2 = -2k$ and gives

$$\partial_k B = -e^k\Phi(d_-), \quad \partial_w B = \frac{\varphi(d_+)}{2\sqrt{w}}.$$

Differentiating once more and evaluating at $(0, w_0)$ gives

$$\begin{aligned}\partial_k B &= -(1 - \Phi(d)), & \partial_w B &= \frac{\varphi(d)}{2\sqrt{w_0}}, \\ \partial_{kk}B &= -(1 - \Phi(d)) + \frac{\varphi(d)}{\sqrt{w_0}}, & \partial_{kw}B &= \frac{\varphi(d)}{4\sqrt{w_0}}, \\ \partial_{ww}B &= -\frac{\varphi(d)}{4w_0^{3/2}} - \frac{\varphi(d)}{16\sqrt{w_0}}.\end{aligned}\tag{2.60}$$

By (2.30)–(2.32),

$$c'(0) = -(1 - u_*), \quad c''(0) = f_* - (1 - u_*).$$

One derivative of the defining identity gives $\partial_k B + \partial_w B w'(0) = c'(0)$, hence

$$w'(0) = \frac{-(1 - u_*) + (1 - \Phi(d))}{\varphi(d)/(2\sqrt{w_0})} = \frac{2\sqrt{w_0}}{\varphi(d)}\Delta.$$

Two derivatives give

$$\partial_{kk}B + 2\partial_{kw}B w' + \partial_{ww}B (w')^2 + \partial_w B w'' = c''.$$

Solving at the money and using (2.60),

$$\begin{aligned} w''(0) &= \frac{2\sqrt{w_0}}{\varphi(d)} \left[f_* + \Delta - \frac{\varphi(d)}{\sqrt{w_0}} - \frac{\varphi(d)}{2\sqrt{w_0}} w'(0) \right. \\ &\quad \left. + \left(\frac{\varphi(d)}{4w_0^{3/2}} + \frac{\varphi(d)}{16\sqrt{w_0}} \right) w'(0)^2 \right] \\ &= \frac{2\sqrt{w_0}}{\varphi(d)} f_* - 2 + \left(\frac{1}{2w_0} + \frac{1}{8} \right) w'(0)^2. \end{aligned}$$

The terms linear in Δ cancel because $w'(0) = 2\sqrt{w_0}\Delta/\varphi(d)$.

Finally, $\sigma = \sqrt{w/\tau}$ gives

$$\sigma' = \frac{w'}{2\sqrt{\tau w}}, \quad \sigma'' = \frac{w''}{2\sqrt{\tau w}} - \frac{(w')^2}{4\sqrt{\tau} w^{3/2}}.$$

Substitution yields (2.39)–(2.40).

2.A.3 Proof of universality and the closure dichotomy

Let g^* satisfy proposition 2.6, write $\lambda_+^* = e^{g^*(1)} < 1$, and fix $\bar{\lambda} \in (\lambda_+^*, 1)$.

Polynomial approximation. By Weierstrass there are polynomials p_N with $\varepsilon_N = \|p_N - g^*\|_\infty \rightarrow 0$. Every polynomial of degree at most N is represented by the endpoint chord and $P_2, \dots, P_N$. For all large N ,

$$p_N(1) \leq g^*(1) + \varepsilon_N < \log \bar{\lambda} < 0,$$

so the approximating slice is admissible.

Transport convergence and a common envelope. Let $\bar{x}_N$ be the corresponding unshifted transport. On every finite z -interval,

$$|\bar{x}_N(z) - \bar{x}^*(z)| \leq (e^{\varepsilon_N} - 1)e^{\|g^*\|_\infty} |z|,$$

and therefore $\bar{x}_N \rightarrow \bar{x}^*$ locally uniformly. Continuity of g^* at the right endpoint gives z_0 such that, for all large N , $e^{p_N(\Lambda(z))} \leq \bar{\lambda}$ when $z \geq z_0$. Hence a constant C_1 , independent of N , satisfies

$$e^{\bar{x}_N(z)} \leq C_1 e^{\bar{\lambda}z} \quad (z \geq 0), \quad e^{\bar{x}_N(z)} \leq C_1 \quad (z \leq 0). \quad (2.61)$$

Normalizers, shifts, and quantiles. The functions $e^{\bar{x}_N} \rho$ converge pointwise and, by (2.61), are dominated by integrable multiples of $e^{(\bar{\lambda}-1)z}$ on the right and e^z on the left. Dominated convergence gives $I_N \rightarrow I^*$, so $m_N = -\log I_N \rightarrow m^*$. Thus $x_N \rightarrow x^*$ locally uniformly and $Q_N \rightarrow Q^*$ uniformly on compact subsets of $(0, 1)$.

Uniform convergence of calls. Couple all laws with one uniform rank and let $Y_N = e^{Q_N(U)}$, $Y^* = e^{Q^*(U)}$. Since positive-part payoffs are one-Lipschitz,

$$\sup_k |c_N(k) - c^*(k)| \leq \mathbb{E}|Y_N - Y^*|.$$

The integrand converges pointwise. The envelope above and boundedness of m_N give a constant C_2 such that $e^{Q_N(u)} \leq C_2\{1 + (1-u)^{-\bar{\lambda}}\}$. This is integrable, so the calls converge uniformly.

Wasserstein convergence. The quantile coupling is optimal on the line:

$$W_p(Y_N, Y^*)^p = \int_0^1 |e^{Q_N(u)} - e^{Q^*(u)}|^p du.$$

For $1 \leq p < 1/\lambda_+^*$ choose $\bar{\lambda} \in (\lambda_+^*, 1/p)$. The integrand is dominated by a multiple of $1 + (1-u)^{-p\bar{\lambda}}$, proving the claim by dominated convergence.

Finite-order exclusions. At finite order, x'_N is positive and continuous. The quantile is a strictly increasing homeomorphism from $(0, 1)$ onto $\mathbb{R}$; it can neither jump, remain flat, nor terminate. This excludes atoms, density gaps, bounded support, and mass at zero. A Gaussian-or-faster tail would require $g(u) \rightarrow -\infty$ at an endpoint, which no polynomial can do.

Uniform closure away from the wall. Suppose admissible g_N converge uniformly to g_∞ . Necessarily $g_\infty(1) \leq 0$. If $g_\infty(1) < 0$, the preceding dominated-convergence argument applies with g_∞ as target. The limit retains a positive continuous density and exponential tails.

The boundary branch. If $g_\infty(1) = 0$, weak convergence may fail. For $g_N \equiv \log(1 - 1/N)$, the solvable slice gives

$$I_N = \frac{\pi(1 - 1/N)}{\sin(\pi/N)} \rightarrow \infty, \quad m_N = -\log I_N \rightarrow -\infty.$$

Every fixed-rank quantile tends to $-\infty$, so mass escapes.

Assume instead that the laws converge weakly. The shifts are bounded above because every normalizer has a positive central contribution. They are bounded below because $m_N \rightarrow -\infty$ along a subsequence would contradict tightness. Every convergent subsequence of shifts then yields

$$Q_\infty(u) = m_\infty + \int_0^{\text{logit } u} e^{g_\infty(\Lambda(t))} dt.$$

Weak convergence identifies this quantile uniquely, so the full shift sequence converges. The limiting speed is continuous and bounded away from zero; hence the limit has full support and a strictly positive continuous density, and all geometric exclusions remain.

The right endpoint speed is one. For each $\delta > 0$ it is eventually at least $1 - \delta$, so the argument of proposition 2.8 gives $\mathbb{E}[e^{rX}] = \infty$ for every $r > 1$. Fatou also gives $\mathbb{E}[e^X] \leq \liminf_N \mathbb{E}[e^{X_N}] = 1$, with equality not guaranteed. This is the boundary law in proposition 2.6.

2.B Numerical implementation

This appendix records the load-bearing parts of a reference implementation. The pseudocode follows the notation of the chapter; tail-sensitivity terms and error handling are described after the listings.

2.B.1 Slice build

Listing 2.1: Build the transport and ledger on one fixed grid.

```
def build(theta, grid):
    L, R, a = theta.L, theta.R, theta.a
    g = (1 - grid.u)*L + grid.u*R + a @ grid.Pmat
    if not finite(g) or max(g) > 700:
        raise InteriorOverflow

    dx = exp(g)                                # x'(z)
```

```

xbar = cumulative_simpson(dx, anchor=grid.mid)
if max(xbar) > 700:
    raise InteriorOverflow

lam_m, lam_p = exp(g[0]), exp(g[-1])
if lam_p >= 1 - 1e-6:
    raise InadmissibleTail

mass = exp(xbar) * grid.rho
I = integrate(mass) \
    + exp(xbar[-1] - grid.zmax)/(1 - lam_p) \
    + exp(xbar[0] - grid.zmax)/(1 + lam_m)
m = -log(I)
x = m + xbar

dG = -exp(x) * grid.rho # G'(z)
G = reverse_integrate(-dG) \
    + exp(x[-1] - grid.zmax)/(1 - lam_p)
return Slice(x=x, dx=dx, G=G, dG=dG,
             m=m, lam=(lam_m, lam_p), theta=theta)

```

The cached density must be computed as $\expit(z)\expit(-z)$. Forming $u*(1-u)$ loses the right tail when u rounds to one. Both x and G carry their exact nodal derivatives, so a cubic Hermite representation needs no global spline solve.

The median is the natural anchor for the cumulative transport. It fixes $\bar{x}(0) = 0$ and lets a single forward cumulative sum build the positive half-line while the same stencil, reversed, builds the negative half-line. The additive freedom is then carried entirely by m . Simpson quadrature is fourth order on the uniform grid; the Hermite interpolant has the same local order because both values and exact nodal derivatives are available.

The grid does not truncate the probability law. From (2.16), on the right

$$e^{\bar{x}(z)}\rho(z) = e^{b_+ + (\lambda_+ - 1)z} \{1 + O(e^{-z})\}.$$

Integrating its leading term from $Z_{\max}$ gives the first correction in (2.53); the relative error is $O(e^{-Z_{\max}})$. The left correction follows from $e^{\bar{x}(z)}\rho(z) \sim e^{b_- + (\lambda_- + 1)z}$. The factor $1 - \lambda_+$ in the right denominator is also a numerical warning: tail mass becomes increasingly sensitive as the finite-forward wall is approached.

2.B.2 Call evaluation

Listing 2.2: Evaluate a normalized call.

```

def call_price(slc, k, grid):
    if k > slc.x[-1]:
        z = grid.zmax + (k - slc.x[-1]) / slc.lam[1]
        return exp(k - z) * slc.lam[1] / (1 - slc.lam[1])
    if k < slc.x[0]:
        return 1 - exp(k) + put_asymptote_left(slc, k, grid)

    z = hermite_root(slc.x, slc.dx, k)
    share = hermite_eval(slc.G, slc.dG, z)
    cash = exp(k - logaddexp(0, z)) # e^{k(1-Lambda(z))}
    return share - cash

```

The log-space cash expression is essential. The algebraically simpler $\exp(k)*(1-\expit(z))$ becomes zero after rank saturation and can create a spurious upward step in the far call wing.

The two branches beyond the return grid are part of the same law. On the right, (2.54) continues the exact power tail. On the left, put asymptotics are computed first and call parity is applied. This avoids subtracting two quantities close to one when a deep in-the-money call is evaluated directly.

2.B.3 Sensitivity pass

Listing 2.3: All raw-coordinate Jacobian columns on one build.

```
def sensitivities(slc, grid):
    # rows: dg/dL, dg/dR, dg/da_2, ..., dg/da_N
    B = vstack([1-grid.u, grid.u, grid.Pmat])
    dxbar = cumulative_simpson(slc.dx * B,
                              anchor=grid.mid, axis=1)
    weight = exp(slc.x) * grid.rho
    dm = -integrate(weight * dxbar, axis=1)
    dx = dm[:, None] + dxbar

    integrand = weight[None, :] * dx
    dG = reverse_integrate(integrand, axis=1)
    dG_slope = -integrand
    dvs = -2 * integrate(dx * grid.rho, axis=1)
    return dG, dG_slope, dvs

def price_jacobian_row(slc, sens, k):
    z = hermite_root(slc.x, slc.dx, k)
    return hermite_eval(sens.dG, sens.dG_slope, z)
```

The derivatives of the two tail corrections in (2.53) must also be included. They use the endpoint sensitivities

$$\partial\lambda_-/\partial\theta = \lambda_-(1, 0, 1, 1, \dots), \quad \partial\lambda_+/\partial\theta = \lambda_+(0, 1, 1, -1, \dots).$$

Derivatives in the endpoint or logistic chart follow by multiplication with the corresponding chart Jacobian.

The analytic pass performs one cumulative integral for each basis column, one normalization dot product, and one reverse cumulative integral. Its work is $O(Pn_{\text{grid}})$ and its storage may be either $O(Pn_{\text{grid}})$, when all columns are retained, or $O(n_{\text{grid}})$ per streamed column. Finite differences have the same formal order but repeat normalization, ledger construction, and every strike root $P + 1$ times. Envelope cancellation removes the roots from the analytic derivative and is therefore responsible for most of the measured constant-factor gain.

2.B.4 Interpolation certificate

Listing 2.4: Fritsch–Carlson margin for a uniform-grid Hermite interpolant.

```
def monotone_margin(y, d, dz, flat_tol=1e-9):
    sec = diff(y) / dz
    d0, d1 = d[:-1], d[1:]
    active = ~((abs(sec) < flat_tol) &
               (abs(d0) < flat_tol) &
               (abs(d1) < flat_tol))
    margin = minimum.reduce([d0, d1,
                            3*sec-d0, 3*sec-d1])
    return min(margin[active])
```

An increasing interpolant is certified when the returned margin is positive. Apply the same routine to $-G$. Segments excluded by `flat_tol` are recorded as flat to tolerance rather than strictly monotone.

For a cubic Hermite segment, let s be its secant slope and d_0, d_1 its endpoint derivatives. The sufficient Fritsch–Carlson box is

$$0 \leq d_0 \leq 3s, \quad 0 \leq d_1 \leq 3s.$$

The routine returns the minimum distance to the four faces of this box. A positive number is therefore not merely a pass flag but a margin. For smooth functions sampled on a fine grid, $d_i/s = 1 + O(h^2)$, so the certificate should clear the boundary comfortably. A small margin is a useful indication of under-resolution even before monotonicity fails.

Monotonicity of x and G is not by itself a convexity certificate for the interpolated call. The separate belly test uses the exact factorization (2.5): analytic w, w', w'' are evaluated at 801 points across the retained strike span and g_D must exceed -10^{-4} . Since a flat smile has $g_D = 1$, the tolerance permits a negative density component no larger than 10^{-4} of the local Black density factor. The threshold is therefore dimensionless and interpretable.

2.B.5 Floating-point failure catalogue

Four failures deserve explicit tests.

1. *Rank saturation.* For z near 37, double precision rounds $\text{expit } z$ to one. Any formula that subsequently forms $1 - u$ loses the right-tail probability. Use $\text{expit}(-z)$ directly and keep the cash leg in log space.
2. *Interior overflow.* Admissible endpoints do not bound a polynomial between them. Large Legendre coefficients can make g exceed the exponent range in the body while $g(0)$ and $g(1)$ remain moderate. The build must inspect the full grid and reject the trial before exponentiation.
3. *Near-wall cancellation.* Normalization and right-tail prices contain $1/(1 - \lambda_+)$. The logistic chart slows parameter steps near the wall, while the hard guard leaves a finite numerical margin.
4. *Lee cancellation.* In the textbook expression for $\Psi(r)$, $\sqrt{r^2 + r} - r$ subtracts nearly equal quantities for large r . Use the rationalized form (2.55), with the limiting value zero handled explicitly as $r \rightarrow \infty$.

Cached, parameter-independent grid arrays should be read-only. An accidental in-place modification of ranks, logistic weights, or basis values otherwise corrupts every later slice without necessarily producing an exception.

2.B.6 Validation stack

The solvable slice tests the most code with the least ambiguity. Across $s \in [0.05, 0.95]$, compare the computed transport with $m + sz$, the shift with $-\log(\pi s / \sin \pi s)$, and the martingale integral with one. The observed worst errors are 1.1×10^{-11} , 3.0×10^{-9} , and 3.3×10^{-16} . The shift error grows toward the wall because the normalizing integrand decays at rate $1 - s$.

The randomized battery serves a different purpose. It samples several basis orders, ordinary slices, near-wall slices, and slices with large interior modes. Strikes are placed between grid nodes, and convexity is checked in normalized strike rather than log-strike. A four-times-finer independent build provides a reference price. Together with the interpolation and belly certificates, this separates algebraic validity, interpolation validity, and grid accuracy.

2.B.7 Reference defaults and timing

The exact timings depend on hardware and caches. The portable facts are the $O(Pn_{\text{grid}})$ analytic complexity and the removal of repeated builds. The reported medians and interquartile ranges come from a single run that alternates analytic and finite-difference solves after warm-up. Alternation reduces drift from cache state and background load; medians avoid a best-of- n claim. A timing table should always travel with its grid size, quote count, order, solver tolerances, and machine context.

Quantity	Default	Quantity	Default
Basis order N	16 (cap 24)	Order guard	(2.50)
Grid	$[-40, 40]$	Nodes	8001 (fit: 2001)
Wall guard	$\lambda_+ \leq 1 - 10^{-6}$	Exponent budget	700
Ridge (α, ν)	$(10^{-6}, 1), n \geq 4$	Vega floor η	10^{-4}
Tail barrier	center 0.90, slope 50	Haircut h	0.005 vol
Solver	trust-region LSQ	Tolerances	10^{-10}
Chart	logistic	Calendar nodes	$\max(49, 4N + 1)$
Belly check	801 points	Cold start	$\sqrt{3w_0}/\pi$

Table 2.3: Reference implementation defaults.

Configuration	Order	Median	IQR	vs. finite difference
Analytic Jacobian	6	15.9 ms	0.5 ms	$1.41 \times$
Finite difference	6	22.5 ms	2.7 ms	—
Analytic Jacobian	12	26.0 ms	2.1 ms	$1.68 \times$
Finite difference	12	43.8 ms	3.2 ms	—
Analytic Jacobian	16	28.9 ms	0.5 ms	$1.76 \times$
Finite difference	16	50.9 ms	3.9 ms	—

Table 2.4: Single-slice calibration times from one interleaved run. The last column is the median finite-difference time divided by the median analytic time for the same configuration.

For a new implementation, the shortest reliable porting sequence is:

1. reproduce the constant-speed normalizer and transport;
2. implement the stable cash leg and tail corrections;
3. keep (x, x', G, G') as one immutable slice object;
4. add the envelope sensitivity pass;
5. run monotonicity, Durrleman, and randomized sub-grid checks on every release.

2.C Data and reproducibility

Every empirical number and figure is generated from one frozen JSON snapshot, captured after the 2026-08-03 US session from a delayed consolidated US listed-options feed. The file contains SPY and NVDA chains for eight expiries each, with spots 757.67 and 206.80, together with forwards, discount factors, prepared quotes, fitted parameters, and display curves. Maturities run from 2 calendar days to 1.3726 years.

The snapshot is self-contained. For each node it records the normalized chain, forward and discount provenance, retained quote set, fitted raw parameters, and diagnostic curves. Reproduction therefore requires neither credentials nor a market-data connection. The capture timestamp, reference date, and session are part of the manifest, so a later book cannot silently replace the one used for the chapter.

The manifest records the common fitting recipe: order 16 under (2.50), logistic chart, haircut $h = 0.005$, equal quote weights, ridge $\alpha = 10^{-6}$ with $\nu = 1$, and calendar rows active. All 16 nodes were fitted; none was tuned individually. The shortest NVDA node has the largest residual, consistent with its sparse and steep quote strip. SPY bands collapse to mids under the haircut rule; NVDA retains interval targets. These states are stored per quote rather than reconstructed from rounded display values.

The figure generator reads only this snapshot and fixed synthetic specifications. Real nodes are rebuilt from the frozen coefficients rather than refitted, and the rebuilt curves agree with the

stored curves to 0.000 vol bp in the worst case. The script emits both the 14 figure PDFs and the LaTeX macro file supplying every empirical value in the chapter. Randomized audits use fixed seeds. Timing values follow a fixed interleaved protocol but vary with the machine.

The synthetic double-hump target is a fixed two-component lognormal mixture, martingale-normalized and sampled on a realistic strike ladder. The constant-speed benchmark requires no data and is evaluated directly from (2.18). Regeneration consists of running the figure generator and then compiling the present manuscript. Refetching the market feed would create a different dataset and is not part of reproduction.

One family-level diagnostic is retained with the data. The two shortest SPY expiries cross in implied-volatility space only far outside their common quote span, at call prices near numerical zero. Their direct normalized call prices remain ordered to 2.2×10^{-15} on the audited grid. Keeping both measurements makes the support-confinement example of section 2.8 reproducible without interpreting the picture by eye.

Chapter 3

Families in the Volatility Chart: SVI-JW and Superposition

Chapter 2 built the law first and read the smile off it. This chapter works the other way, directly in the volatility chart, as the industry does. The gain is legibility: SVI draws one expiry as a tilted hyperbola whose jump-wing reading is the desk's vocabulary, and the Multi-Core Sigmoid adds local body detail by superposing kernels that are provably silent in both tails. The price is that validity must now be policed. The wing slopes are a moment reading capped strictly inside Lee's bound — the bound itself admits negative tail density — while the interior escapes every wing condition, as a classical counterexample shows, and is decided by a dense-grid check instead. The jump-wing coordinates invert exactly on a proved domain and lose exactly the interior curvature on its singular stratum. A structural optimization chart makes the enforced inequalities hold at every iterate by construction. The chapter closes with the three families of this book fitted to the same frozen quotes under one protocol: fit quality, validity margin, and tail behavior, with no family winning every column.

3.1 Working in the volatility chart

Chapter 2 built a smile by choosing a probability law and reading prices off it. Validity was settled before calibration began: every admissible coordinate vector was a law, and the Durrleman function g_D of (2.5) was an identity rather than a constraint. The cost was indirection. The coordinates (L, R, a_n) describe a transport speed in rank, and none of them is a quantity a trading desk quotes.

This chapter takes the opposite route, which is also the industry's. Write the total implied variance $w(k)$ down directly, as a formula in log-moneyness, with parameters chosen for legibility. The gain is immediate: the curve is the object everyone looks at, evaluation is one line, and the parameters can be tied to the vocabulary of the desk — ATM level, skew, wing weights. The cost is exactly the one section 2.1 identified when it introduced g_D : a direct volatility model moves through a space that contains invalid curves, so validity must now be policed — partly by inequalities the optimizer respects, partly by checks it cannot avoid. The purpose of this chapter is to develop that trade honestly on the two families of the reference implementation that live in the volatility chart: SVI with its jump-wing reading, and the Multi-Core Sigmoid built by superposition.

The chart serves two audiences with opposite tastes. An optimizer wants a formula that evaluates in one pass, differentiates in closed form, and behaves sensibly outside the quoted strikes. A desk asks where ATM volatility sits, how steeply the smile leaves the money, how heavy each wing is, and where the curve bottoms. The stochastic-volatility-inspired (SVI) family

of Gatheral [11] answers the first request with a tilted hyperbola,

$$w(k) = a_S + b_S \left\{ \rho_S(k - m_S) + \sqrt{(k - m_S)^2 + s_S^2} \right\}, \quad (3.1)$$

five parameters under the structural conditions $b_S > 0$, $|\rho_S| < 1$, $s_S > 0$. The subscript S keeps the book's ledger unambiguous: bare ρ is the logistic density of section 2.2 and bare m its martingale shift. The desk's request is answered by a second coordinate system on the same curve, the jump-wing handles of section 3.5. The two are exact re-parameterizations of one another on a domain that theorem 3.5 makes precise.

What “valid” means must now be graded, because different conditions are supplied by different mechanisms. Five graded statements recur throughout the chapter, collected once in table 3.1. The first is guaranteed by the optimization chart; the next two are inequalities cheap enough to enforce during calibration; the fourth is the full butterfly condition $g_D \geq 0$ of (2.5), which no inexpensive screen implies — section 3.4 exhibits the classical slice that passes every screen and fails it; the fifth couples expiries and is the subject of section 2.8, whose hinge applies to these families unchanged.

Condition	Mathematical statement	Supplied by
Structural	$b_S > 0, \rho_S < 1, s_S > 0$	the optimization chart (section 3.6)
Positive floor	$\min_k w(k) \geq 0$	a hinge penalty; structural under the structural chart
Wing-admissible	$\max(\beta_L, \beta_R) \leq \beta_{\max} < 2$	a hinge penalty; structural under the structural chart
Butterfly-free	$g_D(k) \geq 0$ for all k	<i>not</i> implied by the rows above; checked on a dense grid (section 3.4)
Calendar-ordered	$w_{\tau_1}(k) \leq w_{\tau_2}(k), \tau_1 < \tau_2$	the coupling of section 2.8, applied model-agnostically

Table 3.1: Five grades of validity for a slice written directly in the volatility chart. In Chapter 2 the first four held by construction; here each must be attributed to the mechanism that supplies it, and the fourth is the one no cheap mechanism reaches.

The chapter runs as follows. Section 3.2 establishes the geometry of the hyperbola (3.1) — including the rigidity that will later motivate superposition. Section 3.3 reads the wings through Lee's moment bound and locates that bound inside g_D ; the boundary case produces the first failure. Section 3.4 shows that the interior escapes every wing condition. Section 3.5 develops the jump-wing coordinates and proves exactly where they invert. Section 3.6 turns to calibration: the measure the objective integrates against, and the chart that makes the enforced inequalities structural. Section 3.7 constructs the Multi-Core Sigmoid by superposition of a convex base and tail-silent kernels. Section 3.8 then puts LQD, SVI, and MCS on the same frozen quotes under one protocol and lets each family's contract show.

3.2 One hyperbola

Everything in the next three sections rests on closed-form facts about (3.1). Write

$$x_S = k - m_S, \quad r_S = \sqrt{x_S^2 + s_S^2}, \quad q_S = \sqrt{1 - \rho_S^2} \quad (3.2)$$

(the subscripts keep the ledger's x and the moment order r of (2.23) untouched).

Proposition 3.1 (Geometry of one slice). *Under $b_S > 0$, $|\rho_S| < 1$, $s_S > 0$, the slice (3.1) is strictly convex with a single minimum:*

$$w'(k) = b_S \left(\rho_S + \frac{x_S}{r_S} \right), \quad w''(k) = \frac{b_S s_S^2}{r_S^3} > 0, \quad (3.3)$$

$$k^* = m_S - \frac{s_S \rho_S}{q_S}, \quad w^* = w(k^*) = a_S + b_S s_S q_S. \quad (3.4)$$

Its far field is two straight rays:

$$w(k) = \beta_L |k| + O(1) \quad (k \rightarrow -\infty), \quad w(k) = \beta_R k + O(1) \quad (k \rightarrow +\infty), \quad (3.5)$$

with

$$\beta_L = b_S(1 - \rho_S), \quad \beta_R = b_S(1 + \rho_S).$$

The curvature at the vertex is $\kappa^* = w''(k^*) = b_S q_S^3 / s_S$.

Proof. Differentiating (3.1) gives (3.3); $b_S, s_S > 0$ make $w'' > 0$ everywhere, so any stationary point is the unique global minimum. Setting $w' = 0$ gives $x_S / r_S = -\rho_S$, hence $x_S^2(1 - \rho_S^2) = \rho_S^2 s_S^2$ with x_S of the opposite sign to ρ_S , i.e. $x_S = -s_S \rho_S / q_S$ and $r_S = s_S / q_S$; substituting yields (3.4) and $\kappa^* = b_S s_S^2 (q_S / s_S)^3$. In either wing $r_S = |x_S| + O(1/|x_S|)$, which gives (3.5). $\square$

The formula reads as a rounded hinge. Far from the vertex the hinge has forgotten its rounding and shows two straight rays, of slope β_L on the left and β_R on the right; the tilt ρ_S divides a fixed total steepness $\beta_L + \beta_R = 2b_S$ between the two sides, with $\rho_S < 0$ — the equity put skew — making the left ray steeper. Near the vertex, s_S decides how sharply one ray turns into the other, m_S places the turn, and a_S raises the whole curve. Figure 3.1 sweeps each role in isolation.




Figure 3.1: The reach of the family (3.1), one parameter at a time around a common slice. (a) The tilt ρ_S redistributes a fixed total steepness between the wings. (b) The width s_S rounds the vertex without moving the rays. (c) The steepness b_S opens or closes both wings together. Every member is one convex curve with one minimum.

Two consequences of proposition 3.1 organize the rest of the chapter.

First, the tails collapse to two numbers. As $|k|$ grows, $w(k)/|k|$ converges to β_L on the put side and β_R on the call side, and no other feature of the slice survives to infinity. Whatever the market's tails are worth, this family records them in (β_L, β_R) alone — which is precisely the pair that Lee's moment bound constrains, and section 3.3 develops that correspondence.

Second, strict convexity is a rigidity. One hyperbola has one belly and one monotone turn of the slope, from $-\beta_L$ up to $+\beta_R$, and this excludes a shape that markets do produce.

Proposition 3.2 (A convex slice cannot carry a W). *A convex function on an interval has no interior strict local maximum. Hence a total-variance curve with a central trough flanked by two shoulders — an interior local maximum on each side — is not representable by any globally convex slice, in particular by (3.1).*

Proof. Convexity makes the one-sided derivatives non-decreasing; a strict interior maximum would force the derivative to pass from positive to negative, contradicting monotonicity. By proposition 3.1, $w'' > 0$ everywhere for the family (3.1). $\square$

On a liquid index chain this rigidity is a virtue: it forbids the wiggles that quote noise would otherwise carve into the curve. Ahead of a scheduled event, or in a name with bimodal positioning, it is a ceiling — the smile genuinely turns more than once, and the family cannot follow. Section 3.7 removes the ceiling without giving up the wings.

3.3 Wings are tails: Lee's bound inside the density factor

Why should a wing slope be capped at all? Because it is a moment barometer. Lee's formula (2.23) ties the asymptotic slope of total variance to the critical moment order of the gross return: on the right wing $\beta_R = \Psi(r_+^*)$, where $1 + r_+^*$ is the largest finite moment of Y and $\Psi(r) = 2 - 4(\sqrt{r^2 + r} - r)$. As the law loses its moments, $r_+^* \downarrow 0$ and $\Psi \uparrow 2$: a wing steeper than 2 is the tail of no probability law at all. For LQD the correspondence was a closed formula in the endpoint speeds (proposition 2.9); for a family whose wings are exactly linear it is a sharp, checkable inequality on two parameters, $\max(\beta_L, \beta_R) \leq 2$.

The cleanest way to watch the bound bite is to follow it into the function that tests for arbitrage. The Durrleman factor (2.5) mixes w , w' , and w'' ; in the wings of the family (3.1) the mixture simplifies to a limit that depends on the slope alone.

Proposition 3.3 (The wing slope sets the sign of the tail density). *For a slice (3.1) with $w > 0$,*

$$\lim_{k \rightarrow -\infty} g_D(k) = \frac{4 - \beta_L^2}{16}, \quad \lim_{k \rightarrow +\infty} g_D(k) = \frac{4 - \beta_R^2}{16}. \quad (3.6)$$

Hence for $\beta < 2$ the corresponding tail of g_D is eventually positive and for $\beta > 2$ eventually negative. At $\beta_R = 2$ the limit vanishes and the sign falls to the next order: the right wing is $w(k) = 2k + (a_S - 2m_S) + O(k^{-1})$, and

$$g_D(k) = \frac{(a_S - 2m_S) - 2}{4k} + O(k^{-2}), \quad (3.7)$$

so the boundary slice carries negative tail density whenever $a_S - 2m_S < 2$. The mirror statement holds on the left with the intercept $a_S + 2m_S$.

Proof. In either wing $w \sim \beta|k|$, $|w'| \rightarrow \beta$, and $w'' \rightarrow 0$ by (3.3). The first bracket of (2.5) tends to $1 - \frac{1}{2} = \frac{1}{2}$, so its square tends to $\frac{1}{4}$; the middle term tends to $\beta^2/16$ because $1/w \rightarrow 0$; the last term vanishes. This gives (3.6). For the boundary order, expand (3.1) on the right with $b_S(1 + \rho_S) = 2$:

$$w(k) = 2k + (a_S - 2m_S) + \frac{b_S s_S^2}{2k} + O(k^{-2}), \quad w'(k) = 2 + O(k^{-2}).$$

Then $kw'/(2w) = \frac{1}{2}\{1 - (a_S - 2m_S)/(2k)\} + O(k^{-2})$, so the first bracket squares to $\frac{1}{4} + (a_S - 2m_S)/(4k) + O(k^{-2})$; the middle term is $\frac{1}{4}w'^2\{1/w + \frac{1}{4}\} = \frac{1}{4} + 1/(2k) + O(k^{-2})$ because $w'^2 \rightarrow 4$ and $1/w = 1/(2k) + O(k^{-2})$; and $w'' = O(k^{-3})$. Subtracting gives (3.7). $\square$

The proposition says two distinct things, and the second is easy to miss. Below the bound, the tail of the implied density is eventually positive — the wing condition genuinely protects the far tail. At the bound, it protects nothing: the limit is zero, and the $1/k$ term decides. A concrete slice makes the boundary failure explicit. Take $(a_S, b_S, \rho_S, m_S, s_S) = (0.04, 2, 0, 0, 0.2)$: its minimum total variance is 0.44, comfortably positive, and both wings sit exactly at $\beta = 2$, so any constraint of the form $\beta \leq 2$ is satisfied with equality and charges nothing. Yet $a_S - 2m_S = 0.04 < 2$, and by (3.7) the density is negative throughout the far tail — $g_D(10) = -0.0485$ on this slice. A constraint written with Lee's constant as its right-hand side therefore admits laws that are not laws.

The reference implementation draws the consequence as a convention, stated here once: the enforced cap is strictly inside the bound,

$$\max(\beta_L, \beta_R) \leq \beta_{\max}, \quad \beta_{\max} = 1.95 < 2. \quad (3.8)$$

What does the buffer exclude? Inverting Ψ gives the *moment budget* at slope β ,

$$r^* = \Psi^{-1}(\beta) = \frac{(2 - \beta)^2}{8\beta}, \quad (3.9)$$

the number of moments beyond the first that a wing of slope β still permits. At the cap, $r^* = 1.6 \times 10^{-4}$: the laws excluded by the buffer are those whose entire moment budget beyond the mean is smaller than two parts in ten thousand. The cap costs essentially nothing in distributional scope and restores the implication “wing-admissible $\Rightarrow$ tail density eventually positive.”

The map (3.9) also says when a fitted wing may be read as a moment estimate, because its conditioning is extreme at both ends. Differentiating, $dr^*/d\beta = -(4 - \beta^2)/(8\beta^2)$, which behaves as $-1/(2\beta^2)$ for small β : a shallow wing pins essentially no upper bound on the moments, and small changes in the fitted slope move the inferred moment violently. As $\beta \rightarrow 2$ the budget collapses to zero. A fitted β is a candidate moment reading only in the middle of the range — and only for a slice that is actually a law, which is the subject of the next section. Figure 3.2 draws all three parts of the story.




Figure 3.2: Lee's bound inside the density factor. (a) The tail limit (3.6) is a downward parabola crossing zero exactly at $\beta = 2$; the dashed line is the working cap $\beta_{\max} = 1.95$. (b) The boundary slice of the text: at $\beta_R = 2$ (orange) the tail density is negative at every plotted strike, approaching zero from below as (3.7) prescribes; the same slice with its wings at the cap (blue) turns positive and settles at the limit $(4 - \beta_{\max}^2)/16$ (dotted). (c) The moment budget (3.9) on a log scale: it explodes as $\beta \rightarrow 0$ and collapses at the cap (marked point, $r^* = 1.6 \times 10^{-4}$).

3.4 The belly escapes

The word “convex” hides two different second derivatives. The family (3.1) has $w'' > 0$ by construction; a butterfly spread tests the convexity of the option *price* in strike, which is the sign of g_D , and g_D mixes w , w' , and w'' . In the wings the two notions agree — proposition 3.3 showed the sign of g_D there is Lee’s — but in the interior they need not. The floor condition of table 3.1 fences the vertex, the cap (3.8) fences the wings, and between them lies a region neither condition reaches.

That the gap is real, and not a defect of any particular fitting code, is a classical example due to Vogt, reproduced as Example 3.1 by Gatheral and Jacquier [13]:

$$(a_S, b_S, \rho_S, m_S, s_S) = (-0.0410, 0.1331, 0.3060, 0.3586, 0.4153).$$

Its minimum total variance is $w^* = 0.0116$, so the floor is clean. Its larger wing slope is 0.174, an order of magnitude under the cap, and by (3.6) both tails of g_D settle near 0.25 — the slice is impeccable everywhere the two inequalities look. Yet g_D dips to -0.033 near $k = 0.88$, squarely between the two clean tails: the implied density is negative there, and the slice prices a butterfly below zero. Figure 3.3 shows both halves. Strict convexity of w did not protect strike-convexity of the price.




Figure 3.3: The Vogt slice. (a) Total variance: the floor is positive and both wing slopes are far below the cap — every inequality of table 3.1 up to wing-admissibility holds strictly. (b) Its Durrleman factor approaches the positive tail limits (3.6) on both sides yet turns negative in the interior. The wing conditions fence the tails; the belly escapes them.

Could the gap be closed by a better inequality? In principle yes: the exact butterfly-free domain of (3.1) has been characterized completely by Martini and Mingone [20]. But the characterization is not one short inequality — it involves a parameter rescaling, root finding, and a numerical minimization — and it is specific to this five-parameter family, with no analogue for the superposition family of section 3.7. The reference implementation therefore adopts the same division of labor for every family in the volatility chart, and it is worth stating as the chapter’s second convention. During optimization, the floor and cap enter as two hinge rows appended to the residual,

$$W \max(-(a_S + b_S s_S q_S), 0), \quad W \max(b_S(1 + |\rho_S|) - \beta_{\max}, 0), \quad (3.10)$$

with a multiplier W large enough to dominate a violation; both rows are exactly zero on a slice satisfying the two inequalities, so a clean fit is identical with or without them. After optimization, the full condition is checked where the inexpensive rows cannot reach: g_D is evaluated from the

model's own derivatives — never by differencing prices — on a dense grid across the traded strike range, and the slice is accepted only if the sampled minimum clears a small numerical tolerance. This is the same posture as the numerical certificates of section 2.7.2: an analytic guarantee where one exists, and where it does not, a computed check on the finished object, with its scope — the sampled grid on the traded range, not the whole line — stated rather than implied. The minimum of g_D over that grid is the *validity margin* that the comparison of section 3.8 reports for every fit of every family.

The division matters in both directions. The condition $g_D \geq 0$ *can* be pushed into the optimizer as one more soft row — a hinge on sampled values of g_D over a grid — and the superposition family of section 3.7, whose kernels can bend the density where a hyperbola cannot, carries exactly such a row (appendix 3.A). But a finite-weight hinge is a fence, not a proof: it is traded against the data term, it vanishes only where the sampled values already satisfy the inequality, and section 3.8 shows fits that fail the final check with the fence active. Conversely, trusting any of the fences alone would publish the Vogt slice. Fence cheaply, then check completely: the fences keep the optimizer in the region where the check usually passes, and the check decides.

3.5 Trader coordinates: the jump-wing handles

The raw quintuple $(a_S, b_S, \rho_S, m_S, s_S)$ is an excellent computational coordinate system and a poor market language: no single coordinate is “the level” or “the skew,” and each of b_S , ρ_S , s_S moves a wing and the interior at once. The jump-wing (JW) coordinates of Gatheral and Jacquier [13] re-read the same curve through five functionals chosen to match the desk's questions, and — the reading this section presses — they sort into *three interior handles and two tail handles*.

Definition 3.4 (Jump-wing handles). At variance time τ , with $w_0 = w(0)$ the ATM total variance, the jump-wing coordinates of a slice are

$$\underbrace{v_J = \frac{w_0}{\tau}, \quad \psi_J = (\sqrt{w})'(0), \quad \tilde{v}_J = \frac{\min_k w(k)}{\tau}}_{\text{interior}}, \quad \underbrace{p_J = \frac{\beta_L}{\sqrt{w_0}}, \quad c_J = \frac{\beta_R}{\sqrt{w_0}}}_{\text{tails}}. \quad (3.11)$$

Substituting (3.1) gives the closed forms

$$v_J = \frac{w_0}{\tau}, \quad \psi_J = \frac{b_S}{2\sqrt{w_0}} \left(\rho_S - \frac{m_S}{\sqrt{m_S^2 + s_S^2}} \right), \quad p_J = \frac{b_S(1 - \rho_S)}{\sqrt{w_0}}, \quad c_J = \frac{b_S(1 + \rho_S)}{\sqrt{w_0}}, \quad \tilde{v}_J = \frac{w^*}{\tau}, \quad (3.12)$$

with w^* from (3.4). Each handle lives in its own units, and misreading them is the most common JW error. With $I(k) = \sqrt{w(k)}/\tau$ the working implied volatility,

$$I(0) = \sqrt{v_J}, \quad I'(0) = \frac{\psi_J}{\sqrt{\tau}}, \quad \min_k I(k) = \sqrt{\tilde{v}_J}, \quad \beta_L = p_J \sqrt{v_J \tau}, \quad \beta_R = c_J \sqrt{v_J \tau}. \quad (3.13)$$

So v_J and $\tilde{v}_J$ are variances — the IV chart shows their square roots; ψ_J is the ATM slope of total *volatility* $\sqrt{w}$, and the plotted tangent is $\psi_J/\sqrt{\tau}$; and p_J, c_J are *normalized* total-variance wing slopes, not slopes of the IV curve. Figure 3.4 annotates a fitted slice — SPY December 2026 from the frozen snapshot, the book's recurring node — in exactly these units: the fit carries $\sqrt{v_J} = 14.89\%$, tangent slope -0.430 , floor $\sqrt{\tilde{v}_J} = 12.41\%$, and wing handles $(p_J, c_J) = (1.384, 0.724)$.

One caution before inverting. The tail handles are *normalized* indicators: by (3.13) the actual asymptotic slopes — the objects Lee's bound and the moment map (3.9) read — carry a




Figure 3.4: The five handles in their honest units, on the SVI fit of the SPY December 2026 quotes (rms 7.5 vol bp). (a) The three interior handles are read on the IV chart: the level $\sqrt{v_J}$, the minimum $\sqrt{\tilde{v}_J}$, and the ATM tangent of slope $\psi_J/\sqrt{\tau}$. (b) The two tail handles are read on the total-variance chart, as the slopes $\beta_L = p_J\sqrt{v_J\tau}$ and $\beta_R = c_J\sqrt{v_J\tau}$ of the asymptotes. Plotting all five on one IV axis would give the wing handles units they do not have.

factor $\sqrt{v_J\tau}$ of the supposedly interior level. Raising v_J with (p_J, c_J) held fixed steepens both actual wings and moves the inferred moments. JW is a quoting convention adapted to the desk, not an orthogonal decomposition of the slice into independent interior and tail degrees of freedom.

3.5.1 Inverting the handles

Reading (3.12) off a fitted slice is evaluation. The interesting direction is the reverse: reconstructing the hyperbola from five quoted handles. The two tail handles fix the scale and the tilt,

$$b_S = \frac{\sqrt{w_0}}{2}(p_J + c_J), \quad \rho_S = \frac{c_J - p_J}{c_J + p_J}, \quad w_0 = v_J\tau; \quad (3.14)$$

the ATM slope fixes the normalized vertex displacement

$$\chi := \frac{m_S}{\sqrt{m_S^2 + s_S^2}} = \rho_S - \frac{4\psi_J}{p_J + c_J}, \quad (3.15)$$

and the gap between level and floor then sets the width. Evaluating (3.1) at $k = 0$ and at the vertex gives the two identities

$$w_0 = a_S + b_S(\sqrt{m_S^2 + s_S^2} - \rho_S m_S), \quad w^* = a_S + b_S s_S \sqrt{1 - \rho_S^2}; \quad (3.16)$$

since $m_S = \chi s_S / \sqrt{1 - \chi^2}$ makes $\sqrt{m_S^2 + s_S^2} = s_S / \sqrt{1 - \chi^2}$, subtracting the second from the first leaves $w_0 - \tilde{v}_J\tau = b_S s_S \mathcal{D}$ with

$$\mathcal{D}(\rho_S, \chi) = \frac{1 - \rho_S \chi}{\sqrt{1 - \chi^2}} - \sqrt{1 - \rho_S^2}, \quad (3.17)$$

so that

$$s_S = \frac{(v_J - \tilde{v}_J)\tau}{b_S \mathcal{D}}, \quad m_S = \frac{\chi s_S}{\sqrt{1 - \chi^2}}, \quad a_S = \tilde{v}_J\tau - b_S s_S \sqrt{1 - \rho_S^2}. \quad (3.18)$$

The denominator's sign is Cauchy–Schwarz. Applied to the unit vectors $(\rho_S, \sqrt{1 - \rho_S^2})$ and $(\chi, \sqrt{1 - \chi^2})$,

$$\rho_S \chi + \sqrt{1 - \rho_S^2} \sqrt{1 - \chi^2} \leq 1, \quad \text{i.e. } \mathcal{D} \geq 0, \quad (3.19)$$

with equality exactly when the vectors coincide, $\chi = \rho_S$. This single inequality gives both the domain and the singularity of the coordinate system.

Theorem 3.5 (The image of the jump-wing map). *Fix $\tau > 0$ and restrict (3.1) to $b_S > 0$, $|\rho_S| < 1$, $s_S > 0$, $w_0 > 0$.*

(i) *A JW point has a unique inverse, given by (3.14)–(3.18), exactly on*

$$v_J > 0, \quad p_J > 0, \quad c_J > 0, \quad -\frac{p_J}{2} < \psi_J < \frac{c_J}{2}, \quad \psi_J \neq 0, \quad \tilde{v}_J < v_J. \quad (3.20)$$

(ii) *The image also contains the stratum $\psi_J = 0$, $\tilde{v}_J = v_J$ (with $v_J, p_J, c_J > 0$), every point of which is attained by infinitely many slices. There are no other image points.*

Proof. If $p_J, c_J > 0$ then (3.14) gives $b_S > 0$ and $\rho_S \in (-1, 1)$; conversely every admissible slice has $p_J, c_J > 0$. From (3.15), using $(\rho_S - 1)(p_J + c_J) = -2p_J$ and $(\rho_S + 1)(p_J + c_J) = 2c_J$,

$$|\chi| < 1 \iff \rho_S - 1 < \frac{4\psi_J}{p_J + c_J} < \rho_S + 1 \iff -\frac{p_J}{2} < \psi_J < \frac{c_J}{2},$$

which keeps m_S finite. By (3.19), $\mathcal{D} > 0$ exactly when $\chi \neq \rho_S$, i.e. $\psi_J \neq 0$; then $s_S > 0$ in (3.18) requires $v_J > \tilde{v}_J$. Each failure breaks a requirement: one of $p_J, c_J \leq 0$ forces $b_S \leq 0$ or $|\rho_S| \geq 1$; ψ_J outside the band sends $|\chi| \geq 1$; $\psi_J = 0$ annihilates $\mathcal{D}$; $\tilde{v}_J \geq v_J$ makes $s_S \leq 0$. For the stratum: $\psi_J = 0$ gives $\chi = \rho_S$, which by (3.4) puts the vertex at $k^* = 0$, so the floor is attained at the money and $\tilde{v}_J = v_J$ automatically. Conversely fix (v_J, p_J, c_J) on the stratum; (3.14) sets (b_S, ρ_S) , and for every $s_S > 0$ the slice with $m_S = \rho_S s_S / \sqrt{1 - \rho_S^2}$ and $a_S = v_J \tau - b_S s_S \sqrt{1 - \rho_S^2}$ reproduces the same five handles. $\square$

Remark 3.6 (The blind spot is the interior curvature). The free width on the stratum is not an abstract degeneracy. There the vertex sits at the money, the handles pin the level and both tails, and what they omit is exactly how sharply the curve turns through its bottom: by (3.4) the vertex curvature $\kappa^* = b_S q_S^3 / s_S$ runs through every positive value as s_S does. Figure 3.5(b) draws three slices sharing one handle quintuple. Nor is the failure confined to the stratum itself: $\mathcal{D}$ vanishes at $\chi = \rho_S$ together with its first χ -derivative, and a direct differentiation of (3.17) gives $\partial^2 \mathcal{D} / \partial \chi^2 = (1 - \rho_S^2)^{-3/2}$ there, so to second order,

$$\mathcal{D} \sim \frac{(\chi - \rho_S)^2}{2(1 - \rho_S^2)^{3/2}} = \frac{8\psi_J^2}{(p_J + c_J)^2(1 - \rho_S^2)^{3/2}}, \quad \psi_J \rightarrow 0, \quad (3.21)$$

so the denominator vanishes *quadratically*: near the stratum, recovering a finite width forces $v_J - \tilde{v}_J = O(\psi_J^2)$, quote noise in the handles is amplified by $1/\psi_J^2$, and a direct evaluation of (3.17) subtracts two nearly equal numbers. An implementation can remove the cancellation by evaluating an algebraically identical form; no rearrangement can restore a curvature the handles never carried. The stratum is not exotic — ATM sitting at the smile's minimum is where a desk would expect the coordinates to be simplest, and fitted index slices pass close to it.

The screens of section 3.4 translate directly into this vocabulary, which is convenient for checking a quoted handle set before any conversion: the floor condition is $\tilde{v}_J \geq 0$; the domain condition is $-p_J/2 < \psi_J < c_J/2$; and since $b_S(1 + |\rho_S|) = \max(\beta_L, \beta_R) = \max(p_J, c_J)\sqrt{w_0}$ by (3.14) and (3.13), the cap (3.8) reads

$$\max(p_J, c_J) \sqrt{v_J \tau} \leq \beta_{\max}. \quad (3.22)$$

None of the three certifies the interior, for the reason section 3.4 established.




Figure 3.5: Where the jump-wing chart tears. (a) The inverse denominator $\mathcal{D}$ against ψ_J at fixed (p_J, c_J) : it closes quadratically at $\psi_J = 0$, following (3.21) (dashed). (b) Three slices with identical handles $(v_J, 0, p_J, c_J, v_J)$ — same level, same asymptotes — differing only in width s_S : the five handles lose the interior curvature (handle agreement across the three slices: 1.1×10^{-17}).

3.6 Calibration in the chart

Chapter 2 settled the anatomy of a slice objective — units, band targets, regularization (section 2.6). Working in the volatility chart changes two of the choices, and this section treats them in turn. The residual can now live directly in implied volatility, since the model produces $w(k)$ without an inversion: with quotes (k_i, σ_i) and weights μ_i , the data term is

$$\sum_i \mu_i (I(k_i; \theta) - \sigma_i)^2, \quad I(k; \theta) = \sqrt{w(k; \theta)/\tau}, \quad (3.23)$$

optionally replaced by the band or haircut targets of (2.48) and (2.49) evaluated in volatility space. And the chart itself becomes a decision, because the coordinates no longer guarantee validity for free.

3.6.1 The measure a fit integrates against

Idealize (3.23) as a continuum problem: minimize $\int (I - \sigma)^2 d\mu(k)$ for some measure μ on the strike axis. A finite quote set turns the integral into a sum, and the weights are the quadrature rule. Equal weights make μ the empirical measure of the *listing* — the exchange's strike schedule, dense at the money and sparse in the wings. That measure was chosen by a listings committee: a tight cluster of near-duplicate ATM quotes outvotes an isolated wing quote by sheer count, and the fit inherits the strike schedule as an accidental prior. Adding more ATM listings, which adds information, then makes the fitted wings *worse*.

A declared alternative weights strike ranges by the time value they carry, $d\mu = \text{TV}(k) dk$, where $\text{TV}(k)$ is the normalized price of the out-of-the-money option at k — the part of the premium that is genuinely optionality rather than intrinsic value, the natural gauge of how much economic content a strike carries. Let C_i be quote i 's cell in the Voronoi partition of the quoted interval — half the gap to each neighbor, one-sided at the ends — with width $|C_i|$. The scheme

$$\mu_i \propto \text{TV}_i |C_i| \quad (3.24)$$

is then a genuine discretization of the declared measure.

Proposition 3.7 (Weights as a quadrature). *For any continuous integrand f , the sum $\sum_i \text{TV}_i |C_i| f(k_i)$ is a Riemann sum of $\int \text{TV}(k) f(k) dk$ over the Voronoi partition, converging under refinement. In particular the weighted loss with (3.24) discretizes the continuum time-value objective; on a uniform strike grid all $|C_i|$ coincide and the scheme reduces to pure time-value weighting.*

Proof. The cells partition the quoted interval and $k_i \in C_i$, so the sum is a Riemann sum over that partition; refinement convergence is the standard statement. The uniform case is arithmetic. $\square$

Two conventions keep the quadrature usable on real ladders, both stated once. The cell of an isolated far-wing quote can be arbitrarily wide, so the width factor is capped (the reference implementation uses ten times the mean spacing): beyond the cap the scheme deliberately under-integrates the deep wing rather than let one quote carry a tenth of the objective. And the weights are normalized to unit mean, so switching schemes redistributes the data term without rescaling it against the regularization. The chapter’s comparisons use equal weights throughout — the point here is not that one measure is right, but that a weight vector *is* a measure, and should be chosen as one.

3.6.2 Two charts

The optimizer never handles a constrained vector. In the historical chart, $\theta \in \mathbb{R}^5$ maps to

$$a_S = \theta_1, \quad b_S = \text{softplus}(\theta_2), \quad \rho_S = \tanh(\theta_3), \quad m_S = \theta_4, \quad s_S = e^{\theta_5}, \quad (3.25)$$

with $\text{softplus}(\theta) = \log(1 + e^\theta)$, so every finite θ is a structurally valid hyperbola — the first row of table 3.1 and nothing more. The floor can still go negative and a wing can still exceed the cap; those are fenced by the hinge rows (3.10), which the optimizer can trade against the data term. This is the same constraint-elimination idea as the logistic chart of section 2.6.3, applied to the weaker guarantee.

The idea extends to the conditions the hinges police. Parameterize the slice by the quantities the inequalities are *about* — the two wing slopes, the vertex location, the floor, and the vertex curvature — each lifted so that every finite coordinate is admissible:

$$\beta_L = \beta_{\max} \Lambda(\theta_1), \quad \beta_R = \beta_{\max} \Lambda(\theta_2), \quad k^* = \theta_3, \quad w^* = \text{softplus}(\theta_4), \quad \kappa^* = e^{\theta_5}, \quad (3.26)$$

with Λ the logistic function of section 2.1. The raw parameters are recovered exactly by inverting proposition 3.1:

$$b_S = \frac{\beta_L + \beta_R}{2}, \quad \rho_S = \frac{\beta_R - \beta_L}{\beta_R + \beta_L}, \quad s_S = \frac{b_S(1 - \rho_S^2)^{3/2}}{\kappa^*}, \quad m_S = k^* + \frac{s_S \rho_S}{\sqrt{1 - \rho_S^2}}, \quad a_S = w^* - b_S s_S \sqrt{1 - \rho_S^2}. \quad (3.27)$$

Under this *structural chart* every finite iterate has $w(k) \geq w^* > 0$ and both wings strictly inside $(0, \beta_{\max})$: the second and third rows of table 3.1 hold by construction, and the hinge rows (3.10) are identically zero along the entire optimization path. What the chart does *not* deliver is the fourth row — the Vogt slice of section 3.4 is reachable in either chart, and the dense-grid check remains the deciding instance. The two charts describe the same family, and on the SPY December quotes they reach the same slice to 0.00 vol bp (fig. 3.6); what changes is the geometry the solver works in, with the kinks of the hinge rows removed exactly where a trust-region step would otherwise stall on them.

3.6.3 The residual Jacobian, in layers

The least-squares solvers of both charts consume an analytic Jacobian, and it assembles by the chain rule in three layers. The raw-parameter partials of the hyperbola follow from (3.1) with




Figure 3.6: The structural chart. (a) The wing lift $\beta = \beta_{\max}\Lambda(\theta)$: every finite coordinate lands strictly between 0 and $\beta_{\max}$, itself strictly below Lee’s bound — admissibility is a property of the chart, not a constraint on the step. (b) The five structural coordinates read on the SPY December fit: vertex $(k^*, w^*) = (0.108, 0.0058)$, osculating parabola of curvature $\kappa^* = 0.397$, and the two asymptote slopes. The optimizer moves exactly the quantities the inequalities constrain.

(3.2):

$$\frac{\partial w}{\partial a_S} = 1, \quad \frac{\partial w}{\partial b_S} = \rho_S x + r, \quad \frac{\partial w}{\partial \rho_S} = b_S x, \quad \frac{\partial w}{\partial m_S} = -b_S \left(\rho_S + \frac{x}{r} \right), \quad \frac{\partial w}{\partial s_S} = \frac{b_S s_S}{r}. \quad (3.28)$$

The data rows live in volatility, so each is scaled by

$$\frac{\partial I}{\partial w} = \frac{1}{2\sqrt{w\tau}} = \frac{1}{2\tau I}. \quad (3.29)$$

Finally the lift: for (3.25) the three nontrivial factors are

$$\frac{db_S}{d\theta_2} = 1 - e^{-b_S}, \quad \frac{d\rho_S}{d\theta_3} = 1 - \rho_S^2, \quad \frac{ds_S}{d\theta_5} = s_S, \quad (3.30)$$

while the structural chart pushes the same raw-space rows through the closed-form 5×5 matrix $\partial(a_S, b_S, \rho_S, m_S, s_S)/\partial\theta$ obtained by differentiating (3.27) along (3.26). Hinge rows contribute their linear part where the violation is strictly positive and a zero row otherwise. Nothing here is deep; the point of recording the layers is that the chart is a change of variables like any other, and its cost is one matrix product per iterate rather than a constrained solver.

3.7 Superposition: capacity without touching the wings

Proposition 3.2 left a gap: some smiles — ahead of a binary event, or in a name with bimodal positioning — carry a central trough with a shoulder on each side, and no globally convex slice can follow them. The design problem is sharper than “add flexibility.” The wings are the solved part of the curve: Lee’s bound caps them, the quoted wing strikes pin them, and section 3.3 showed the tail of the density reads them alone. What is missing is detail in the *body*, and the correction that supplies it must not move the wings at all.

The Multi-Core Sigmoid (MCS) solves this by superposition: a convex base that owns the wings, plus localized corrections that are silent — in value, slope, *and* curvature — in both tails. The construction rests on one elementary mechanism, developed first.

3.7.1 The base

Work in the standardized moneyness

$$\zeta = \frac{k}{\sigma_{\text{ref}} \sqrt{\tau}}, \quad (3.31)$$

where σ_{ref} is a reference volatility fixed before the fit (the quoted volatility nearest the money), so ζ measures moneyness in standard deviations and the parameters are roughly scale-free across expiries. Displacements along this axis are written $\bar{\zeta}$ (measured from a center fixed by context, so the ledger's rank u stays untouched). The convexity primitive is the log-cosh

$$\mathcal{A}_\kappa(\bar{\zeta}) = \frac{4}{\kappa^2} \log \cosh \frac{\kappa \bar{\zeta}}{2}, \quad \mathcal{A}'_\kappa(\bar{\zeta}) = \frac{2}{\kappa} \tanh \frac{\kappa \bar{\zeta}}{2}, \quad \mathcal{A}''_\kappa(\bar{\zeta}) = \text{sech}^2 \frac{\kappa \bar{\zeta}}{2}, \quad (3.32)$$

a smooth, even, convex softening of the absolute value. Two facts about it carry the section. Near the origin $\mathcal{A}_\kappa(\bar{\zeta}) = \bar{\zeta}^2/2 + O(\bar{\zeta}^4)$: a rounded vertex. Far out it is *asymptotically affine*,

$$\mathcal{A}_\kappa(\bar{\zeta}) = \frac{2}{\kappa} |\bar{\zeta}| - \frac{4}{\kappa^2} \log 2 + o(1), \quad (3.33)$$

since $\log \cosh x = |x| - \log 2 + \log(1 + e^{-2|x|})$. The base models annualized variance $V(\zeta) = w/\tau$ as

$$V_{\text{base}}(\zeta) = V_0 + V_1(\zeta - \zeta_0) + V_2 \mathcal{A}_{\kappa_i}(\zeta - \zeta_0), \quad \kappa_i = \begin{cases} \kappa_P, & \zeta < \zeta_0, \\ \kappa_C, & \zeta \geq \zeta_0, \end{cases} \quad (3.34)$$

six parameters: level V_0 , slope V_1 , and convexity amplitude $V_2 \geq 0$ at the center ζ_0 , with separate put and call steepnesses. Because $\mathcal{A}_\kappa$ and $\mathcal{A}'_\kappa$ vanish at the origin and $\mathcal{A}''_\kappa(0) = 1$ whatever κ , the curve is C^2 across ζ_0 even though the steepness switches there: the asymmetry lies in how fast each wing straightens, not in a kink or a curvature jump. By (3.33) the asymptotic slopes of V_{base} are $V_1 - 2V_2/\kappa_P$ on the left and $V_1 + 2V_2/\kappa_C$ on the right; converting through (3.31) and $w = \tau V$, the total-variance wing slopes in k are

$$\beta_P = \frac{\sqrt{\tau}}{\sigma_{\text{ref}}} \left(\frac{2V_2}{\kappa_P} - V_1 \right), \quad \beta_C = \frac{\sqrt{\tau}}{\sigma_{\text{ref}}} \left(V_1 + \frac{2V_2}{\kappa_C} \right). \quad (3.35)$$

These six numbers set the wings; the design goal is that nothing added later moves them.

3.7.2 A kernel that vanishes to second order

The correction kernel is the normalized centered second difference of the same primitive: with $\bar{\zeta} = \zeta - \zeta_r$,

$$\mathcal{B}_{\zeta_r, \ell_r, \kappa_r}(\zeta) = \frac{\mathcal{A}_{\kappa_r}(\bar{\zeta} - \ell_r) - 2\mathcal{A}_{\kappa_r}(\bar{\zeta}) + \mathcal{A}_{\kappa_r}(\bar{\zeta} + \ell_r)}{2\mathcal{A}_{\kappa_r}(\ell_r)}, \quad (3.36)$$

centered at ζ_r , of half-width ℓ_r and steepness κ_r , normalized so that $\mathcal{B}(\zeta_r) = 1$ (the numerator at $\bar{\zeta} = 0$ is $2\mathcal{A}_{\kappa_r}(\ell_r)$ because $\mathcal{A}$ is even with $\mathcal{A}(0) = 0$). Its derivatives are the same second difference of $\mathcal{A}'$ and $\mathcal{A}''$ over the same normalizer.

Lemma 3.8 (Zero wings). $\mathcal{B} \rightarrow 0$, $\mathcal{B}' \rightarrow 0$, and $\mathcal{B}'' \rightarrow 0$ as $\zeta \rightarrow \pm\infty$.

Proof. A centered second difference annihilates affine functions: for any affine L , $L(\bar{\zeta} - \ell) - 2L(\bar{\zeta}) + L(\bar{\zeta} + \ell) = 0$ identically. By (3.33), $\mathcal{A}_\kappa$ is affine plus $o(1)$ in each tail, so its second difference tends to zero there. The same argument applies to $\mathcal{A}'$, which tends to the constants $\pm 2/\kappa$ (affine of zero slope), and to $\mathcal{A}''$, which decays. $\square$

The lemma is the whole design, and it does not depend on the log-cosh: *any* asymptotically affine primitive yields a tail-silent kernel by the same second difference. The log-cosh is a convenience — the model and its Jacobian share one set of derivatives, the kernel has unit height so its amplitude reads directly as a variance displacement, and the half-width and steepness are separate dials where a Gaussian bump has one. At its center, $\mathcal{B}''(\zeta_r) = \{\text{sech}^2(\kappa_r \ell_r/2) - 1\}/\mathcal{A}_{\kappa_r}(\ell_r) < 0$, so a positive multiple is locally concave — a shoulder — and a negative multiple digs a notch.

3.7.3 The model, and what the superposition preserves

The MCS slice is the base plus M signed kernels,

$$V_M(\zeta) = V_{\text{base}}(\zeta) + \sum_{r=1}^M \gamma_r \mathcal{B}_{\zeta_r, \ell_r, \kappa_r}(\zeta), \quad w(k) = \tau V_M\left(\frac{k}{\sigma_{\text{ref}} \sqrt{\tau}}\right), \quad (3.37)$$

with $6 + 4M$ parameters. By lemma 3.8, every member of the family has the same asymptotic wings as its base, whatever M : expressiveness in the body and control of the tails are decoupled by construction. Figure 3.7 shows the mechanism end to end on a target built to be globally admissible — a hyperbolic total-variance base with two Gaussian shoulders, asymptotically linear wings, and $\min g_D = 0.277$ verified from the analytic derivatives on a dense grid to $|k| = 12$, with the positive tail limit (3.6) taking over beyond. The convex base alone ($M = 0$) misses the shoulders by 118.2 vol bp, and the residual it leaves is structured: two localized features, not noise. Two kernels reduce the miss to 0.7 vol bp, and the fitted slice remains butterfly-free ($\min g_D = 0.420$ on the wide diagnostic grid).




Figure 3.7: Superposition in three steps. (a) The log-cosh primitive merges with its affine asymptote (3.33); only the rounded vertex differs. (b) The centered second difference (3.36): a unit-height bump whose value, slope, and curvature all decay — the correction is silent in the wings to all three orders (lemma 3.8). (c) A W-shaped target (dots): the convex base misses both shoulders by 118.2 vol bp exactly as proposition 3.2 requires; two signed kernels (centers dotted) seat them, reducing the miss to 0.7 vol bp while the wings — and the fitted slice’s admissibility — are preserved.

Two honest boundaries of the construction, each one sentence. *Wing-neutral is weaker than wing-admissible*: the kernels preserve the base slopes (3.35), but nothing in the family forces those slopes into $[0, 2]$ — admissibility of the base is a separate condition, checked like any other. And *the family is not arbitrage-free by construction*: a signed kernel can drive variance negative or break price convexity locally, so the dense-grid check of section 3.4 decides for MCS exactly as it does for SVI, and the parameters $(\gamma_r, \zeta_r, \ell_r, \kappa_r)$ are means rather than objects — two overlapping kernels can trade amplitude, so only the fitted curve is comparable across fits.

3.7.4 Capacity and its governance

The number of cores M is the family’s capacity dial, and counting degrees of freedom already bounds it: $6 + 4M$ parameters against n_q quotes requires

$$M \leq \left\lfloor \frac{n_q - 6}{4} \right\rfloor \quad (3.38)$$

for the fit to be determined at all, the analogue for this family of the order guard (2.50). But identification is the weaker constraint. A kernel is a local feature detector, and on a smile that has no local features it will find quote noise instead: the fit error against the quotes falls while the error against the underlying curve rises, and the manufactured curvature appears exactly where section 3.4 taught us to look — in the sparsely quoted region, as a negative dip of g_D . Section 3.8 measures this on the frozen snapshot: at $M = 2$ on a liquid index chain the family fits the mid quotes better than SVI on every node while failing the validity check on most of them; at $M = 0$ the margins are uniformly clean and the fit quality is SVI’s. The capacity is real and so is its price. The reference implementation therefore treats M as a governed dial — capped by (3.38), ridged in amplitude, and reserved for smiles that are genuinely non-convex — rather than a default. The calibration itself is layered to match the model’s structure: the base is fitted first, the kernels are then seeded on the largest remaining residuals (one feature each), and a final joint refinement adjusts all $6 + 4M$ parameters together.

3.8 Three families on the same quotes

The chapter closes by putting the three families side by side on the same data: the frozen snapshot of appendix 2.C, all 16 nodes, under one protocol. Each family is calibrated by the reference implementation to the same mid quotes with equal weights — LQD at order 16 under the guard (2.50) in its logistic chart, SVI in the structural chart of section 3.6.2, and MCS at its capacity $M = 2$ under (3.38) — and each fit reports the same three quantities: the rms distance to the quotes, the validity margin $\min g_D$ over a dense grid on the traded range, and its wing slopes. Every Durrleman factor is evaluated from the family’s own analytic structure; appendix 3.A records the details. The point of the exercise is not a ranking. It is to watch each family’s contract — what it guarantees, what it must have checked — show up in numbers.

3.8.1 The December node and the SPY gallery

Figure 3.8 shows the SPY December 2026 slice. The three fitted curves are indistinguishable at chart scale; the information is in the residual panel. LQD tracks the mid quotes at 3.4 vol bp rms with an unstructured residual. SVI carries 7.5 vol bp with the long-wavelength wave of a five-parameter curve absorbing a seventeen-strike skew — the visible cost of parsimony. MCS at $M = 2$ sits between, 5.6 vol bp, its two kernels spent on the largest features of the SVI-shaped residual. All three agree on the wings to within a few hundredths of slope ($\beta_R = 0.034, 0.066$, and 0.054 respectively), which is the pattern section 2.4 led us to expect: no finite strip identifies an asymptotic slope by itself, and three different conventions continue similar evidence three different ways.

Figure 3.9 extends the reading to the eight SPY expiries, and its second panel is the one to study. On fit quality the ordering is stable: LQD, then MCS, then SVI, with the gap widening at the long end where the skew’s shape departs furthest from a hyperbola. On the validity margin the picture inverts. LQD’s margin is positive on every node — proposition 2.4 made that structural, and the panel is simply the numerical trace of a theorem. SVI’s margin is positive on every node as well, not by theorem but by rigidity: five parameters fitted to a liquid convex smile have no capacity left to bend the density negative. MCS at $M = 2$ fails the check on most of the SPY expiries — its own sampled hinge on g_{mD} notwithstanding, the concrete form of




Figure 3.8: Three families, one protocol, on SPY December 2026. (a) The fitted curves overlaid at chart scale inside the bid-ask band. (b) The residuals against the mid quotes separate them: LQD (3.4 bp rms, 17 parameters), MCS (5.6 bp, 14), SVI (7.5 bp, 5). More coordinates buy a whiter residual.

section 3.4’s warning that a finite-weight fence bends the optimum without bounding it. Its kernels, having no genuine local features to explain on a liquid index smile, chase quote noise, and the manufactured curvature surfaces as a negative g_D dip — typically in the thinly quoted region where no data resists it. Across all 16 nodes the medians read:

	median rms (bp)	worst rms (bp)	median min g_D	nodes with margin < 0
LQD (order 16, guarded)	7.3	34.0	0.184	0 of 16
SVI (structural chart)	19.0	68.0	0.173	0 of 16
MCS, $M = 2$	11.8	38.5	-0.026	10 of 16
MCS, $M = 0$ (base only)	18.1	—	0.179	0 of 16

The last row isolates the dial: withdrawing the kernels returns the margins to uniformly positive and the fit quality to SVI’s level. Capacity is not free and not fatal — it is a choice, and on these chains the correct choice is zero.

3.8.2 Expressiveness measured on a legal target

The gallery shows what excess capacity does where it is not needed. The complementary question — what rigidity costs where capacity *is* needed — must be asked against a target that an arbitrage-aware family is allowed to reproduce. A two-lognormal mixture supplies one: weights 0.65 and 0.35, component volatilities 14% and 42%, second component forward 0.85 and the first fixed by the martingale identity ($f_1 = 1.081$), quarter year to expiry. Its law is genuine by construction, and the exact density gives $\min g_D = 0.369$ on the quoted range — bimodal, but strictly admissible. Figure 3.10 fits the three families to its smile. LQD follows the two modes at 3.6 vol bp with margin 0.371. SVI misses by 166.4 vol bp — proposition 3.2 in numbers, one belly against two modes — while remaining perfectly valid ($\min g_D = 0.431$): the failure is expressiveness, not arbitrage. MCS halves SVI’s miss to 87.3 vol bp but overdrives its kernels into a small violation (-0.066): on this target its capacity is necessary yet still needs its checks. No family wins all three columns, which is the result.




Figure 3.9: The SPY gallery under one protocol. (a) Fit quality by expiry: more parameters fit better, uniformly. (b) The validity margin by expiry: LQD is positive by construction, SVI by rigidity, and MCS at $M = 2$ — its kernels chasing noise on a smile with no genuine local features — fails the dense-grid check on most nodes. The two panels together are the chapter’s thesis in one figure: in the volatility chart, fit quality and validity are separate accounts.




Figure 3.10: The rigidity question asked legally: a bimodal but strictly admissible mixture law ($\min g_D = 0.369$). (a) The smile and the three fits. (b) Signed errors: SVI’s single belly leaves a $\pm$ two-vol-point structured miss; MCS halves it; LQD follows the modes at quote scale.

	LQD (Ch. 2)	SVI	MCS
Validity	structural: every admissible vector is a law (proposition 2.4)	policed: floor and cap structural in the chart; butterfly checked on a grid	policed: wings neutral by lemma 3.8; butterfly checked on a grid
Tails	exponential class, slopes from endpoint speeds (proposition 2.9)	exactly linear wings; slopes capped at $\beta_{\max} < 2$	base-owned wings (3.35); admissibility a separate check
Interior capacity	polynomial order, guarded by (2.50)	one belly, one turn (proposition 3.2)	M kernels, guarded by (3.38); governed, not default
Interpretability	endpoint speeds and ledger functionals	JW handles on the domain of theorem 3.5; curvature lost on the stratum	curve-level only: kernels are non-unique by design
What convention supplies	order, ridge, tail chart	$\beta_{\max}$, chart, weights	M , ridge, weights

Table 3.2: What each family guarantees, what must be checked, and what the modeler chooses. The families do not dominate one another; they allocate the same difficulty differently.

3.8.3 The chapter's contract

The comparison, the mixture, and the gallery say one thing three ways, and it is the chapter's summary. Working in the volatility chart does not remove the constraint that a smile be a probability law; it moves the constraint from the coordinates into the process — inequalities the chart can absorb, a check the chart cannot, and capacity that must be governed because nothing structural governs it. What the chart buys in exchange is legibility: coordinates a desk can quote, wings that are two readable numbers, corrections that are local by construction. Chapter 2 and this one are therefore not competitors but the two poles of one design axis — validity built into the coordinates versus coordinates built for the reader — and the choice between them allocates difficulty; it does not eliminate it. The next chapter moves to a third position on the same axis: local volatility, where the modeled object is neither a law nor a curve but the field consistent with all of them at once.

3.A Comparison protocol and numerical notes

Protocol

Every number and figure in this chapter is produced by one deterministic script from the frozen snapshot of appendix 2.C and the fixed synthetic constructions below; nothing is typed by hand. The three-family comparison of section 3.8 calibrates each family to the same mid quotes of each node with equal unit-mean weights:

- **LQD:** order 16 reduced by the guard (2.50), logistic tail chart, ridge $\alpha = 10^{-6}$ with $\nu = 1$ — the recipe of section 2.6 with the mid objective.
- **SVI:** the structural chart (3.26) with $\beta_{\max} = 1.95$ and hinge multiplier $W = 10^3$ (the rows are structurally zero in this chart; the multiplier matters only for the historical chart (3.25)).
- **MCS:** $R = 2$ under the cap (3.38), amplitude ridge 10^{-2} , and a one-sided hinge on sampled g_D over a grid extending two ζ -units past the quoted range, weighted double on the put

side — a soft fence with the same status as (3.10): zero on any admissible fit. Kernel centers may roam half a unit past the quoted range; half-widths are confined to $[0.15, 1.5]$ and steepnesses to $[1, 12]$ in ζ -units. The base is fitted first, kernels are seeded greedily on the largest residuals with a masked neighborhood so two kernels cannot seed one feature, and a joint refinement follows.

The validity margin is $\min g_D$ over 801 equally spaced points spanning exactly the node's quoted log-moneyness range.

Evaluating g_D per family

Each family's Durrleman factor is computed from its own analytic structure, never by differencing prices — a finite difference of reconstructed prices carries curvature noise at exactly the scale the diagnostic measures. For SVI, (w, w', w'') are the closed forms (3.3). For MCS, the ζ -space derivatives of (3.37) are analytic (second differences of (3.32)) and convert by $w' = \sqrt{\tau} V' / \sigma_{\text{ref}}$, $w'' = V'' / \sigma_{\text{ref}}^2$. For LQD no derivatives are needed at all: (2.5) states $f_X(k) = g_D(k) \varphi(d_-) / \sqrt{w}$, so

$$g_D(k) = \frac{f_X(k) \sqrt{w(k)}}{\varphi(d_-(k, w(k)))}, \quad d_- = -\frac{k}{\sqrt{w}} - \frac{\sqrt{w}}{2},$$

with f_X the slice's explicit density — the identity that defined g_D is run backwards. For the mixture target of fig. 3.10 the same identity is used with the exact mixture density, which is how the target's own margin is known rather than estimated.

Synthetic constructions

The W-shaped target of fig. 3.7 is built in total variance so that its wings are asymptotically linear and its derivatives are exact:

$$w(k) = 0.008 + 0.012\sqrt{k^2 + 0.15^2} + 0.0025(e^{-\{(k-0.28)/0.11\}^2} + e^{-\{(k+0.28)/0.11\}^2}), \quad \tau = 0.25,$$

sampled at 61 equally spaced quotes on $[-0.5, 0.5]$. Gaussians vanish with all derivatives, so both wings carry the asymptotic slope 0.012 — far under Lee's bound — and the tail limit of g_D is $(4 - 0.012^2)/16 = 0.250$; the script verifies $\min g_D = 0.277 > 0$ on a dense grid to $|k| = 12$ from the analytic derivatives before any figure is written. A target built carelessly in volatility space with linear volatility tails would violate Lee far outside the plotted window: total variance would grow quadratically. Certify a synthetic target the same way as a fit, globally, or it certifies nothing.

The mixture target is the two-lognormal law of section 3.8: weights (0.65, 0.35), volatilities (0.14, 0.42), second component forward 0.85 and the first fixed by the martingale identity $0.65 f_1 + 0.35 \cdot 0.85 = 1$ (so $f_1 = 1.0808$), $\tau = 0.25$, quoted at 25 strikes on $k \in [-0.5, 0.4]$ through exact Black inversion of the mixture call.

Floating-point strictness

The structural chart's guarantee — wings strictly inside $(0, \beta_{\max})$, floor strictly positive — is a statement about real numbers, and float64 does not honor it automatically. Two boundary effects deserve record. A logistic lift evaluated at a moderately large coordinate rounds to exactly 1, parking a wing exactly at the cap; the implementation clips the lift one unit in the last place inside 1. And when one wing saturates against an ordinary other, the quotient for ρ_S in (3.27) can round to ± 1 , making $1 - \rho_S^2$ an exact zero and m_S undefined; the implementation computes it by the algebraically identical product form

$$1 - \rho_S^2 = \frac{4 \beta_L \beta_R}{(\beta_L + \beta_R)^2},$$

which never rounds to zero for admissible wings. A chart whose strictness silently fails at the boundary reproduces exactly the failure mode it was built to remove.

Chapter 4

Local Volatility

Chapters 2 and 3 fitted one law per expiry and let a calendar condition relate neighbors. This chapter asks for more: one diffusion whose marginal laws are all of the surface's laws at once. The Dupire equation supplies it, and in the volatility chart its coefficient is a ratio of the two objects the static chapters policed — the calendar increment over the Durrleman factor — so the field exists exactly where the surface is valid. Read as a formula, the ratio divides two derivatives of noisy data and fails structurally; read as a pricing map, it poses a well-behaved inverse problem: a piecewise-affine sheet whose positivity is a box constraint on finitely many numbers, priced by a monotone implicit march whose step matrix is an M -matrix, and differentiated exactly behind the same factorization. The quotes do not determine the sheet — two surfaces ten volatility points apart reprise the same quote set within a fraction of a spread — so calibration returns a declared representative of an equivalence class, and what the discretization leaks is measured, never assumed.

4.1 One diffusion for the whole surface

The static chapters treated an expiry as a law. Chapter 2 built the law first and read the smile off it; Chapter 3 worked in the volatility chart and policed validity there. Either way, a fitted surface is at this point a *family* of laws, one per expiry, related pairwise by the calendar comparison of section 2.8: at equal log-moneyness, later normalized calls dominate earlier ones, which is convex order of the normalized returns. Convex order is exactly the condition under which the family is jointly realizable: by Kellerer's theorem [18], *some* martingale has these marginals. But Kellerer's theorem is an existence statement. It does not say which martingale, it does not give the process any structure a computation could use, and it leaves open what additional information a particular choice smuggles in.

Dupire's observation [8] resolves all three questions inside one class. Adopt, as a standing hypothesis, the deterministic-carry setting of section 2.8, under which $Y_\tau = S_\tau / F_{0,\tau}$ is a single positive martingale whose time- τ marginal is Chapter 2's normalized return for the expiry- τ slice. Consider diffusions of the form

$$dY_\tau = \sqrt{v(\tau, Y_\tau)} Y_\tau dW_\tau, \quad Y_0 = 1, \quad (4.1)$$

where $v(\tau, y) \geq 0$ is a deterministic field, the *local variance*; its square root $\sigma_{\text{loc}} = \sqrt{v}$ is the local volatility. Within this class the correspondence between coefficient and marginals is exact in both directions: the field determines every marginal law by solving a forward equation, and the marginals determine the field by a pointwise formula, wherever the density is positive. A valid surface is therefore not merely *consistent with* some model — it *is* the coefficient field of one canonical model, written in different coordinates. This chapter is about that change of coordinates: how to compute it, why the obvious way of computing it fails on real data, and what a market's finitely many quotes do and do not say about the field.

Two remarks bound the object's ambitions before we start. First, the diffusion (4.1) is a *representative*, not a discovery. By Gyöngy's theorem [16], any Itô martingale (under mild regularity) with the same marginals is mimicked by (4.1) with $v(\tau, y)$ equal to the conditional expectation of its instantaneous variance given $Y_\tau = y$: the local-volatility model is the Markovian projection common to every dynamics the statics allow. Statements that depend only on marginal laws — every European price, every quantity of Chapters 2–3 — are safe; statements about joint laws, such as how the smile moves when the spot moves, are not determined by the surface at all and are postponed to Chapter 9. Second, the class (4.1) contains no jumps. A genuine discontinuity — an earnings gap, a scheduled announcement — can only be imitated by a burst of diffusive variance; scheduled events belong to a change of clock, which is Chapter 8's subject. Throughout this chapter τ is the variance time used since Chapter 2, $y = e^k$ the normalized strike, and $c(\tau, y)$ the normalized call: rates, borrow and dividends live inside the forward (Chapter 6), and the field v is variance per unit variance time.

The plan follows the two directions of the correspondence. Section 4.2 derives the forward equation, shows that positivity of the field propagates both static no-arbitrage families at once, and identifies the inverse formula as a ratio of two objects the reader already owns. Section 4.3 demonstrates why that formula must not be applied to data. Sections 4.4 and 4.5 build the objects that replace it: a finite-dimensional field whose validity is a box constraint, and a discretization that inherits the continuous theorem instead of merely approximating it. Section 4.6 asks what the quotes actually determine, and sections 4.7 and 4.8 assemble the calibration and run it on the book's frozen data.

4.2 The equation the surface must obey

4.2.1 Forward: from the field to every marginal

Fix a bounded measurable field $v \geq 0$ and let Y solve (4.1). Everything follows from watching the call price evolve in maturity.

Proposition 4.1 (The forward Dupire equation). *Suppose Y_τ has a density $f(\tau, \cdot)$ such that vy^2f and $y \partial_y(vy^2f)$ are integrable and vanish at $y = 0$ and as $y \rightarrow \infty$. Then the normalized call $c(\tau, y) = \mathbb{E}[(Y_\tau - y)^+]$ satisfies*

$$\partial_\tau c(\tau, y) = \frac{1}{2} v(\tau, y) y^2 \partial_{yy} c(\tau, y), \quad c(0, y) = (1 - y)^+. \quad (4.2)$$

Proof. By Itô's formula, the density satisfies the Fokker–Planck equation $\partial_\tau f = \frac{1}{2} \partial_{yy}(vy^2f)$ in the weak sense. Integrate it against the payoff:

$$\partial_\tau c(\tau, y_0) = \int_0^\infty (y - y_0)^+ \partial_\tau f(\tau, y) dy = \frac{1}{2} \int_0^\infty (y - y_0)^+ \partial_{yy}(vy^2f) dy.$$

Integrating by parts twice moves the two derivatives onto the payoff, whose derivatives in y are the step $\mathbf{1}_{\{y > y_0\}}$ and then the point mass δ_{y_0} , while the boundary terms vanish by the decay hypothesis. What remains is $\frac{1}{2} v(\tau, y_0) y_0^2 f(\tau, y_0)$, and $f(\tau, y_0) = \partial_{yy} c(\tau, y_0)$ is the density identity of section 2.1. The initial condition is $Y_0 = 1$. $\square$

Equation (4.2) is a *forward* parabolic equation in the strike variable with the payoff as initial data: heat flow, with $\frac{1}{2}vy^2$ as a spatially varying conductivity. One field, one linear march, and every marginal law of the surface comes out at once — this is the computational engine of the whole chapter. The analogy with heat is not decoration; it is the source of every structural property we will use. Heat flow spreads bumps and never sharpens them, and spreading, translated into option language, is exactly the no-arbitrage direction:

Proposition 4.2 (Positive local variance propagates no-arbitrage). *Let c solve (4.2) with $v \geq 0$ bounded. Then for every $\tau > 0$: (i) $\partial_{yy}c(\tau, \cdot) \geq 0$ — each slice is butterfly-free; (ii) $\partial_\tau c(\tau, y) \geq 0$ — the normalized call is non-decreasing in maturity at fixed y , which is the calendar order of section 2.8.*

Proof sketch. Differentiating (4.2) twice in y shows that $f = \partial_{yy}c$ satisfies $\partial_\tau f = \frac{1}{2} \partial_{yy}(vy^2 f)$ — the Fokker–Planck equation again — with initial datum $\partial_{yy}(1 - y)^+ = \delta_1 \geq 0$. Such equations preserve nonnegativity: probabilistically, $f(\tau, \cdot)$ is the density of the diffusion (4.1) started at 1; analytically, apply the parabolic maximum principle to the adjoint equation. With $f \geq 0$ established, (ii) is (4.2) itself: $\partial_\tau c = \frac{1}{2}vy^2 f \geq 0$. $\square$

Pause on what this proposition buys, because it is the entire strategy of the chapter. In Chapters 2 and 3, validity was work: the LQD chart made butterfly freedom an identity within one slice, the volatility chart had to police $g_D \geq 0$, and calendar order across slices was in both cases a separate condition. Here, positivity of *one* function of two variables propagates *both* families across the whole surface, for free, forever. The rest of the chapter is the discipline of not squandering that theorem: choose a parametrization in which $v > 0$ is trivial to enforce (section 4.4), choose a discretization that inherits the maximum principle rather than merely approximating it (section 4.5), and then measure what little the finite grid leaks instead of assuming the theorem carried over.

4.2.2 Backward: from the surface to the field

Read (4.2) the other way. If the surface is given, the field is determined pointwise,

$$v(\tau, y) = \frac{\partial_\tau c(\tau, y)}{\frac{1}{2}y^2 \partial_{yy}c(\tau, y)}, \quad (4.3)$$

wherever the density in the denominator is positive. This is the promised uniqueness: within the class (4.1), at most one field is consistent with a given surface, and (4.3) computes it. The same formula becomes most instructive in the volatility chart.

Proposition 4.3 (The field in the volatility chart). *Write the surface in implied form, $c = B(k, w(k, \tau))$ with $w(\cdot, \tau)$ the total-variance slice. Then (4.3) reads*

$$v(\tau, e^k) = \frac{\partial_\tau w(k, \tau)}{g_D(k)}, \quad (4.4)$$

with g_D the Durrleman factor of (2.5), evaluated on the slice at maturity τ [12].

Proof. Numerator: at fixed k , $\partial_\tau c = \partial_w B \cdot \partial_\tau w$. With $d_\pm = -k/\sqrt{w} \pm \sqrt{w}/2$ one has $d_+^2 - d_-^2 = (d_+ + d_-)(d_+ - d_-) = -2k$, hence $\varphi(d_+) = e^k \varphi(d_-)$, and

$$\partial_w B = \varphi(d_+) \partial_w d_+ - e^k \varphi(d_-) \partial_w d_- = \varphi(d_+) \partial_w (d_+ - d_-) = \frac{e^k \varphi(d_-)}{2\sqrt{w}}.$$

Denominator: $\partial_{yy}c = f(\tau, y) = f_X(k)/y$ by the change of variable $y = e^k$, and (2.5) gives $f_X(k) = g_D(k) \varphi(d_-)/\sqrt{w}$. Substituting both into (4.3),

$$v = \frac{e^k \varphi(d_-) \partial_\tau w / \sqrt{w}}{y^2 f_X(k) / y} = \frac{\varphi(d_-) \partial_\tau w}{\sqrt{w} f_X(k)} = \frac{\partial_\tau w}{g_D(k)}. \quad \square$$

The reader has met both halves of this ratio before, as the two families of static validity. The numerator is the calendar increment: nonnegative exactly when adjacent slices are ordered in the sense of section 2.8. The denominator is positive density written in the volatility chart, the object Chapter 3 policed slice by slice. So (4.4) is a compatibility theorem disguised as

a formula: *a local-volatility field exists precisely where the surface is valid*, and each mode of static arbitrage corresponds to its own sign failure — a calendar violation makes the numerator negative, a butterfly violation the denominator. This chapter therefore takes as input hypothesis what Chapters 2–3 produced as output: slices that are butterfly-free and calendar-ordered, in the normalized units of section 2.8. Where the hypothesis fails, no field exists and (4.4) says so by changing sign.




Figure 4.1: The field as a ratio, on the frozen SPY surface (stored per-expiry fits of the snapshot, total variance interpolated linearly in τ between the eight expiries; horizontal bars mark each expiry’s quoted span). (a) The calendar increment $\partial_\tau w$: piecewise constant in τ by construction, steepening toward the short-dated put side. (b) The Durrleman factor of each slice. (c) Their ratio, as local volatility: the put wall, the quiet call side, and the short-dated structure. Grey cells are inadmissible (a negative numerator or denominator): they amount to 0.5% of the displayed mesh, all in extrapolated territory — restricted to the common quoted support of each expiry pair the share is 0.00%, and the extracted local volatility there ranges from 9.0% to 61.5%. Validity where the data speaks, sign failures where conventions extrapolate: the ratio is already a diagnostic instrument.

Figure 4.1 evaluates the ratio on the book’s frozen SPY surface — the eight fitted slices of the snapshot, interpolated linearly in total variance between expiries. The exercise makes three points at once. The field inherits recognizable structure from the smile: a steep put-side wall, a quiet call side, short-dated sharpness. The two sign conditions are visibly the binding ones: every inadmissible cell in the picture sits outside the common quoted support of the expiries that bracket it, in territory owned by extrapolation conventions rather than prices. And the ratio was computed here from *fitted, smooth* slices — an operation this chapter endorses. What it must never be applied to is the market’s raw quotes, and the next section shows why.

4.3 The wrong direction

Formula (4.3) invites a procedure: interpolate the quoted smile, differentiate it twice in strike and once in maturity, divide. On a continuum of exact prices this is a theorem. On a real chain — a few dozen noisy quotes per expiry — it fails, and the failure is structural, not bad luck.

Quantify it. Suppose each quoted total variance carries an error of size ε , and strikes are spaced Δk apart. A second difference estimates curvature as

$$\partial_{kk} w \approx \frac{w(k - \Delta k) - 2w(k) + w(k + \Delta k)}{\Delta k^2},$$

and the worst configuration of errors — alternating signs, which is exactly what a bid–ask bounce produces — enters this estimate at size $4\varepsilon/\Delta k^2$. The noise is amplified by the *inverse square* of the strike spacing: a finer quote grid makes the estimate worse, without bound. There

is no stable limit to converge to; numerical differentiation of noisy data is the canonical ill-posed operation. And the place this amplified noise lands could not be worse chosen: by (4.4) the curvature sits inside g_D , the sign-carrying factor of the risk-neutral density. The first time the noise pushes the estimated density through zero, the extraction returns a *negative variance* — not a large error but a non-object, which no pricing equation, no simulation and no risk system downstream can absorb.




Figure 4.2: The same synthetic market, read both ways. Quotes are generated from a known smooth field (dashed) and carry a deterministic alternating ± 50 bp volatility ripple — a stylized bid-ask bounce. (a) The extraction (4.4) from the noisy quotes at $\tau = 0.25$. From 13 quotes per expiry it sawtooths between 13% and 73% around a true ATM level of 23%; from 31 quotes of the *same* market — finer Δk , same ± 50 bp — it fails outright at 14 strikes (crosses: negative variance or negative density), exactly the $4\epsilon/\Delta k^2$ scaling at work. (b) The same 31-quote noisy chains handed to the forward calibration of this chapter: an admissible surface by construction, within 12.9 volatility points of the truth along the plotted cross-sections, with reprice residuals of rms 49 bp — the scale of the ripple itself. The noise stays in the residual, where it belongs.

Figure 4.2 performs the experiment on a synthetic market whose answer is known (the field and quote recipe are in appendix 4.A). The left panel is the naive pipeline at two quote densities and needs little commentary beyond its caption: refining the strike grid, which improves every well-posed computation, makes this one *worse*, quadratically, until half the strikes report that no diffusion exists. The right panel previews the alternative developed in the rest of the chapter. The same noisy quotes are handed to a calibration that never differentiates them; every iterate it produces is a genuine positive field; and the fitted surface tracks the truth through the quoted region while its residuals absorb the ripple. One caveat, quantified now and explained in section 4.6: at the edges of the shortest expiry’s quoted span, where a single rippled quote is the only witness, the fitted field wanders by as much as 22 volatility points. Noise rejection is a property of dense data and declared smoothing.

The repair is a change of role, not of formula, and it is due to Dupire’s equation being read in its *forward* direction, as Andreasen and Huge taught [1]. Decide that the unknown is the field v itself. Parametrize it. Price it through (4.2) — which differentiates only the smooth numerical solution, never the data — and ask an optimizer for the positive field whose prices meet the quotes. Positivity is imposed on the *parameters*, so every iterate of the search is an admissible diffusion before any quote is consulted. Readers of Chapters 2–3 will recognize the design: it is the same constraint-elimination idea that built the LQD chart and the structural SVI chart (sections 2.6.3 and 3.6.2), one dimension up — choose the coordinates in which the constraint set is trivial, and let the pricing map do the hard work. Three questions remain,

and they organize the rest of the chapter: what parametrization makes positivity trivial in two dimensions (section 4.4); what discretization of (4.2) keeps proposition 4.2 alive on a finite grid (section 4.5); and what, exactly, the quotes determine about the field once the formula (4.3) is off the table (section 4.6).

4.4 A sheet with finitely many parameters

4.4.1 Positivity as a box

The unknown is now a positive function of two variables — an infinite-dimensional object with a one-sided constraint, which sounds no easier than where we started. The move that makes it finite is standard in finite elements, and its payoff here is larger than convenience.

Definition 4.4 (P_1 local-variance sheet). Fix vertices (τ_i, y_j) , $i = 1..n_\tau$, $j = 1..n_y$, on a tensor grid, triangulated by splitting each rectangular cell along one diagonal (Delaunay, remark 4.6). The local variance is the continuous piecewise-affine interpolant

$$v_\theta(\tau, y) = \sum_{\ell=1}^{n_\theta} \theta_\ell \mathcal{H}_\ell(\tau, y), \quad n_\theta = n_\tau n_y, \quad (4.5)$$

where $\mathcal{H}_\ell$ is the hat function of vertex ℓ — barycentric on each triangle, equal to 1 at its own vertex and 0 at every other — and the parameters $\theta_\ell = v(\tau_i, y_j) > 0$ are the *nodal local variances*, indexed flat as θ_ℓ or by row and column as $\theta_{i,j}$, whichever is convenient.

Why affine, and why triangles? Because affine interpolation is a *convex combination*: on each triangle, v_θ is a weighted average of its three vertex values with nonnegative barycentric weights summing to one. Averages cannot escape the range of what they average, which is the whole theorem:

Proposition 4.5 (Nodal bounds are surface bounds). *If $v_{lo} \leq \theta_\ell \leq v_{hi}$ for all ℓ , then $v_{lo} \leq v_\theta(\tau, y) \leq v_{hi}$ everywhere on the triangulation's hull.*

Positivity of a surface has been reduced to positivity of finitely many numbers — a box constraint, which every least-squares solver handles natively, with no penalty terms. This is the promised chart, one dimension up: as in the LQD chart of Chapter 2 and the structural chart of section 3.6.2, the coordinates are chosen so that every point of the parameter box is a valid object, and the optimizer never needs to know a constraint existed. Combined with proposition 4.2, the chain of guarantees is complete before any data arrives: box $\Rightarrow$ positive field $\Rightarrow$ butterfly-free, calendar-ordered surface. Figure 4.3 shows the object itself — the sheet this chapter will calibrate to the frozen SPY chains in section 4.8 — and fig. 4.4 takes it apart.

Remark 4.6 (Which diagonal, and does it matter). On a tensor grid every cell is a rectangle, and a rectangle's corners are cocircular — the degenerate case in which both diagonals are equally Delaunay. A triangulation routine therefore tie-breaks cell by cell, deterministically but by no designed convention, and figs. 4.3 and 4.4 draw the resulting patchwork faithfully. The choice is second-order small: the two candidate interpolants of a cell agree at all four corners and along all four edges, and differ in the interior by at most half the cell's twist, $(\theta_{i,j} + \theta_{i+1,j+1} - \theta_{i,j+1} - \theta_{i+1,j})/2$, at the center — zero whenever the nodal values are locally bilinear. It is nonetheless a convention, and the reference implementation pins it: the triangulation is computed once, cached, and the same triangles price every iterate.

4.4.2 Beyond the hull: the wing contract

The grid must end somewhere, and pricing does not. Below the lowest strike column the sheet continues local variance *linearly*, with slope a times the slope of the first interior cell; above




Figure 4.3: The model, drawn: the calibrated SPY local-volatility sheet of section 4.8 (9 maturity rows $\times$ 24 strike columns, 216 vertices), with the triangulation the pricing basis actually uses. Every triangle is a region where local variance is an affine function of (τ, y) ; every dot is one calibration parameter. The put-side wall, the short-end structure and the quiet call-side plain are all just nodal values: the whole surface is this list of positive numbers, nothing else.




Figure 4.4: The anatomy of the sheet, on a small demonstration grid. (a) The triangulated vertex grid; the shaded star of triangles is the support of one vertex's hat function — the entire region of the surface its parameter influences. The cell diagonals are not uniformly oriented; that is real (remark 4.6). (b) A strike cross-section along a vertex row: each hat is 1 at its own vertex, 0 at every other, affine between, and they sum to 1 identically — the partition of unity behind proposition 4.5.

the highest column it is clamped flat. Three regimes for a are meaningful: $a = 0$, the default, clamps the put wing flat as well; a fixed multiple keeps an unquoted put wing rising; and a is released to the optimizer exactly when a variance-swap quote is present — the one instrument in the problem that genuinely carries information about the deep-put tail, as Chapter 5 will show.

Remark 4.7 (The extrapolation sits outside the positivity guarantee). Proposition 4.5 stops at the hull, because outside it interpolation becomes extrapolation and the basis weights turn *signed*: an extrapolated value is no longer an average. With $a > 0$ and a first cell whose variance decreases toward the boundary, the linear continuation can reach zero before $y = 0$ (pricing floors it, so the priced curve kinks rather than blows up), and a rising continuation exceeds v_{hi} freely. Every guarantee in this chapter that leans on the box is therefore scoped to the hull and to the flat-clamped wings. The wing is a disclosed convention — the same epistemic status as the tail conventions of Chapters 2–3, and the chapter's contract table records it.

Remark 4.8 (Lee's bound, inherited). Bounded local variance bounds the smile. The marched call is *monotone in the field*: for $\delta v \geq 0$ the exact difference $c_{v+\delta v} - c_v$ solves $\partial_\tau s = \frac{1}{2}(v + \delta v)y^2\partial_{yy}s + \frac{1}{2}\delta v y^2\partial_{yy}c_v$, whose source is nonnegative by proposition 4.2 and whose propagator preserves sign — so the difference is nonnegative. When the wings are clamped, $v \leq v_{hi}$ holds on the whole strike line, so the surface is dominated by the constant-variance model: $w(k, \tau) \leq v_{hi} \tau$ for *every* k . Total implied variance is then bounded in strike, its wing slope decays to zero, and the surface sits strictly inside Lee's cap of 2 — automatically, where Chapter 3 had to legislate a cap (section 3.3). A steep free- a continuation escapes v_{hi} and with it this bound.

4.5 Pricing on the lattice

4.5.1 The implicit march and its M-matrix

The forward equation is solved on a strike grid over $[0, y_{\max}]$ with Dirichlet boundaries $c(\cdot, 0) = 1$ and $c(\cdot, y_{\max}) = 0$, by *fully implicit* Euler steps, the field evaluated at the new time level. On a possibly non-uniform grid with gaps $\Delta y_i^- = y_i - y_{i-1}$ and $\Delta y_i^+ = y_{i+1} - y_i$, the discrete generator

$\mathcal{L}$ encoding $\frac{1}{2}vy^2\partial_{yy}$ has the three-point stencil

$$\mathcal{L}_{i,i\mp 1} = \frac{v_i y_i^2}{(\Delta y_i^- + \Delta y_i^+) \Delta y_i^\mp} \geq 0, \quad \mathcal{L}_{ii} = -(\mathcal{L}_{i,i-1} + \mathcal{L}_{i,i+1}), \quad (4.6)$$

and each step solves one tridiagonal system,

$$\mathcal{M}^{n+1}U^{n+1} = U^n + (\text{boundary terms}), \quad \mathcal{M}^{n+1} = I - \Delta\tau \mathcal{L}^{n+1}. \quad (4.7)$$

Now the promised discipline: does this discretization *inherit* proposition 4.2, or merely approximate it? The answer runs through one classical linear-algebra fact, and it is worth seeing how little is needed.

Proposition 4.9 (The step matrix is an M-matrix). *For any $v \geq 0$ and any $\Delta\tau > 0$, the matrix $\mathcal{M} = I - \Delta\tau \mathcal{L}$ of (4.7) has positive diagonal, non-positive off-diagonals, and each diagonal entry exceeds the absolute sum of its row's off-diagonals by at least 1 (exactly 1 at fully interior rows). Hence $\mathcal{M}$ is a strictly diagonally dominant nonsingular M-matrix, and $\mathcal{M}^{-1} \geq 0$ elementwise.*

Proof. Read the sign pattern off (4.6): the off-diagonals of $\mathcal{M}$ are $-\Delta\tau \mathcal{L}_{i,i\mp 1} \leq 0$, the diagonal is $1 + \Delta\tau(\mathcal{L}_{i,i-1} + \mathcal{L}_{i,i+1}) \geq 1$, and the dominance gap is the 1 of the identity at interior rows, more where a stencil arm reaches a boundary column. This is the defining pattern of a nonsingular M-matrix, and M-matrices are inverse-positive [3]. (Concretely: Jacobi iteration for $\mathcal{M}u = b$ with $b \geq 0$ keeps nonnegative iterates and converges by dominance, so $u \geq 0$.) $\square$

$\mathcal{M}^{-1} \geq 0$ is the discrete ghost of the heat kernel being positive, and everything follows from it:

Corollary 4.10 (Discrete maximum principle, unconditionally). *For any $\Delta\tau > 0$ the scheme (4.7) is monotone — $U^n \geq V^n$ with the same boundary data implies $U^{n+1} \geq V^{n+1}$ — and every step keeps the values between the extremes of the previous step and the boundary data: the march is unconditionally ℓ^∞ -stable, with no step-size restriction, and cannot create new extrema.*

Proof. Monotonicity is $\mathcal{M}^{-1} \geq 0$ applied to $U^n - V^n \geq 0$. For the upper bound, let M_* be the larger of $\max_i U_i^n$ and the boundary values: because the interior rows of $\mathcal{M}$ sum to 1 once the boundary columns are included, the constant vector obeys $\mathcal{M}(M_*\mathbf{1}) \geq M_*\mathbf{1}$ against the right-hand side of (4.7), and inverse positivity turns that into $U^{n+1} \leq M_*\mathbf{1}$. The lower bound is symmetric. $\square$

Remark 4.11 (What the M-matrix argument does and does not prove). Proposition 4.9 and corollary 4.10 control the call *values*: order preservation, boundedness, no invented extrema. They do not by themselves give nonnegative second differences — a bounded monotone vector can still be locally concave — so “the scheme cannot manufacture negative densities” would require a separate convexity-preservation proof for the actual stencil and boundary treatment, which this chapter does not claim. What can be claimed is what the reference implementation does: the scheme is monotone and stable by theorem, and the butterfly and calendar cleanliness of the *output* is measured on the price grid with every fit. Section 4.8 reports the measured minima for the chapter's two calibrated sheets; they sit at the scale of solver rounding.

4.5.2 Why not Crank–Nicolson

First-order accuracy in time looks like a price worth negotiating, and Crank–Nicolson — replace (4.7) by $(I - \frac{\Delta\tau}{2}\mathcal{L})U^{n+1} = (I + \frac{\Delta\tau}{2}\mathcal{L})U^n$ — is the standard second-order offer. Look at what it costs. The implicit factor is still an M-matrix, but monotonicity now also needs the *explicit*

factor to be nonnegative, and its diagonal is $1 - \frac{\Delta\tau}{2}(\mathcal{L}_{i,i-1} + \mathcal{L}_{i,i+1})$: nonnegative only under the CFL-like bound

$$\Delta\tau \leq \frac{2}{\mathcal{L}_{i,i-1} + \mathcal{L}_{i,i+1}} = \frac{2\Delta y^2}{v_i y_i^2} \quad (\text{uniform grid}). \quad (4.8)$$

Refining the strike grid — which section 4.6 will show short-dated smiles *require* — tightens this bound quadratically. At a 40% local volatility on a $\Delta y = 0.005$ lattice the bound is 0.0003 years, more than thirty times smaller than the standard 0.01 marching step: monotone Crank–Nicolson on realistic lattices is simply not on offer. When the bound is violated the scheme remains second-order *accurate* but is no longer *monotone*, and the payoff kink at $y = 1$ excites the classic oscillation. Figure 4.5 shows what that means for the object we care about: the discrete density. Damped start-up steps (a few implicit steps before switching to Crank–Nicolson) are the standard repair [15]; the reference implementation’s default declines the negotiation altogether and trades one order of time accuracy for corollary 4.10.




Figure 4.5: Monotonicity is a property of the scheme, not of luck. The same flat 40% surface, the same $\Delta y = 0.005$ lattice, marched to $\tau = 0.1$ in ten steps by the same solver. Implicit Euler: a clean bell of second differences, minimum 0.0012 in the plotted window — corollary 4.10 at work. Undamped Crank–Nicolson: the kink oscillation swings the discrete density to -36 — butterflies invented out of nothing by a scheme that is, in the textbook sense, the more accurate one.

4.5.3 Reading out, and the operator’s blind spot

The observation operator is plain: the normalized call at a quote’s (τ, y) is read off the marched solution by interpolation, at maturities that are forced to be lattice time-nodes. One subtlety deserves a principle of its own, because it is general to inverse problems. An optimizer facing a *fixed* discretization will happily bend its parameters to cancel the operator’s own discretization error, so the in-operator residual systematically flatters the fit — validating a reconstruction with the operator that produced it is the inverse-problems “inverse crime”. The chapter’s protocol therefore reprises every fit on a *refined* operator (half the strike step, four times the time steps) and reports both numbers. The gap between them is not noise; it is the measured size of the compensation, and section 4.8 will show it is the dominant error term at the short end.

4.6 What the quotes determine

Before writing an objective function, a chapter should ask whether the problem it is about to pose has one answer. This one does not, in three distinct ways — one about resolution,

one about averages, one about whole regions of the field — and each dictates a piece of the calibration’s design.

The length scale. Everything interesting about a smile at maturity τ happens within a few multiples of the diffusive width $\sigma\sqrt{\tau}$ of the money. This single quantity dictates where both grids of sections 4.4 and 4.5 must put their points. *Vertex placement:* strike columns are spaced by delta — a fixed ladder of exercise probabilities mapped through Φ^{-1} , rather than fixed steps of strike — so parameter density follows where optionality lives. *Vertex coverage:* an axis sized by the longest expiry undersamples a short one, whose entire smile can fall between two columns — a two-day smile at 20% volatility is about 1.5% wide, an order of magnitude narrower than a one-year smile — so every expiry must be guaranteed several columns *within its own* range, spread evenly across it. *Lattice steps:* the same width bounds the strike step from above, and if the *quote* spacing is finer than the lattice step, the fitted smile oscillates at quote frequency until the lattice out-resolves it. A grid that does not resolve the data’s length scale fits an alias, not the data; the exact placement rules used in this chapter are collected in appendix 4.A.

An averaging fact, and the floor. At the money, implied and local variance are related as average to integrand. If the field near the money depends only on time, $v(\tau, y) \approx v(\tau, 1)$, the model is Black–Scholes with a time-dependent volatility and *exactly* $w(0, \tau) = \int_0^\tau v(s, 1) ds$; in general this holds to leading order near the money. An average can only be matched by an integrand that crosses it: whenever the ATM term structure is *increasing*, early local variance must sit strictly *below* the shortest implied variance. A parameter floor set carelessly — say at a round 5% volatility — therefore makes some low-volatility short-dated term structures literally unfittable, with the optimizer pinned to the box and the misfit looking like a mystery. The box of proposition 4.5 must be generous and *adaptive*: a cap a fixed multiple of the largest implied volatility in view, a floor a fraction of the *smallest ATM* implied — keyed to ATM quotes deliberately, so that one noisy deep-wing print cannot lower the floor where it does real stabilizing work.

Rows the quotes cannot see. The quotes at the first expiry τ_1 pin — through the averaging fact — only the *integral* of local variance over $[0, \tau_1]$. Any vertex row strictly inside that interval is individually unidentified: the optimizer can ring one row up and the next down at no cost in fit, and under noise it will (this is precisely how the rippled fit of fig. 4.2 spent its freedom at the shortest expiry, wandering 22 volatility points at a span edge held by a single quote). The cure is not more data but a declared prior: tie each sub-front row to the first identified row — the *front tie* of section 4.7 — turning an unidentified subspace into a stated constant continuation.

Two sheets, one quote set. The three paragraphs above are instances of one geometry, and it is better seen than told. A fit like this chapter’s carries a couple of hundred nodal parameters, and even when the quotes outnumber them several-fold — nearly five to one on the frozen SPY chain of section 4.8, between two and three on NVDA — the quotes cluster near the money of a handful of expiries, and a vanilla constrains local variance only through a smoothing integral along diffusion paths. Whole regions of the sheet (deep wings, sub-front rows, times beyond the last expiry) are therefore weakly determined no matter how the counting goes. Figure 4.6 makes it concrete on the synthetic running example: pushing the sheet’s unquoted deep-put column (at $y = 0.45$, below every quoted strike) down by 10 full volatility points changes the reprice of the entire quote set by at most 22.6 bp — 2.5 bp rms — concentrated in the longest expiry’s lowest strikes, the only quotes whose diffusion reaches the moved column at all. A thousand basis points of surface, twenty of quotes.

The consequence is philosophical and operational at once. What a calibration returns is one *representative* of an equivalence class of sheets the data cannot distinguish, and every device in




Figure 4.6: What sparse quotes do not see. (a) The calibrated synthetic sheet at $\tau = 0.5$ (solid) and the same sheet with its unquoted deep-put column pushed down by 10 volatility points (dashed); the shaded band is the quoted region, where the two are indistinguishable. (b) The absolute reprice change at every quote of every expiry: at most 22.6 bp, and only where the longest expiry’s lowest strikes feel the moved column. Low quote error is *not* surface recovery.

the next section’s objective is a declared choice of representative: the roughness term selects the smoothest member near a reference, the box clips the unidentified extremes, the front tie pins the rows below the first expiry, and the seed decides which basin a non-convex search lands in. None of this is hidden tuning; it is the “convention” column of the book’s thesis, operating inside a calibration. The worked example of section 4.8 accordingly reports quote error and surface error as two separate numbers.

4.7 The objective and its derivatives

4.7.1 The objective

The pieces assemble into a bound-constrained least-squares problem over the nodal variances:

$$\min_{\theta \in [v_{lo}, v_{hi}]^{n_\theta}} \frac{1}{2} \left(\sum_q r_q(\theta)^2 + \lambda \|\Gamma(\theta - \theta_{\text{ref}})\|^2 + \lambda_0 \|\theta_{1,\cdot} - \theta_{2,\cdot}\|^2 \right), \quad (4.9)$$

whose three blocks have each been argued for already.

The *data block* compares the marched call c_θ at quote q ’s coordinates with the quoted one, in vega units: $r_q = (c_\theta(\tau_q, y_q) - c_q)/\eta_q$, with η_q one volatility percent of the quote’s vega, floored (the same device as Chapter 2’s vega floor η ; this chapter’s constant is in appendix 4.A). This is the residual design of section 2.6 — solve in price space, measure in volatility space — carried over; when band edges rather than mids are trusted, the band objective of that section applies verbatim. The *roughness block* is Tikhonov regularization with a curvature seminorm: Γ takes second divided differences along strike within each maturity row and along time within each strike column, normalized by the local vertex spacing so that a tightly spaced money column is not penalized more than a wide wing one (it reduces to the $(1, -2, 1)$ stencil on a uniform grid). It pulls toward a reference sheet θ_{ref} — flat, at the median ATM variance, in this chapter’s protocol — and it is the declared answer to section 4.6: of the equivalence class the data allows, return the smoothest member near the reference. The *front tie* at weight λ_0 pins the $\tau = 0$ row to the first quoted row, closing the unidentified sub-front subspace. Two further row types plug into the same stack when their instruments are present, and are only named here: a variance-swap

quote adds one scalar residual per quoted expiry, constraining the deep-put tail integral that vanillas cannot see (Chapter 5 develops the integral, and it is exactly the instrument that justifies releasing the wing slope a of remark 4.7); and prior structures from a previous fit enter as reference terms in the same Γ -and- θ_{ref} position (Chapter 10 develops when and how much to trust them).

4.7.2 Every derivative from the same factorization

A least-squares solver lives on Jacobians, and a PDE-implicit pricing map sounds like finite-difference territory: n_θ rebuilds of the march per iteration. It is not, for the same structural reason that served the earlier chapters: the map from parameters to operator is *linear*. By (4.5), $\mathcal{L}(\theta)$ is linear in θ with $\partial\mathcal{L}/\partial\theta_\ell = \mathcal{L}[\mathcal{H}_\ell]$ — the stencil (4.6) with the field v replaced by the hat function $\mathcal{H}_\ell$.

Proposition 4.12 (Tangent: same factor, extra right-hand sides). *The sensitivities $s_\ell^n = \partial U^n / \partial \theta_\ell$ obey*

$$\mathcal{M}^{n+1} s_\ell^{n+1} = s_\ell^n + \Delta\tau \mathcal{L}[\mathcal{H}_\ell] U^{n+1}, \quad s_\ell^0 = 0, \quad (4.10)$$

where $\mathcal{L}[\mathcal{H}_\ell] U^{n+1}$ applies the boundary-inclusive stencil to the solution extended by its fixed boundary values. This is the discretization of the continuous sensitivity equation $\partial_\tau s_\ell = \frac{1}{2} v y^2 \partial_{yy} s_\ell + \frac{1}{2} \mathcal{H}_\ell y^2 \partial_{yy} c$.

Proof. Write one step over the interior unknowns as $\mathcal{M}^{n+1} U^{n+1} = U^n + b^{n+1}$, where b^{n+1} carries the Dirichlet data through the stencil's boundary columns — and therefore depends on θ through the field at the boundary-adjacent nodes. Differentiating in θ_ℓ and using $\partial_{\theta_\ell} \mathcal{M}^{n+1} = -\Delta\tau \mathcal{L}[\mathcal{H}_\ell]$ (exact by linearity) gives (4.10), the interior product $\mathcal{L}[\mathcal{H}_\ell] U^{n+1}$ and the derivative of b^{n+1} combining into the boundary-inclusive product, because the boundary values themselves do not move with θ . $\square$

Read (4.10) twice. As a computation: every sensitivity marches behind the *same* tridiagonal factorization as the value solve, so the entire n_θ -column Jacobian costs one factorization plus n_θ extra back-substitutions per step — the columns form a block of right-hand sides that a solver processes together. As probability: the source term says vertex ℓ influences a price exactly through the diffusion's visits to ℓ 's star of triangles, weighted by the convexity it finds there. Figure 4.7 shows both readings on the synthetic running example: the influence of one interior vertex is a cone — zero at the expiry *before* the vertex's time, which the march cannot reach backward; sharpest near its own maturity; spreading and flattening as the diffusion forgets where it collected its variance — and every analytic column agrees with a central finite difference of the full march to 1.1×10^{-8} in relative terms.

The optimization itself is then standard machinery used plainly: a trust-region reflective least-squares iteration on (4.9), box bounds handled natively [5], the analytic Jacobian of proposition 4.12 supplied at every iterate, and a flat seed at the reference sheet. For the tangent route the Jacobian cost grows linearly with n_θ ; the classical adjoint alternative — one backward sweep per objective, cost independent of the parameter count — is derived in appendix 4.A for the reader who will need it at larger grids.

4.8 Examples and the contract

All fits in this section use one protocol — the grids of section 4.6, the objective of section 4.7, every constant recorded in appendix 4.A — applied unchanged to a synthetic market whose answer is known and to the book's frozen SPY and NVDA chains. No fit is tuned separately.




Figure 4.7: The tangent system, seen and audited. (a) Sensitivity of the priced call curve to the single vertex at $(\tau, y) = (0.25, 1)$, across the four quoted expiries: zero at $\tau = 0.10$ (an earlier expiry cannot feel a later vertex), a sharp cone at $\tau = 0.25$, spreading thereafter. (b) Every analytic column of the Jacobian against a central finite difference of the full march, over all 50 vertices: worst relative disagreement 1.1×10^{-8} .

4.8.1 The round trip, and its two separate numbers

The synthetic market prices a known smooth field through an *independent* dense scheme (so the exercise cannot flatter the chapter's march by validating it against itself), quotes it at four expiries, and hands the 124 clean quotes to the calibration from a flat seed. Per section 4.6, the result is two numbers, not one. The *quotes* reprice to 0.6 bp rms and 1.8 bp at worst — the fit is essentially exact in the space the optimizer sees. Separately, the *field itself* agrees with the truth to 0.78 volatility points rms on the quote-covered region, with a maximum of 3.10 points — more than a hundred times the quote error, concentrated exactly where fig. 4.6 said the quotes do not look: between expiries and near the span edges (fig. 4.8b). Both numbers are honest; only reporting the first would not be.




Figure 4.8: The synthetic round trip. (a) Target quotes (dots) and the recovered fit (lines) at the four expiries: rms 0.6 bp, max 1.8 bp. (b) The second number: $|\text{fitted} - \text{true}|$ local volatility over the quote-covered region, quote positions dotted. The error lives between the expiries and at the edges of the quoted spans — largest where the nearest quote is farthest.

4.8.2 The frozen chains

The real test fits one sheet per underlier to every prepared mid quote of the snapshot’s eight expiries at once: 1026 quotes against 216 nodal variances for SPY, 506 against 207 for NVDA. Figure 4.3 already showed the SPY sheet itself; fig. 4.9 shows four of its expiries against the quotes, priced by the *one* calibrated diffusion — a two-day smile and a year-and-a-half LEAP from the same list of positive numbers.




Figure 4.9: Four SPY expiries priced by the single calibrated sheet, against the frozen quotes (dots: mids; whiskers: bid–ask). Panel titles give the expiry and its distance; annotations the per-expiry rms on the fitting operator.

Figure 4.10 reports the chapter’s two metrics side by side. On the fitting operator the surfaces read 2.7 bp (SPY) and 16.1 bp (NVDA) rms across all quotes — NVDA’s wider spreads and harder short-dated smiles carry more residual everywhere. Repriced on the refined operator the same surfaces read 13.6 and 46.1 bp, and the worst single expiry — in both names the one-day front — reads 28.2 and 116.9 bp. The gap between the paired bars is the point, and it is the operator’s blind spot of section 4.5 made quantitative: the optimizer partially cancels the coarse march’s error at the expiries it is fitted on, the in-operator residual flatters accordingly, and the number to quote is the refined one. Where the two agree (the long end) the fit is genuinely converged; where they diverge (day-scale maturities) the remaining error is discretization, not disagreement with the quotes. The measured cleanliness of remark 4.11 completes the report: across every expiry of both calibrated sheets, the minimum divided second difference of the marched prices is -8.5×10^{-13} (SPY) and -2.1×10^{-13} (NVDA), and the minimum adjacent-expiry increment is -8.9×10^{-16} and -7.8×10^{-16} — solver rounding, on both counts, as proposition 4.2 predicts and remark 4.11 declines to assume.




Figure 4.10: Per-expiry rms for SPY (a) and NVDA (b), on the fitting operator (grey) and on the refined operator (blue), log scale. The refined bars are the ones to read: the gap between the pairs is the compensated discretization error of section 4.5, and it owns the short end.

4.8.3 The contract

Structural	Numerically controlled	Outside the claim
Box bounds imply a positive field on the hull; a positive field implies butterfly-free, calendar-ordered marginals (propositions 4.2 and 4.5); every optimizer iterate is an admissible diffusion	Butterfly, calendar and bound cleanliness of the marched prices, measured per fit; refined-operator reprice beside the in-operator residual	Convexity preservation of the discrete scheme as a theorem; the extrapolation region beyond the hull, where basis weights are signed
Monotone, unconditionally ℓ^∞ -stable march (proposition 4.9 and corollary 4.10); exact tangent Jacobian from the same factorization (proposition 4.12)	Tangent columns audited against finite differences; an independent discretization cross-checks the march	Uniqueness of the calibrated sheet: the data determines an equivalence class, the conventions of section 4.7 a representative
Bounded field $\Rightarrow$ implied variance bounded in strike, strictly inside Lee's cap on the hull (remark 4.8)	Grid placement rules resolving each expiry's $\sigma\sqrt{\tau}$; adaptive variance box keyed to ATM quotes	Jumps and scheduled-event spikes (the clock, Chapter 8); spot-volatility dynamics (Chapter 9); the deep-put tail's shape without a variance-swap quote (Chapter 5)

Table 4.1: The local-volatility model contract: theorem, numerical control, and excluded scope.

The chapter's separation thesis reads off the table. *Validity* is structural and cheap: a box. *Information* is scarcer than the formula (4.3) pretends: quotes determine the field only through smoothing integrals, and only where they exist. *Convention* fills the rest — roughness, reference, tie, wing, seed. The calibration does not eliminate those choices; it declares them, measures what the discretization leaks, and reports the error twice: once as the optimizer saw it, once as it is.

4.A Numerical facts and the chapter protocol

4.A.1 The protocol

Every fit in this chapter — synthetic and real — uses the following recipe; no constant is tuned per name.

Vertex grid. Maturity rows at $\tau = 0$ and at every quoted expiry. Strike columns delta-spaced: the ladder of put deltas $\{1, 2, 5, 10, 25, 40\}\%$, their call mirrors, and the money, mapped to strikes through $y = \exp(\sigma_A \sqrt{\tau_{\max}} \Phi^{-1}(\cdot))$ with σ_A the longest expiry’s ATM volatility; the two end columns are extended to cover the pooled quoted range, so no quote falls outside the hull. Coverage rule: every expiry must own at least 6 columns inside its own traded range, enforced by splitting the widest boundary-augmented gap until none exceeds one fifth of the range (section 4.6); this is what densifies the short-dated money in fig. 4.3.

Lattice. Strike step $\Delta y = \min(0.01, 0.15 \min_e \sigma_e \sqrt{\tau_e})$, capped at 800 nodes over $[0, \max(2.5, 1.05 y_{n_y})]$ with y_{n_y} the highest vertex column, and $y = 1$ a node; every quoted expiry is a time node and each maturity interval receives $\max(8, \lceil \Delta \tau_{\text{int}} / 0.01 \rceil)$ implicit steps. The refined (audit) operator halves the strike step and quadruples each interval’s step count.

Box and objective. Adaptive box $v_{\text{lo}} = \min(0.0025, (\frac{1}{2} \min_e \sigma_{\text{ATM},e})^2)$, $v_{\text{hi}} = \min(\max(0.36, (3 \sigma_{\max})^2), 16)$ — the floor a fraction of the smallest ATM implied per the averaging fact, the cap a multiple of the largest implied in view. Seed and reference $\theta^0 = \theta_{\text{ref}}$ flat at the median ATM variance. Weights $\lambda = 10^{-2}$, $\lambda_0 = 1$; residual scale η_q equal to one volatility percent of the quote’s Black vega, floored at 10^{-3} . The snapshot carries no variance-swap quotes, so the wing multiple stays at its default $a = 0$ (flat clamp) throughout.

Solver. Trust-region reflective least squares [5] with the analytic Jacobian of proposition 4.12, termination tolerances 10^{-8} , at most 200 objective evaluations; the fits of section 4.8 converged in 116 (SPY), 94 (NVDA) and 20 (synthetic) evaluations from their flat seeds.

4.A.2 The synthetic market

The running example’s truth is the smooth skewed field

$$\sigma_{\text{loc}}(\tau, y) = 0.21 - 0.10 \tanh\left(\frac{y-1}{0.22}\right) + 0.03 e^{-2\tau}, \quad v = \sigma_{\text{loc}}^2,$$

priced by an *independently implemented* dense march (uniform $\Delta y = 0.0025$ on $[0, 3]$, $\Delta \tau = 0.002$, fully implicit) — a second discretization on different grids, so the round trip of fig. 4.8 cross-checks the chapter’s march rather than validating it against itself. Quotes are read at the four expiries $\tau \in \{0.10, 0.25, 0.50, 1.00\}$, at 31 strikes per expiry uniform in the standardized moneyness $\zeta \in [-2.2, 2.2]$ of section 3.7, with $\sigma_{\text{ref}} = 0.24$, i.e. $k = 0.24 \sqrt{\tau} \zeta$ (the coarse arm of fig. 4.2a uses 13). The “noise” of section 4.3 is a deterministic alternating ± 50 bp volatility ripple — no randomness anywhere in the chapter’s figures. The synthetic sheet has rows $\{0\} \cup \{\text{expiries}\}$ and the delta-spaced columns at scale 0.24 plus one deliberate deep-put column at $y = 0.45$ for the experiment of fig. 4.6; its box is $[0.005, 0.36]$.

4.A.3 The adjoint

The tangent route of proposition 4.12 prices the Jacobian at $O(N_\tau N_y n_\theta)$ per evaluation, for a march of N_τ steps on N_y lattice nodes. The classical alternative makes the cost of a *gradient* independent of n_θ ; deriving it costs a few lines, so the chapter records it for the grids on which it will one day pay.

Proposition 4.13 (Adjoint: one backward sweep). *Let $\mathcal{J}(\theta)$ be a scalar objective depending on the march through the observed values $\mathcal{O}_n U^n$, with $\mathcal{O}_n$ fixed observation rows. Define adjoint*

states by the backward recursion

$$(\mathcal{M}^n)^\top p^n = \mathcal{O}_n^\top \frac{\partial \mathcal{J}}{\partial (\mathcal{O}_n U^n)} + p^{n+1}, \quad p^{N+1} = 0. \quad (4.11)$$

Then

$$\frac{\partial \mathcal{J}}{\partial \theta_\ell} = \sum_n \Delta \tau (p^n)^\top \mathcal{L}[\mathcal{H}_\ell] U^n, \quad (4.12)$$

with the boundary-inclusive product of proposition 4.12, at the cost of one backward solve, $O(N_\tau N_y)$, plus the $O(n_\theta)$ accumulation of (4.12).

Proof. Form the Lagrangian $\mathfrak{L} = \mathcal{J} - \sum_n (p^n)^\top (\mathcal{M}^n U^n - U^{n-1} - b^n)$, with b^n the boundary terms of proposition 4.12. Stationarity in each U^n — which appears in step n 's system and in step $n+1$'s right-hand side — gives (4.11); the θ -derivative of $\mathfrak{L}$ then retains the explicit dependence of $\mathcal{M}^n$ and b^n , which combine as in proposition 4.12 into the boundary-inclusive (4.12). Structurally this is summation by parts: the transpose of a lower block-bidiagonal system is upper block-bidiagonal, i.e. a backward march with transposed tridiagonal factors. $\square$

4.A.4 Measured cleanliness

The diagnostics quoted in section 4.8 are computed on the fitting operator's price grid: the butterfly proxy is the minimum divided second difference $(U_{i+1} - U_i)/\Delta y_i^+ - (U_i - U_{i-1})/\Delta y_i^-$ over all interior nodes and all expiries; the calendar proxy is the minimum of $U^{(e+1)} - U^{(e)}$ over adjacent expiries at fixed y (the lattice prices are already in the normalized units of section 2.8, so the fixed- y comparison is the right one); and the bounds check verifies $0 \leq U \leq 1$ to the same tolerance. For both calibrated sheets all three sit at solver rounding (section 4.8); they are measured per fit precisely because remark 4.11 declines to promote them to a theorem.

Bibliography

- [1] Jesper Andreasen and Brian Huge. Volatility interpolation. *Risk*, pages 76–79, March 2011.
- [2] Shalom Benaim and Peter Friz. Regular variation and smile asymptotics. *Mathematical Finance*, 19(1):1–12, 2009.
- [3] Abraham Berman and Robert J. Plemmons. *Nonnegative Matrices in the Mathematical Sciences*. SIAM, Philadelphia, 1994.
- [4] Fischer Black. The pricing of commodity contracts. *Journal of Financial Economics*, 3(1–2):167–179, 1976.
- [5] Mary Ann Branch, Thomas F. Coleman, and Yuying Li. A subspace, interior, and conjugate gradient method for large-scale bound-constrained minimization problems. *SIAM Journal on Scientific Computing*, 21(1):1–23, 1999.
- [6] Douglas T. Breeden and Robert H. Litzenberger. Prices of state-contingent claims implicit in option prices. *Journal of Business*, 51(4):621–651, 1978.
- [7] Kresimir Demeterfi, Emanuel Derman, Michael Kamal, and Joseph Zou. A guide to volatility and variance swaps. *Journal of Derivatives*, 6(4):9–32, 1999.
- [8] Bruno Dupire. Pricing with a smile. *Risk*, 7(1):18–20, 1994.
- [9] Matthias R. Fengler. Arbitrage-free smoothing of the implied volatility surface. *Quantitative Finance*, 9(4):417–428, 2009.
- [10] Frederick N. Fritsch and Ralph E. Carlson. Monotone piecewise cubic interpolation. *SIAM Journal on Numerical Analysis*, 17(2):238–246, 1980.
- [11] Jim Gatheral. A parsimonious arbitrage-free implied volatility parameterization with application to the valuation of volatility derivatives. Presentation at Global Derivatives & Risk Management, Madrid, 2004.
- [12] Jim Gatheral. *The Volatility Surface: A Practitioner’s Guide*. Wiley, Hoboken, 2006.
- [13] Jim Gatheral and Antoine Jacquier. Arbitrage-free SVI volatility surfaces. *Quantitative Finance*, 14(1):59–71, 2014.
- [14] Warren G. Gilchrist. *Statistical Modelling with Quantile Functions*. Chapman & Hall/CRC, Boca Raton, 2000.
- [15] Michael B. Giles and Rebecca Carter. Convergence analysis of Crank–Nicolson and Ranacher time-marching. *Journal of Computational Finance*, 9(4):89–112, 2006.
- [16] István Gyöngy. Mimicking the one-dimensional marginal distributions of processes having an Itô differential. *Probability Theory and Related Fields*, 71(4):501–516, 1986.

- [17] Thomas W. Keelin. The metalog distributions. *Decision Analysis*, 13(4):243–277, 2016.
- [18] Hans G. Kellerer. Markov-Komposition und eine Anwendung auf Martingale. *Mathematische Annalen*, 198:99–122, 1972.
- [19] Roger W. Lee. The moment formula for implied volatility at extreme strikes. *Mathematical Finance*, 14(3):469–480, 2004.
- [20] Claude Martini and Arianna Mingone. No arbitrage SVI. *SIAM Journal on Financial Mathematics*, 13(1):227–261, 2022.
- [21] Anthony Neuberger. The log contract. *Journal of Portfolio Management*, 20(2):74–80, 1994.
- [22] Emanuel Parzen. Nonparametric statistical data modeling. *Journal of the American Statistical Association*, 74(365):105–121, 1979.
- [23] Alexander Petersen and Hans-Georg Müller. Functional data analysis for density functions by transformation to a Hilbert space. *Annals of Statistics*, 44(1):183–218, 2016.
- [24] Volker Strassen. The existence of probability measures with given marginals. *Annals of Mathematical Statistics*, 36(2):423–439, 1965.